

3150 - Operating Systems

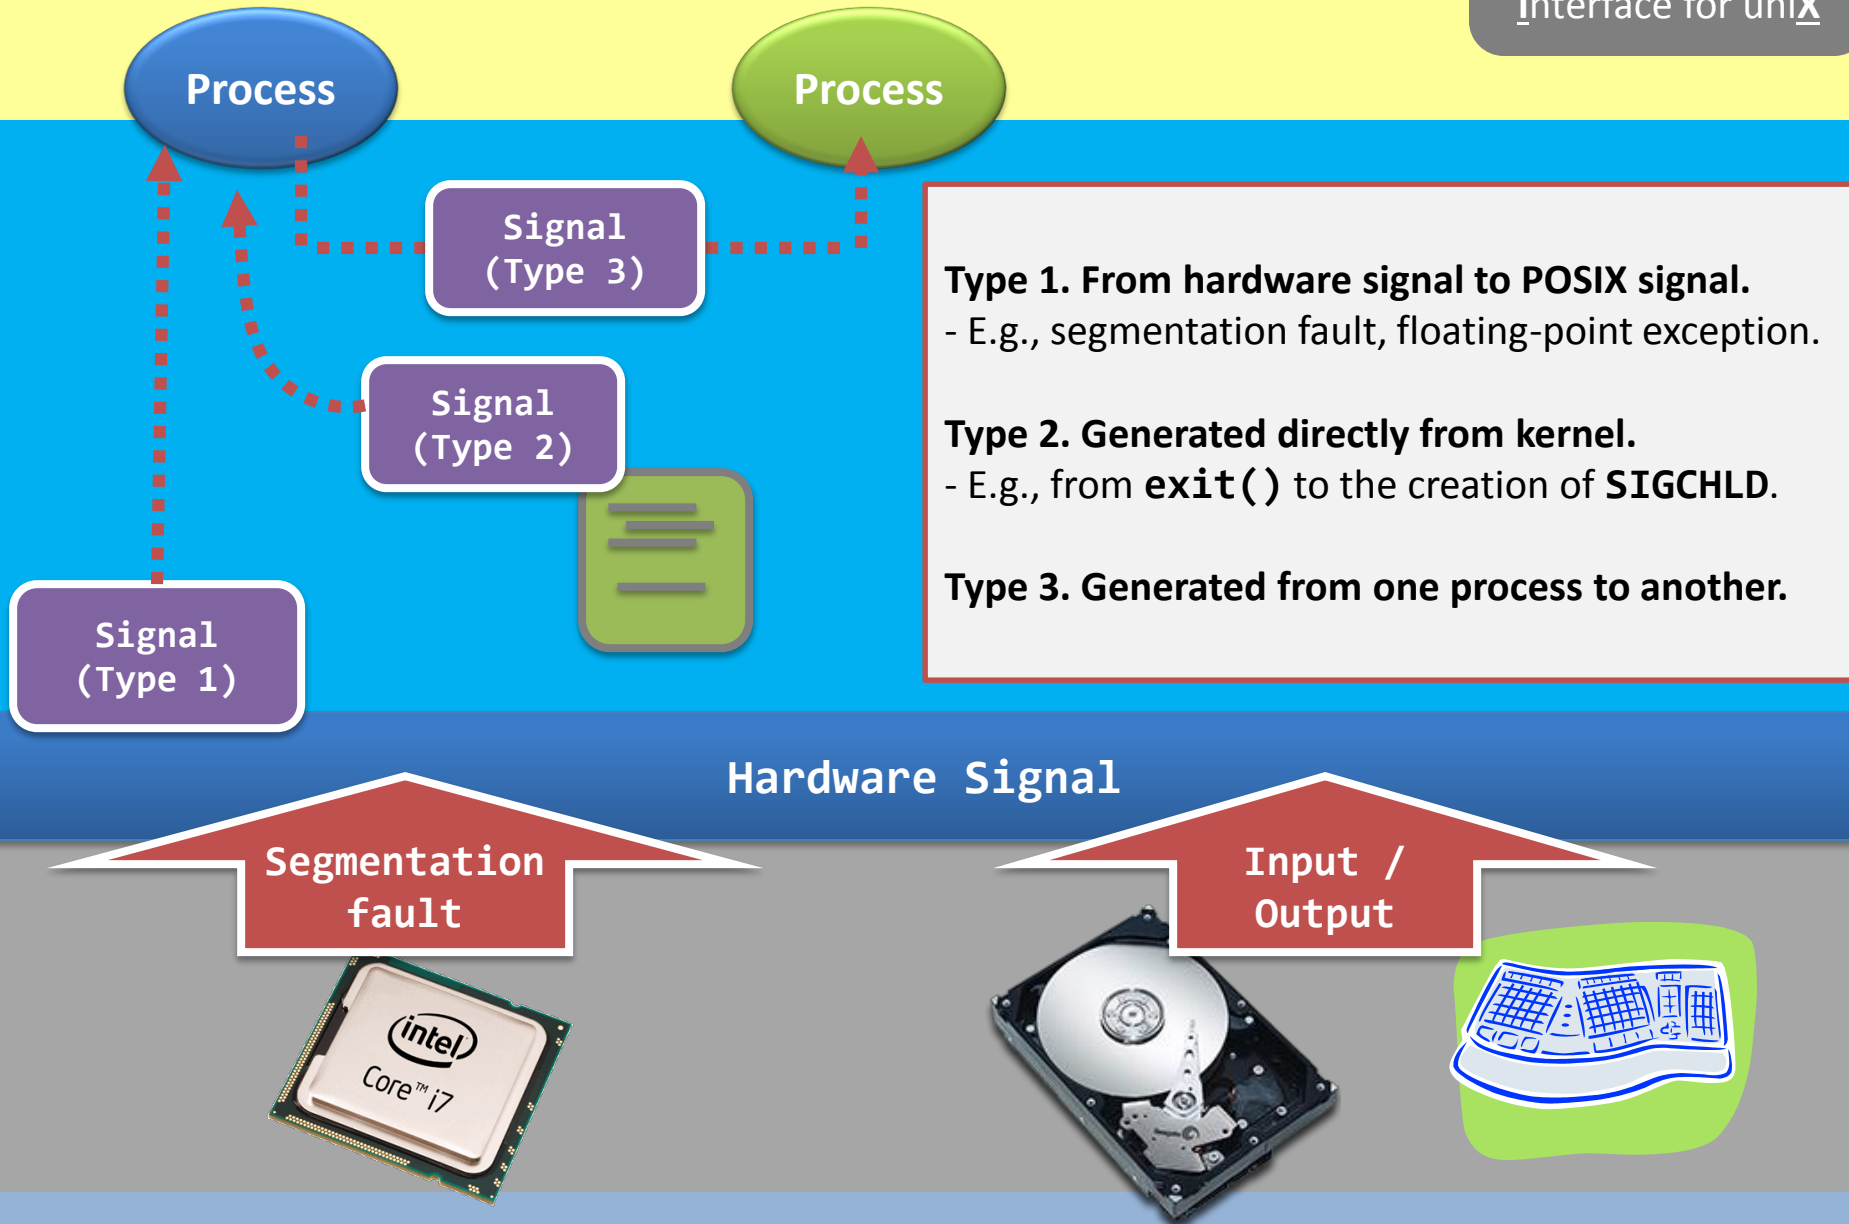
Dr. WONG Tsz Yeung

Chapter 2, part 3 – POSIX Signals

- Let's learn how to interrupt a process (and have fun)

Outline

POSIX – Portable
Operating System
Interface for uniX



Signals

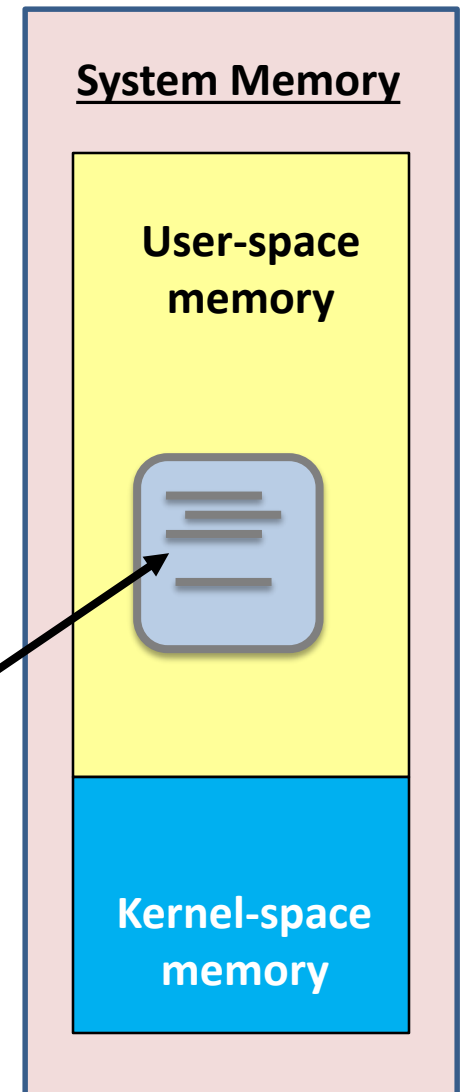
-preface: hardware signals;



Hardware signal

- It is the **hardware interrupt**.
 - CS students may not know what an interrupt is...

Originally, the CPU is working on a program code.



Hardware signal

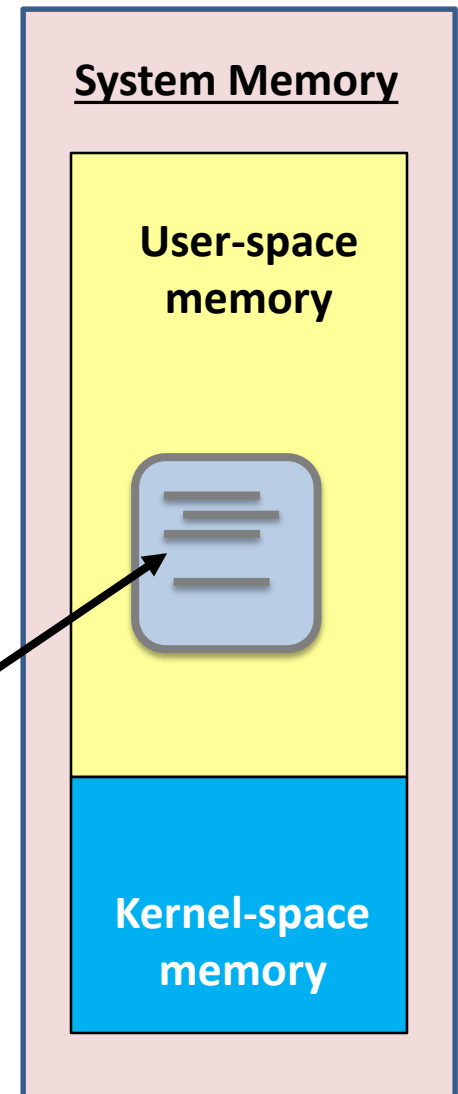
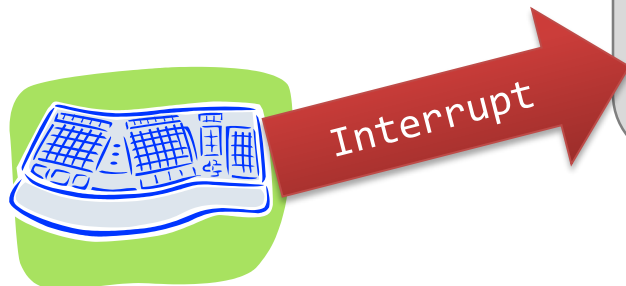
- It is the **hardware interrupt**.
 - CS students may not know what an interrupt is...

Suddenly, someone types things on to the keyboard.

This generates a hardware-**interrupt** that *interrupts* the execution of the CPU.

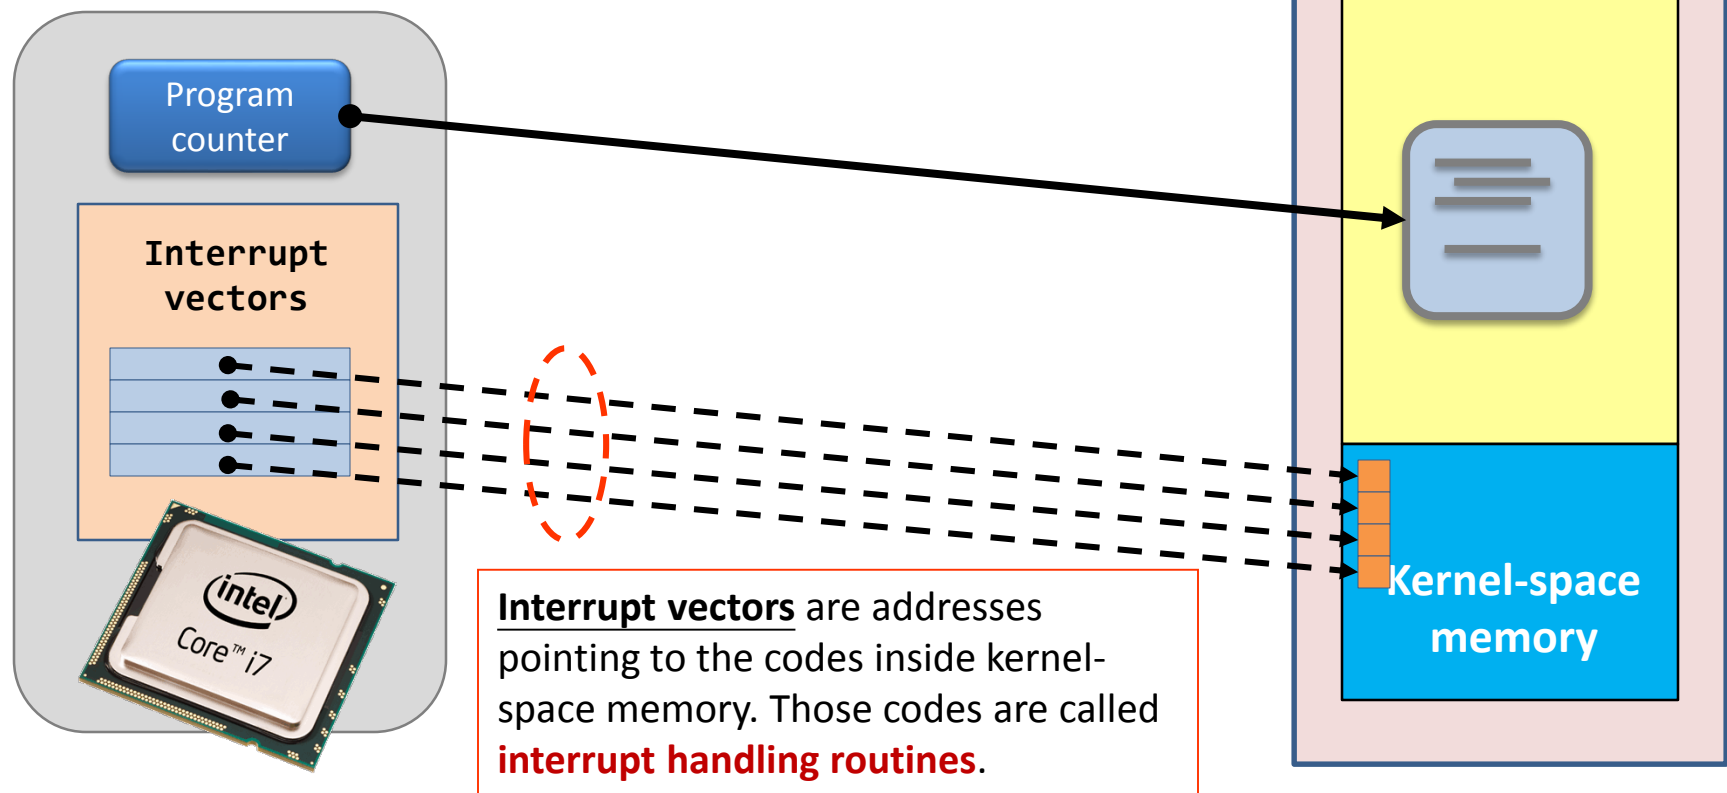
noun

verb



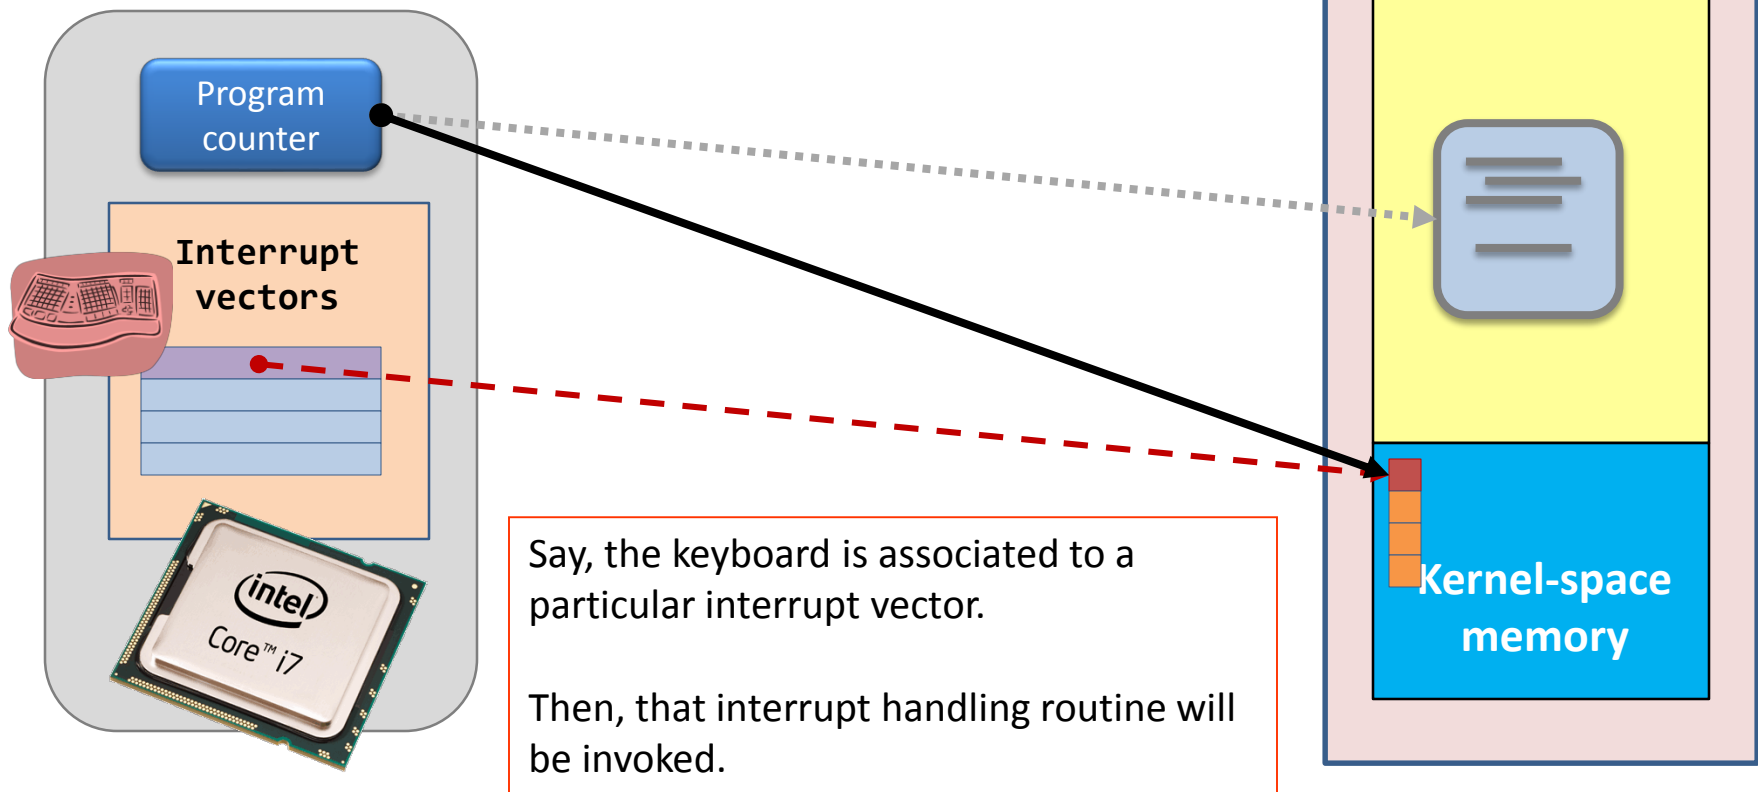
Hardware signal

- It is the **hardware interrupt**.
 - CS students may not know what an interrupt is...



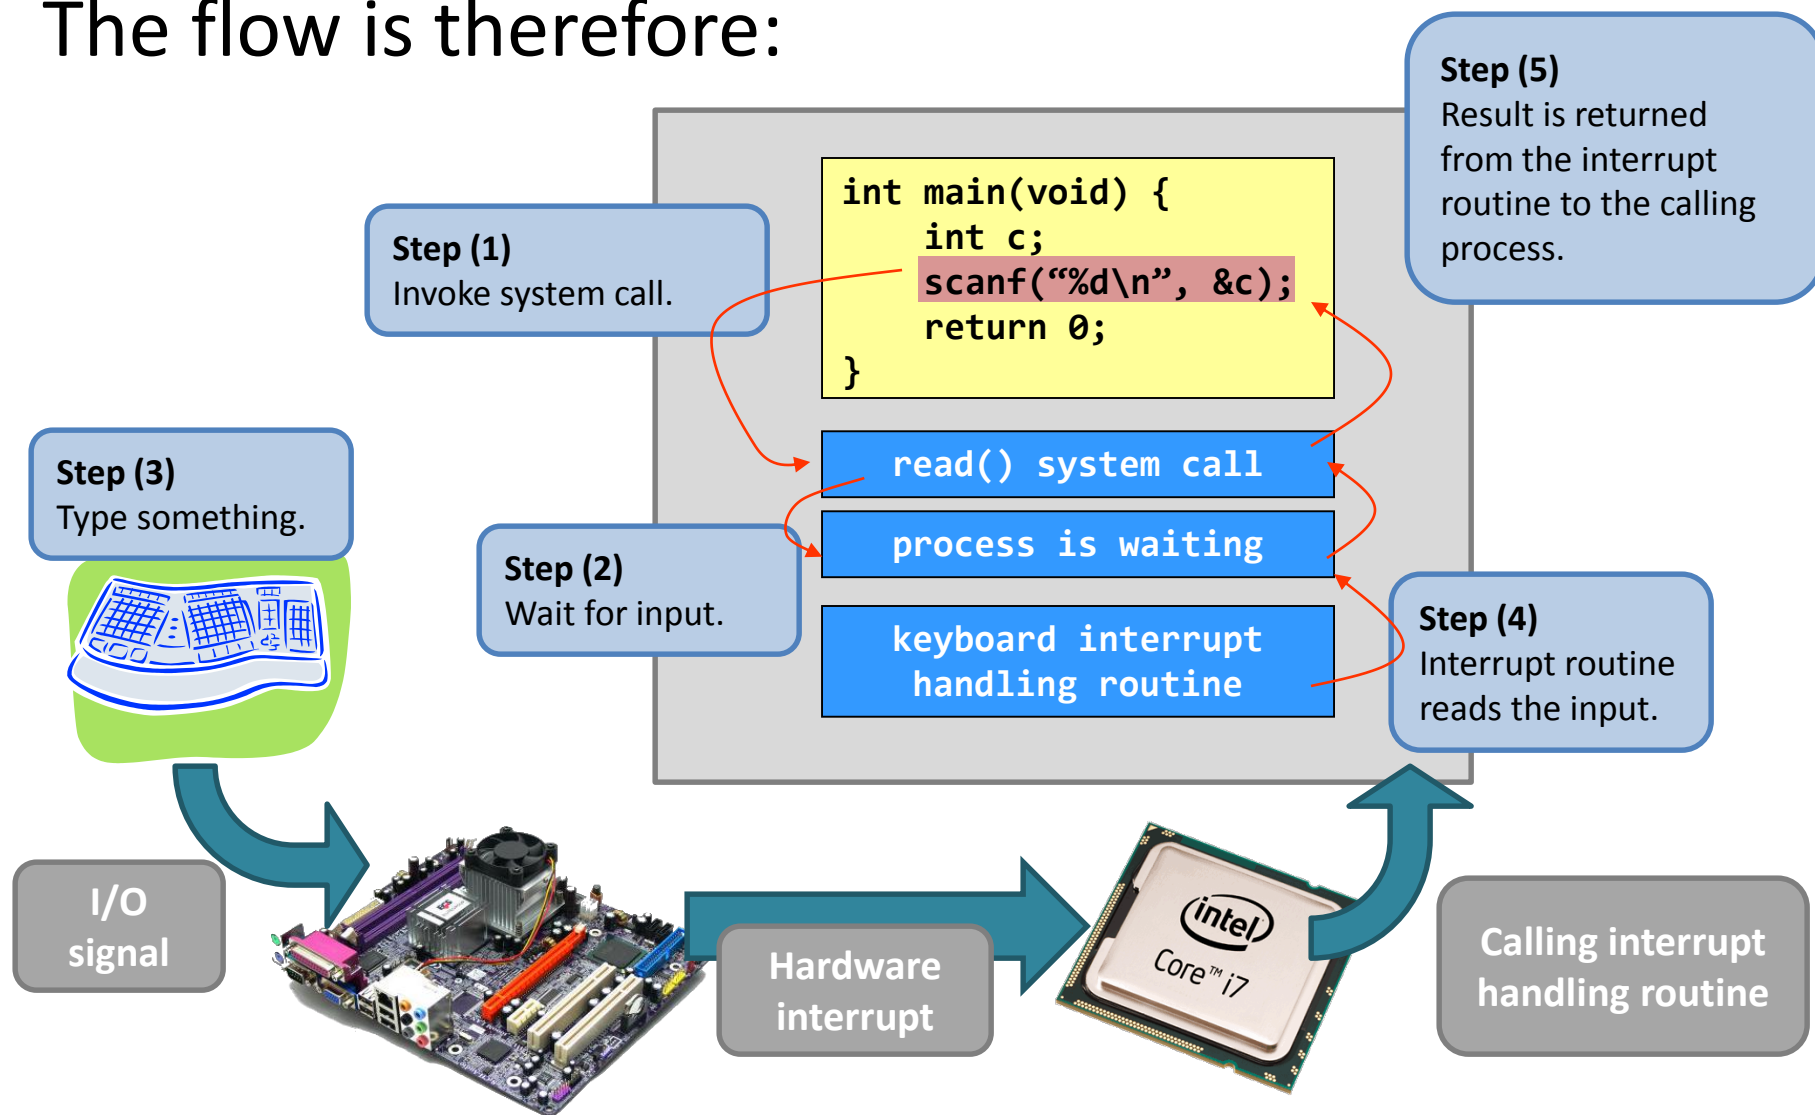
Hardware signal

- It is the **hardware interrupt**.
 - CS students may not know what an interrupt is...



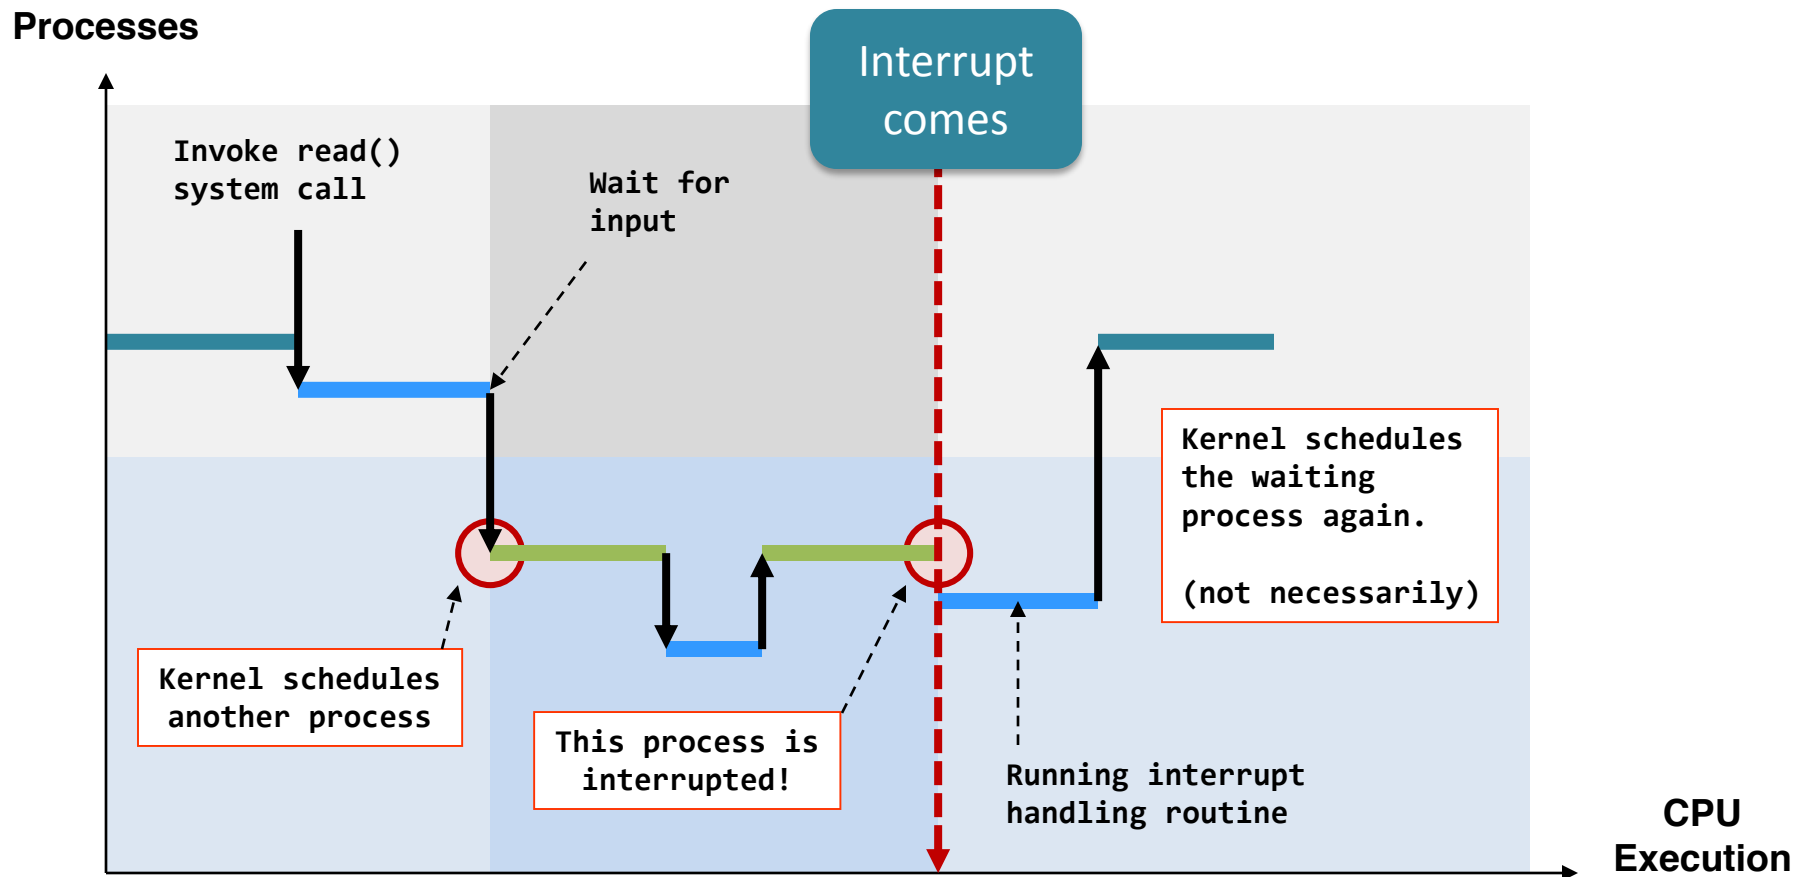
Hardware signal

- The flow is therefore:



Hardware signal

- Don't think that the kernel will be waiting for the keyboard input **forever**!



Signals

- preface: hardware signals;
- POSIX signals basics;**

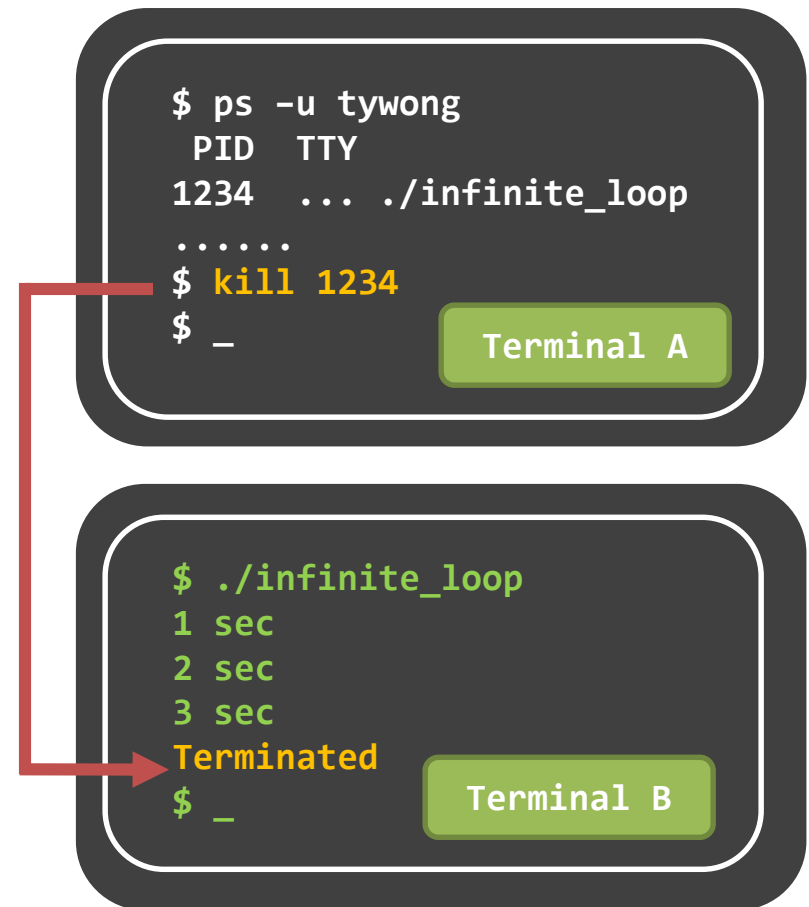
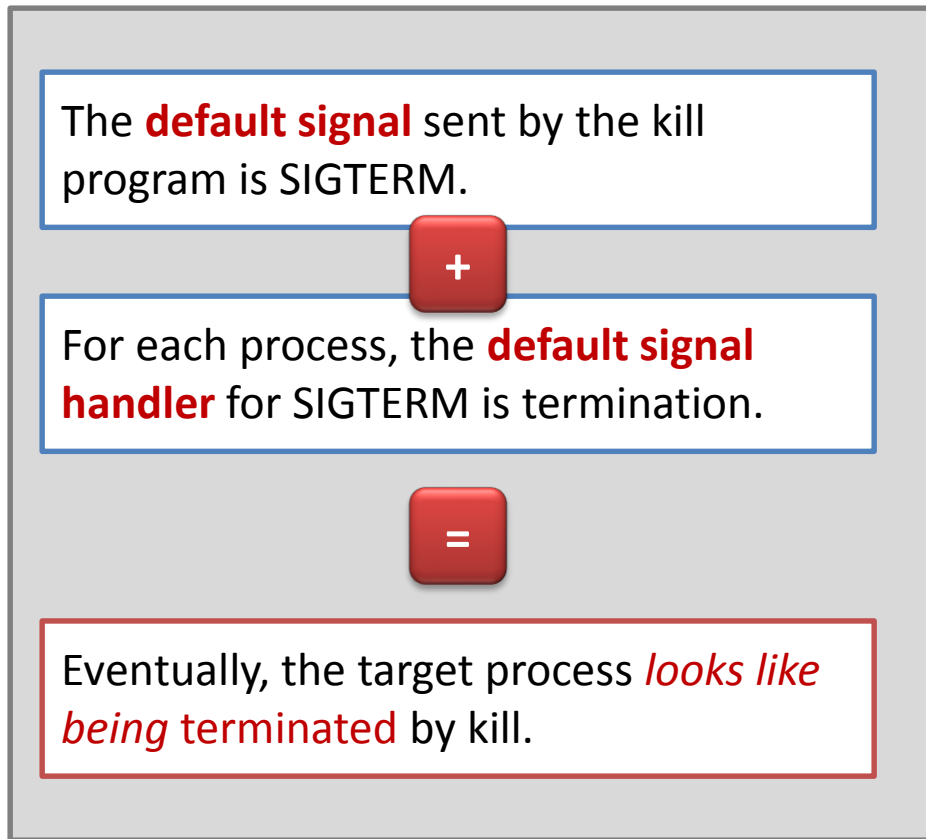


What is a POSIX signal?

Generated from one process to another	Generated from CPU to kernel, or CPU to processes
From terminals: E.g., “ Ctrl + C ”, “ Ctrl + Z ” and “ Ctrl + \ ”	<u>Example 1: Segmentation fault.</u> The signal is labeled SIGSEGV , which comes from CPU to kernel then to processes.
Using programs: E.g., “ kill ”, “ top ”, etc.	<u>Example 2: Floating point exception.</u> The signal is labeled SIGFPE , which is coming from CPU when a user-level code makes things wrong: division by zero.
Using the “ kill() ” system call.	<u>Example 3: Child process termination.</u> The signal is labeled SIGCHLD , which is coming from the kernel; CPU is not involved.

Send signals using kill program

- The kill program is to **send signals** to target processes.



Some default things...

Signal	Description	Default signal handler
SIGINT	Its name is the interrupt signal . Can be generated by “ Ctrl + C ”.	Target process termination.
SIGTERM	Its name is the termination signal . The default signal sent by the “ kill ” program.	Target process termination.
SIGTSTP	Its name is the terminal stop signal . Can be generated by “ Ctrl + Z ”.	Target process suspension.
SIGCONT	Its name is the continuation signal . Will discuss this later.	Target process resumes execution if it is previously suspended.
SIGCHLD	(No special name). It is sent to the parent process to a terminated child.	Ignore by default. <i>unless use wait()</i>
SIGKILL	Its name is the kill signal . If sent, the process MUST DIE .	Target process termination. (Plus, no one could stop the termination.)

What are SIGTSTP & SIGCONT?

```
$ ps -u tywong
PID  TTY
1234  ... ./infinite_loop
.....
$ kill -TSTP 1234
$ kill -CONT 1234
```

SIGTSTP

SIGCONT

The “terminal stop signal” stops the process.

Terminal A

```
$ ./infinite_loop
1 sec
2 sec
3 sec
[1] Stopped ./infinite_loop
$ 4 sec
5 sec
6 sec
```

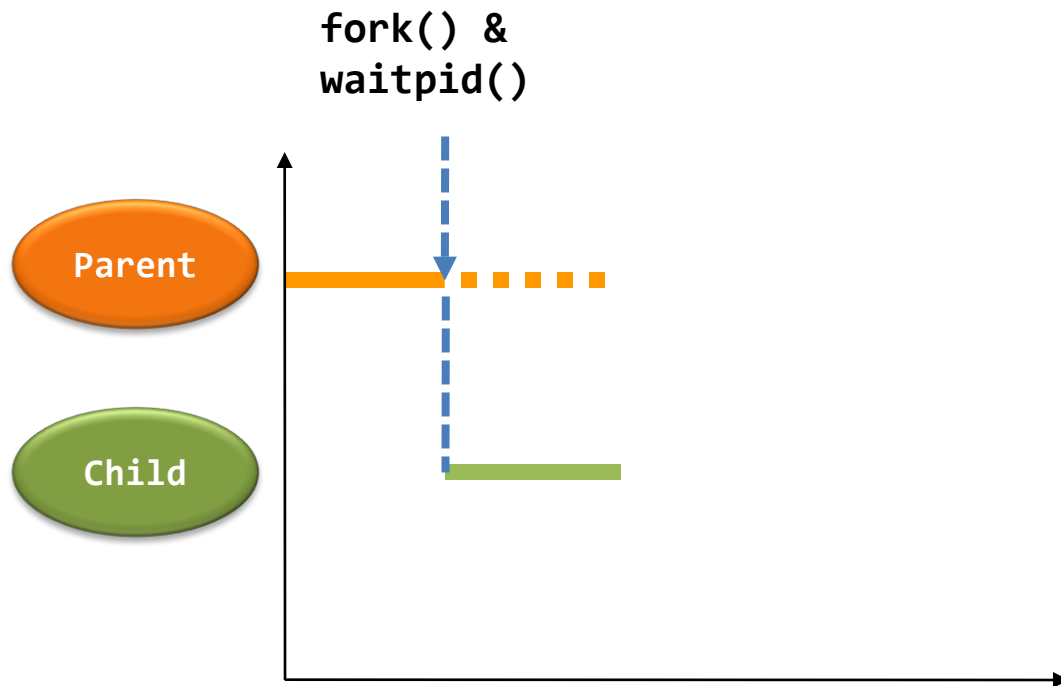
What happened to the shell?!

The “continuation signal” resumes the process. But, wait....

Terminal B

Foreground and background jobs

A characteristic of a foreground job is shell is:
The shell is waiting for the job to change state.



```
$ ./infinite_loop
```

```
1 sec
```

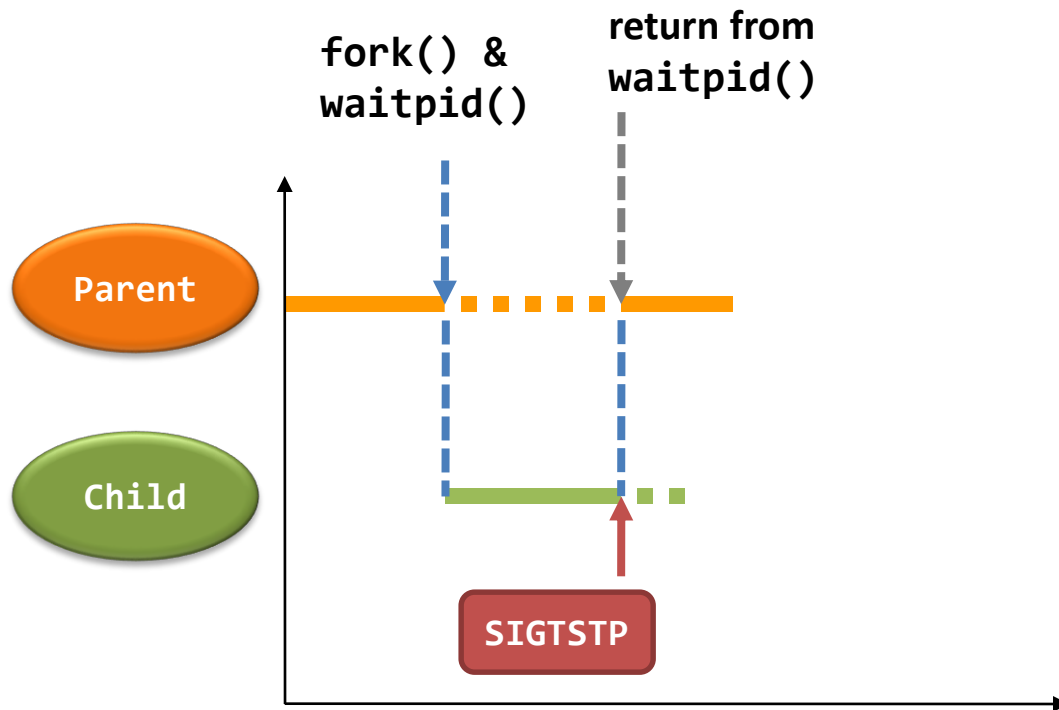
```
2 sec
```

```
3 sec
```

We name this a
foreground job.

Foreground and background jobs

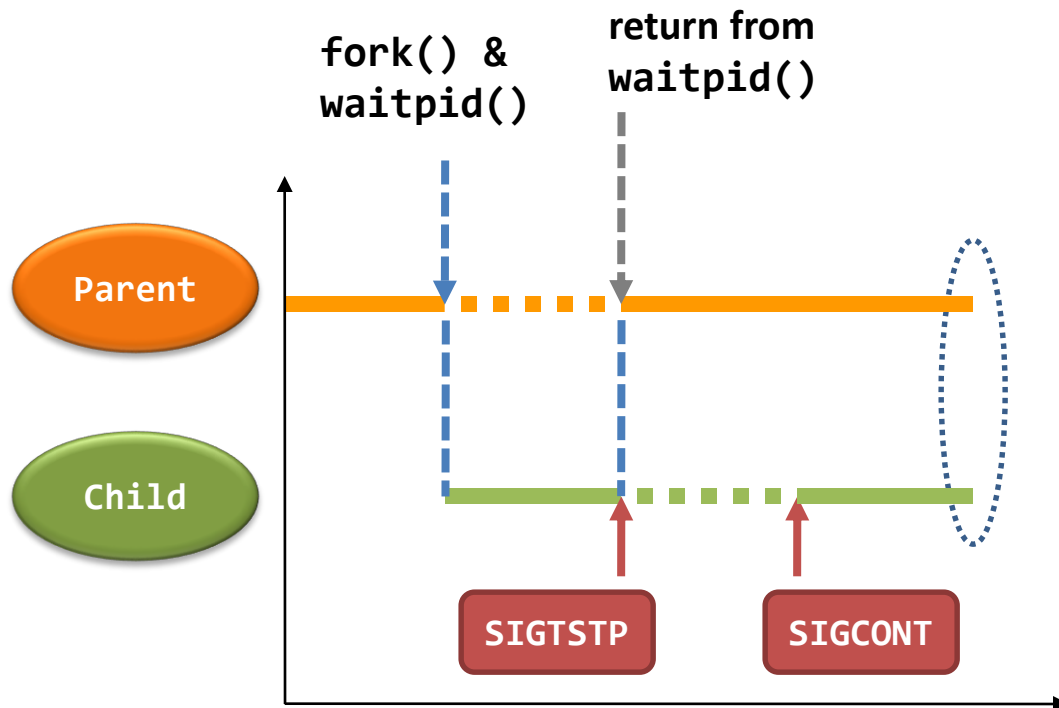
Now, the shell wakes up. By the way, “[1] Stopped ...” is actually printed out by the shell!



```
$ ./infinite_loop
1 sec
2 sec
3 sec
[1] Stopped ...
$ _
```


Foreground and background jobs

SIGCONT is fired from another shell. Since the parent shell doesn't know about this fact, they both just run in parallel.



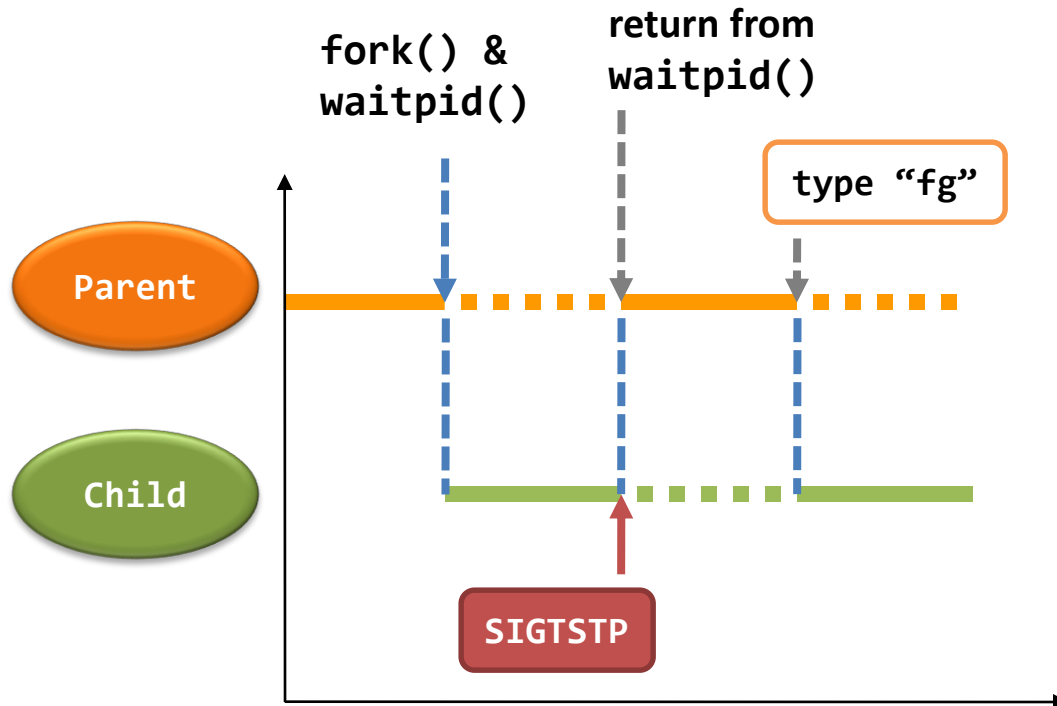
```
$ ./infinite_loop
1 sec
2 sec
3 sec
[1] Stopped ...
$ 4 sec
5 sec
6 sec
```

Now, the job becomes
a background job.

Foreground and background jobs

DISCUSSION

The question is simple: how to implement the command “fg”?



```
$ ./infinite_loop
1 sec
2 sec
3 sec
[1] Stopped ...
$ fg
4 sec
5 sec
6 sec
```

Now, the job becomes
a foreground job.

Sending signals in a process

- Remember, **kill()** is not intended to kill anybody, but to send signals.

```
int kill(pid_t pid, int sig)
```

PID of
the target
process

The desired signal. For the complete
list, please read the manpages:

On Linux:

man 7 signal

On Unix:

man -s 3HEAD signal

```
1 int main(void) {
2     int i, sum;
3     srand(time(NULL));
4     while(1) {
5         sum = 0;
6         for(i = 0; i < 3; i++)
7             sum += (rand() % 6) + 1;
8         if(sum == 18)
9             kill(getpid(), SIGTERM);
10    }
11    return 0;
12 }
```

An alternative: **raise(SIGTERM);**

[examples@3150] cat suicide.c

Signals

- Preface: hardware signals;
- POSIX signals basics;
- Handling POSIX signals;**



Setting new signal handlers

- What is the meaning?
 - The process no longer executes the default handler...
 - The signal handlers can be changed to an user-level function!

```
signal( int signum, void (* function)(int) )
```

The signal number, e.g., **SIGINT** or **SIGTERM**.

A function pointer to the new handler. This means you should provide the name of a function, with:

- return type of “**void**” and
- one argument with type “**int**”

Two special values:

- **SIG_DFL**: set to the default handler.
- **SIG_IGN**: ignore this signal completely.

Setting new signal handlers

```
1 void sig_handler(int sig) {  
2     if(sig == SIGINT)  
3         printf("\nCtrl + C\n");  
4 }  
5  
6 int main(void) {  
7     signal(SIGINT, sig_handler);  
8     printf("Press enter\n");  
9     getchar();  
10    printf("End of program\n");  
11 }
```

Line 7 registers the signal handler when **SIGINT** is received.

Lines 1-4 together define the signal handler.

[examples@3150] cat handle_int.c

Setting new signal handlers

```
1 void sig_handler(int sig) {  
2     if(sig == SIGINT)  
3         printf("\nCtrl + C\n");  
4 }  
5  
6 int main(void) {  
7     signal(SIGINT, sig_handler);  
8     printf("Press enter\n");  
9     getchar();  
10    printf("End of program\n");  
11 }
```

```
$ ./handle_int  
Press enter  
^C  
Ctrl + C
```

```
[examples@3150] cat handle_int.c
```

Setting new signal handlers

EXTRA & IMPORTANT

- An important point to note:
 - Apparently, when a signal handler returns, **the process goes back to where it was executing.**
 - But...

```
1 void sig_handler(int sig) {
2     printf("\nSignal received.\n");
3 }
4
5 int main(void) {
6     signal(SIGINT, sig_handler);
7     printf("Sleep for 24 hours\n");
8     sleep(24 * 60 * 60);
9     printf("Wake up and die.\n");
10 }
```

```
$ ./break_sleep
Sleep for 24 hours
^C
Signal received.
Wake up and die
$_
```



[examples@3150] cat break_sleep.c

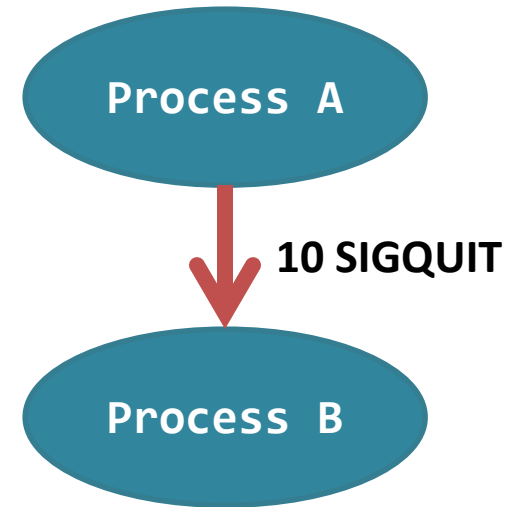
- An important point to note:
 - **Apparently**, when a signal handler returns, **the process goes back to where it was executing**.
 - But...this only happens when **the involved system/library call can be restarted automatically**.

Can be restarted	Cannot be restarted
<p>[file related] <code>open(), read(), write();</code></p> <p>[process related] <code>wait(), waitpid();</code></p>	<p><code>sleep(); pause();</code></p> <p>With dozens of calls that you may not meet before...</p>

Counting Signal Received?

EXTRA & IMPORTANT

```
1 void handler(int sig) {
2     static int count = 0;
3     printf("count = %d\n", ++count);
4 }
5
6 int main(void) {
7     int i;
8     if( fork() == 0 ) {
9         printf("Press Enter...\n");
10        while(getchar() != '\n');
11        for(i = 0; i < 10; i++)
12            kill(getppid(), SIGQUIT);
13    }
14    else {
15        signal(SIGQUIT, handler);
16        sleep(1000);
17        wait(NULL);
18    }
19    return 0;
20 }
```



```
$ ./many_signal
Press Enter...
[ENTER]
count = 1
$_
```




[examples@3150] cat many_signal.c

Counting Signal Received?

EXTRA & IMPORTANT

- Another important note:
 - Signals are not queued!
 - An array is used for indicating a signal has received or not.



SIG	HUP	INT	QUIT	...	KILL	SEGV	...
BIT	0	0	1	...	0	0	

- In the previous example, the bit (or mask) is always 1 no matter how many signals are sent.
- The mask will be set to 0 when the signal is handled.
- **Guess** what will happen with **1 million** SIGQUIT sent?

- Write programs that try to handle the signals:
 - SIGSEGV; SEGFPE;
 - Question: are the errors really go away?
 - SIGCHLD;
 - Question: Is there any zombie left in the system?

Misc. Topics


- Waiting for signals;**
- Breaking out of loops;**
- Timers and periodic signals;**

(1) - Waiting for a signal

- The **pause()** system call suspends the calling process until...
 - a signal which is **handled** by the process is received, or
 - a signal which **terminates** the process is received.

It suspends the execution of the program until a signal is caught...

Of course, **pause()** is designed not to be restarting after a signal handler.



```
1 void sig_handler(int sig) {  
2 }  
3  
4 int main(void) {  
5     signal(SIGINT, sig_handler);  
6     pause();  
7     printf("Ctrl+C received. Bye!\n");  
8     return 0;  
9 }
```

[examples@3150] cat pause.c

(2) – Breaking out of loop

```
int stop = 0;

void sig_handler(int sig) {
    stop = 1;
}

int main(void) {
    unsigned int i = 0;
    signal(SIGINT, sig_handler);

    while( !stop ) {
        sleep(1);
        printf("%d sec\n", ++i);
        fflush(stdout);
    }

    printf("Exit peacefully\n");
    return 0;
}
```

The while loop is an infinite one until the user presses “**Ctrl+C**”. We usually call that feature: the **graceful termination**.

It allows the program to **exit normally** so that it has a chance to do things such as:

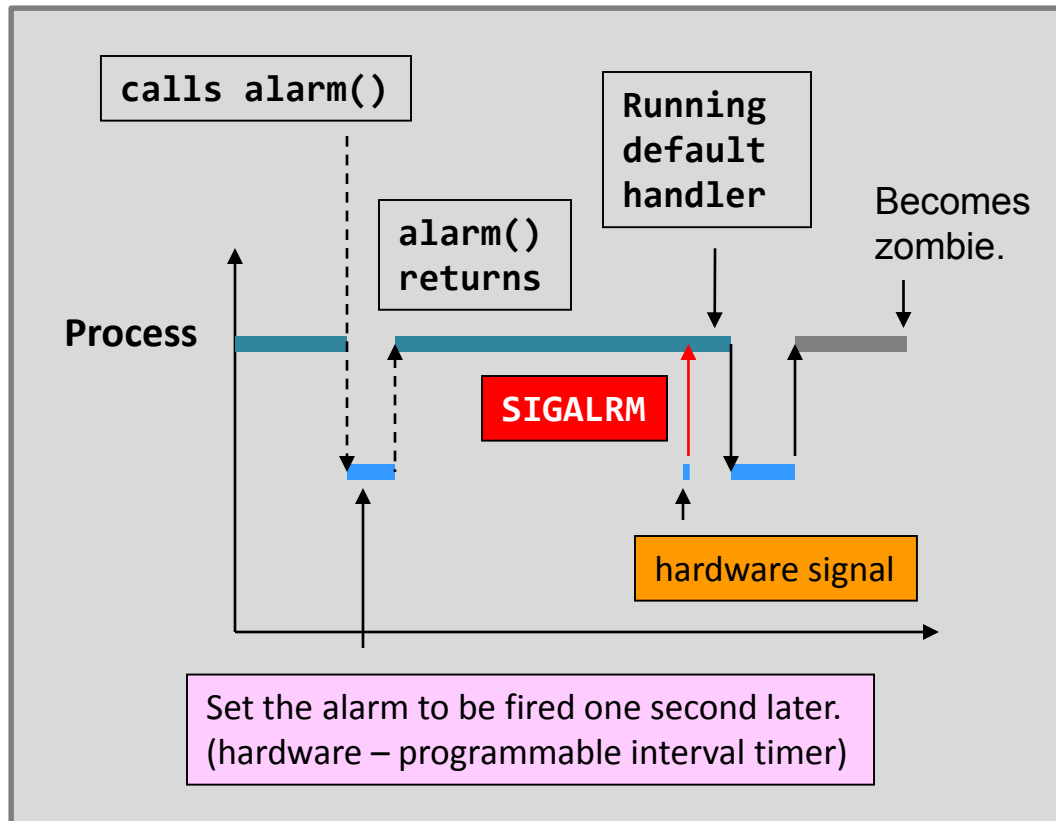
- closing network connections,
- committing database changes,
- etc.

rather than terminating the process forcefully.

[examples@3150] cat infinite_loop.c

(3) – Timer and alarm()

- **alarm()** is a system call that allows **asynchronous timing** for a process.



```
int main(void) {  
    alarm(1);  
    while(1);  
    return 0;  
}
```

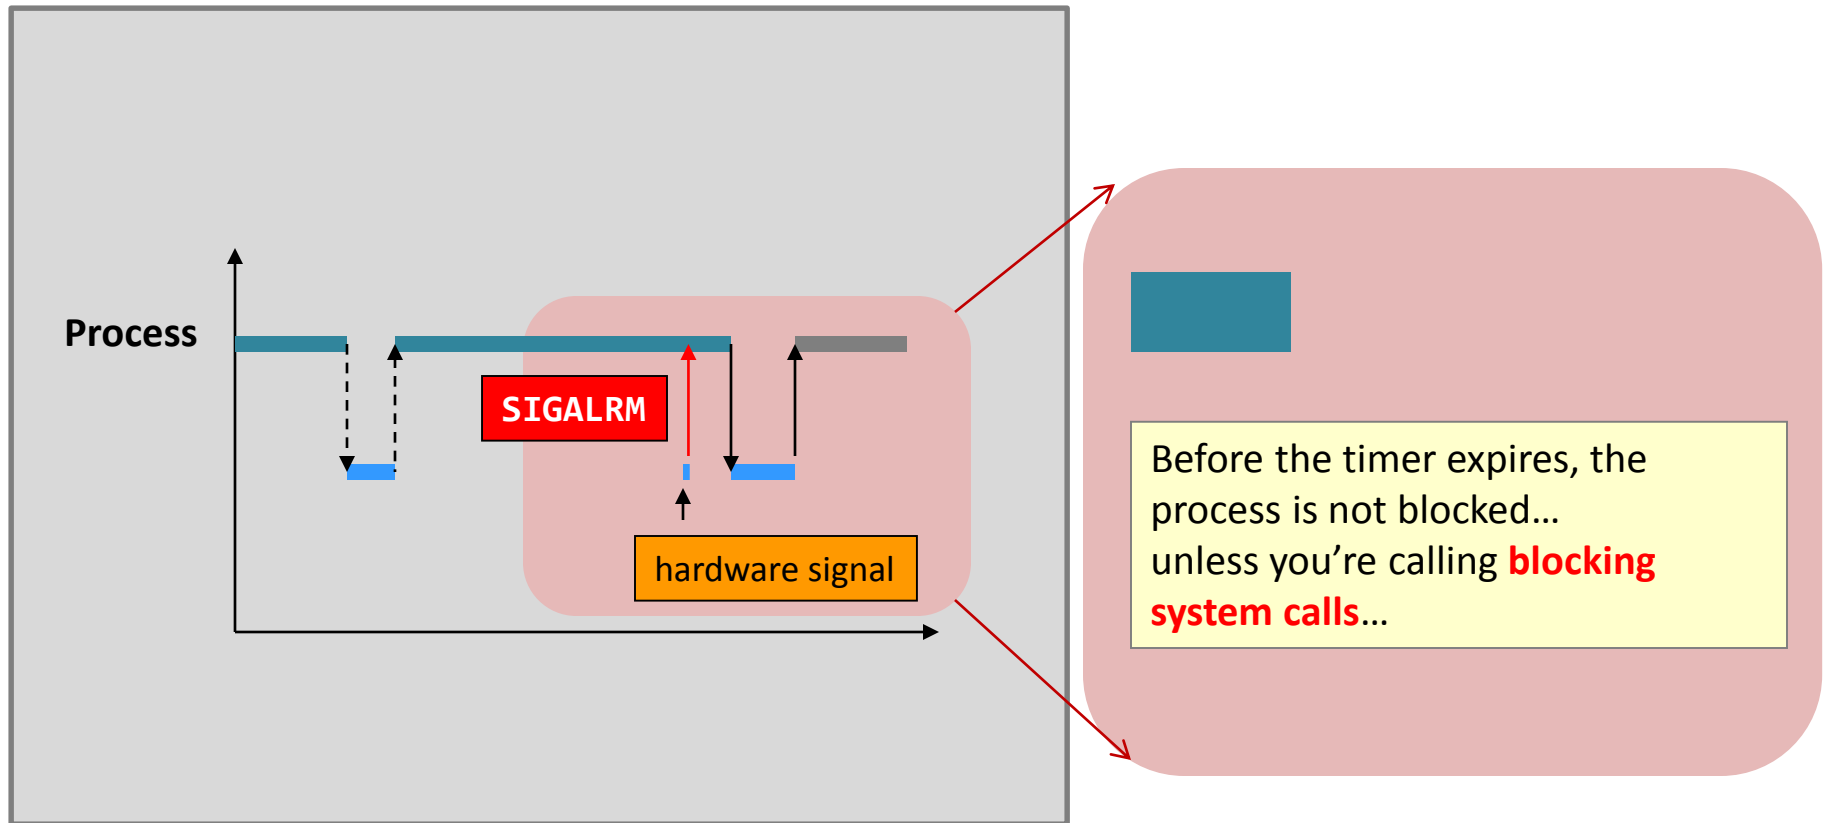
```
$ ./alarm  
Alarm clock  
$ _
```

1 sec later

[examples@3150] cat alarm.c

(3) – Timer and alarm()

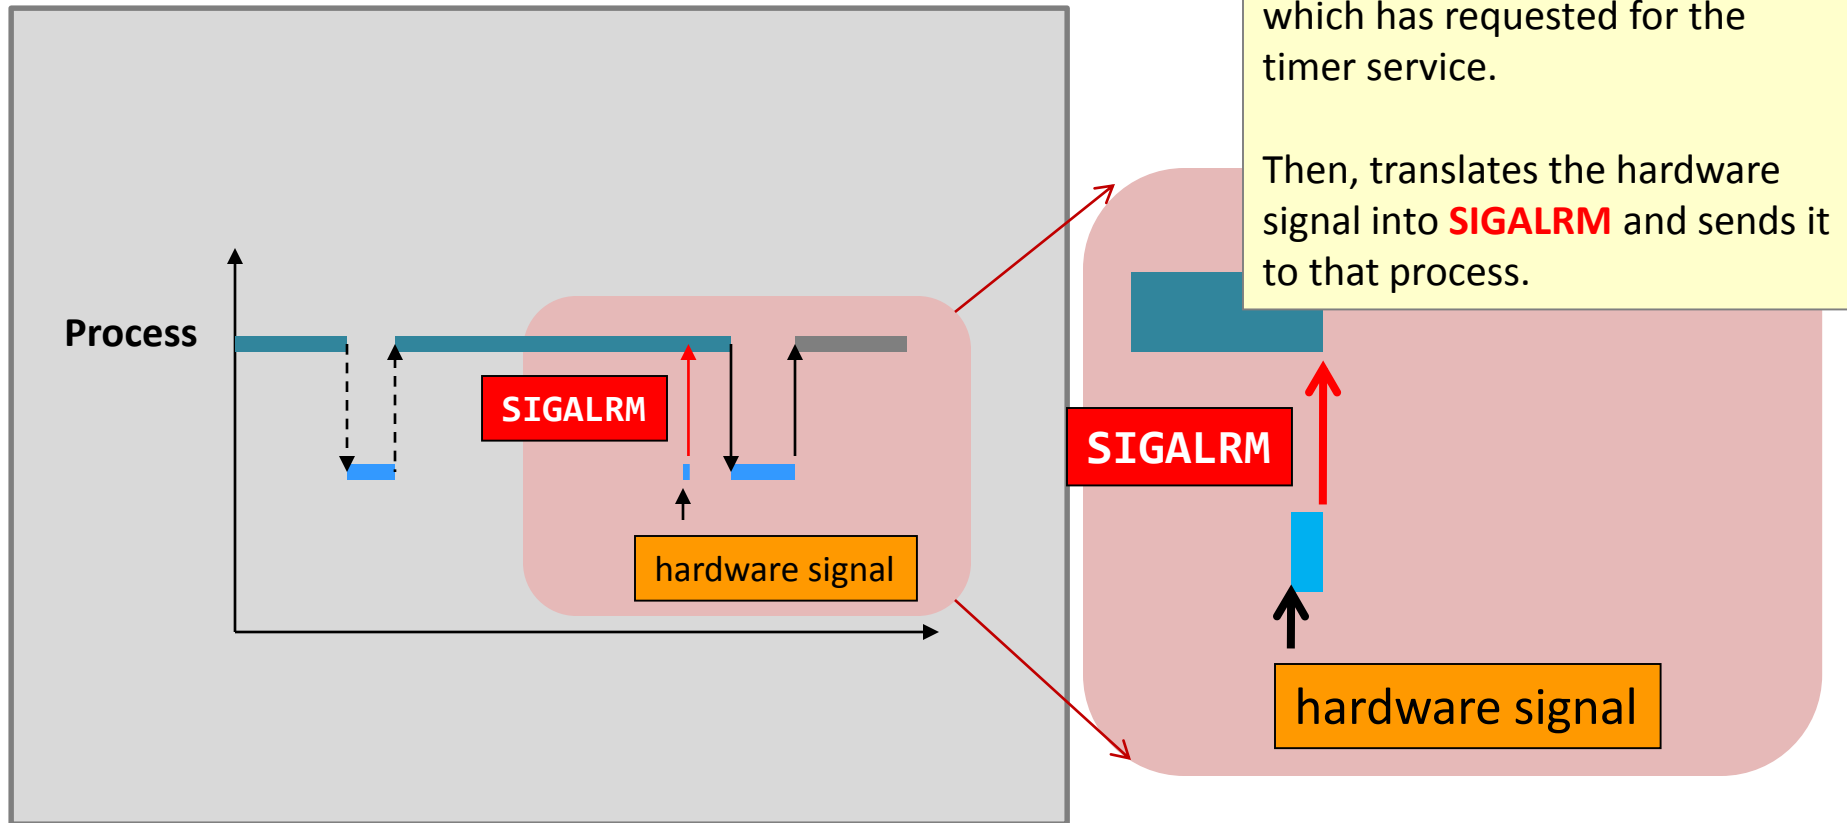
- **alarm()** is a system call that allows **asynchronous timing** for a process.



```
[examples@3150] cat alarm.c
```

(3) – Timer and alarm()

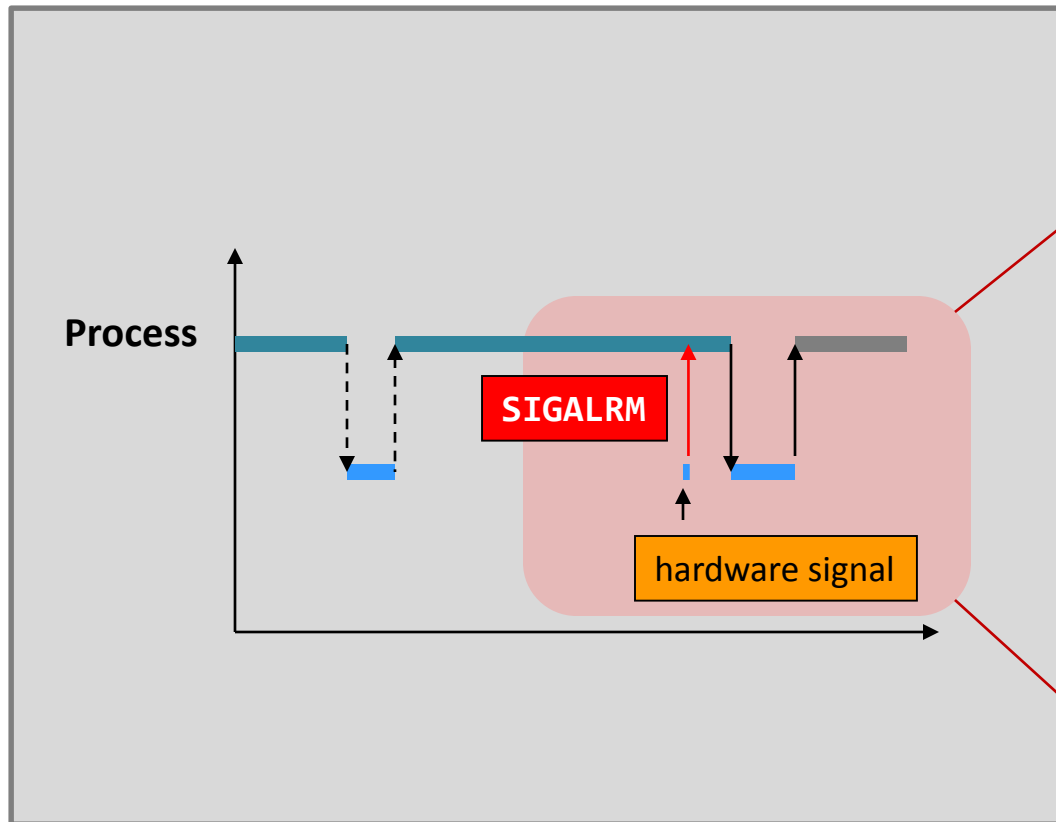
- **alarm()** is a system call that allows **asynchronous timing** for a process.



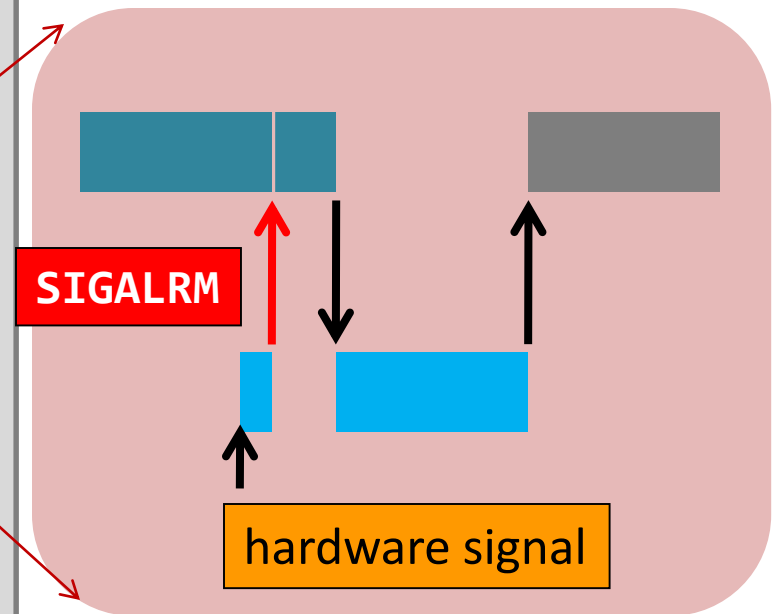
```
[examples@3150] cat alarm.c
```

(3) – Timer and alarm()

- **alarm()** is a system call that allows **asynchronous timing** for a process.



Upon the arrival of the signal, the default signal handling routine is called, which is **termination**!



```
[examples@3150] cat alarm.c
```

(3) – Timer and alarm()

- Of course, you can implement something more meaningful.

Guess: what will happen with this **exit()** call.

Listen! You've only 5 seconds to finish your typing!

This **cancels** the scheduled clock interrupt!

```
void sig_handler(int sig) {  
    printf("\nTimeout! Goodbye!\n");  
    exit(0);  
}  
  
int main(void) {  
    char buf[1024];  
    signal(SIGALRM, sig_handler);  
    alarm(5);  
    if(fgets(buf, 1024, stdin) == NULL) {  
        printf("No input. Goodbye!\n");  
        exit(0);  
    }  
    alarm(0);  
    printf("Your input: %s", buf);  
}
```

[examples@3150] cat alarm_fgets.c

(3) – Timer and alarm()

- Remember, “**alarm()**” only fires once!
 - What if I want **periodic signals**?
- How about calling **alarm()** again in a signal handler? Or...
- “**setitimer()**” (set interval timer) can help you.
 - Her sibling is “**getitimer()**”.
 - Read the manpage by yourself.

Summary

- Signal is a kind of interrupts...
 - This is quite hard to master...
 - and the course instructor still needs to refer to man pages before explaining special behaviors of some signals.
 - It is the source of the many fancy (or evil) scenarios.
- You may need a reference book if you want to go deeper:
 - *Advanced Programming Environment in UNIX*;
 - “**man 7 signal**” is a vast resource.