

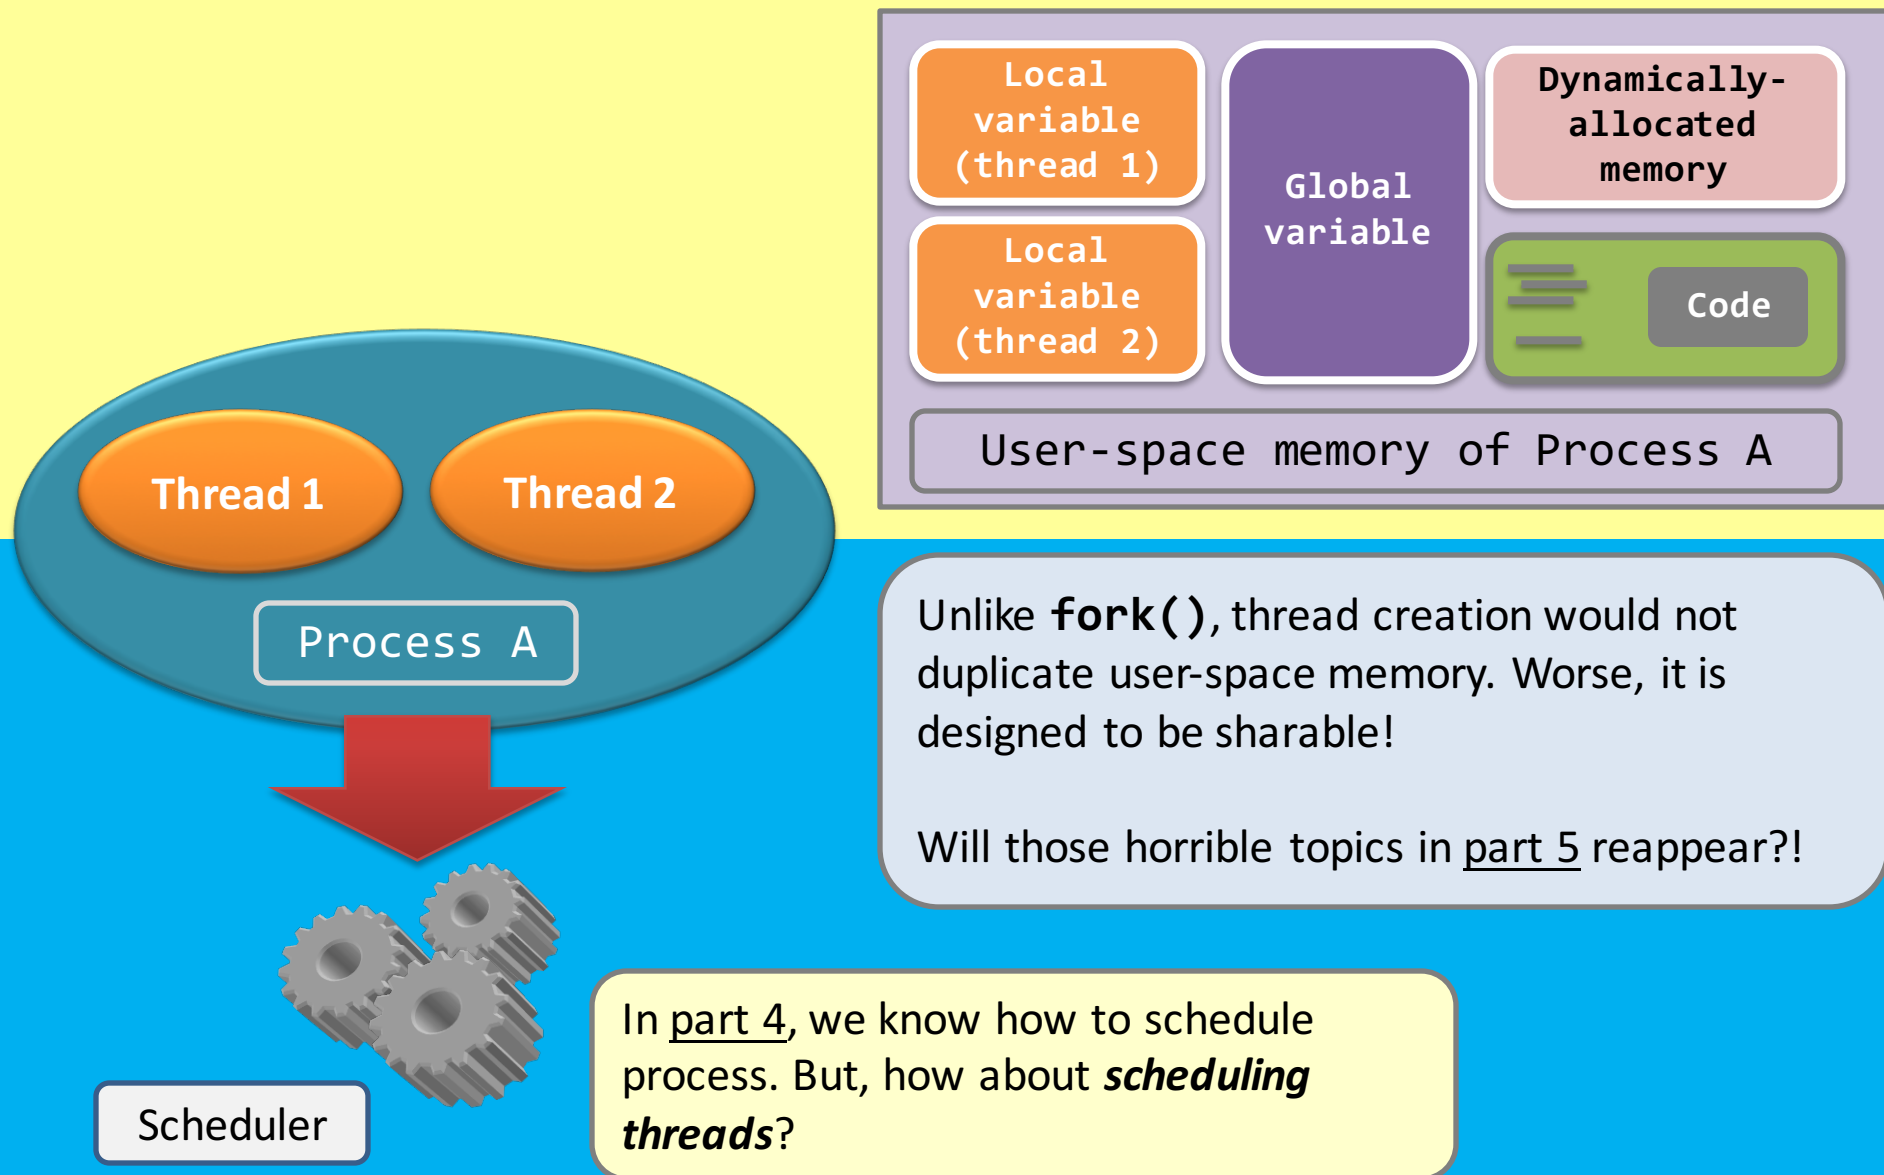
3150 - Operating Systems

Dr. WONG Tsz Yeung

Chapter 2, part 6 Light-Weight Processes: Threads

*-threads & processes: like spaghetti & meatballs mixed together...
they are delicious but complicated...*

Outline



Multi-threading

- Introduction.



Multi-thread – an example

- Threading is everywhere:
 - Software with GUI;
 - Network programs;
- Basically, we need threads for:
 - Executing parallel tasks for easy data sharing;
 - Distributing blocking calls to different threads.



Thread 1: Networking.

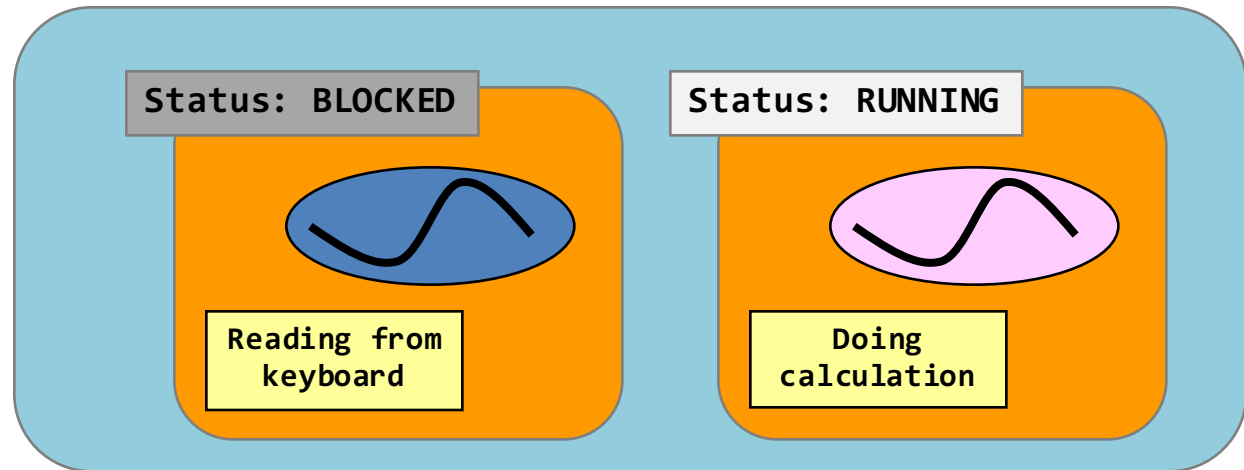
Thread 2: GUI control.

Thread 3: Flash player.

Multi-thread – merits

- Responsiveness and multi-tasking.
 - Multi-threading design allows an application to do **parallel tasks simultaneously**.
 - See from the following example?
 - Although a thread is blocked, the process can still depend on another thread to do other things!

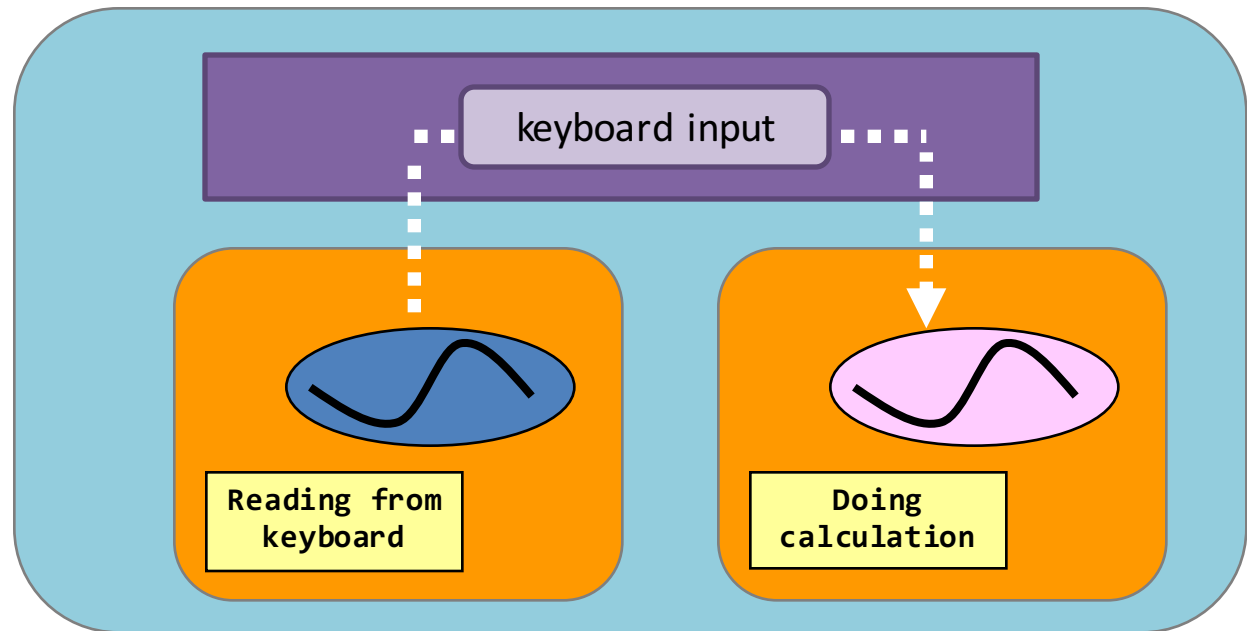
It'd be nice to assign **one thread for one blocking system/library call**.



Multi-thread – merits

- Ease in data sharing.
 - Data sharing can be done using:
 - global variables, and
 - dynamically allocated memory.

Of course, this leads to the **mutual exclusion** & the **synchronization** problems.

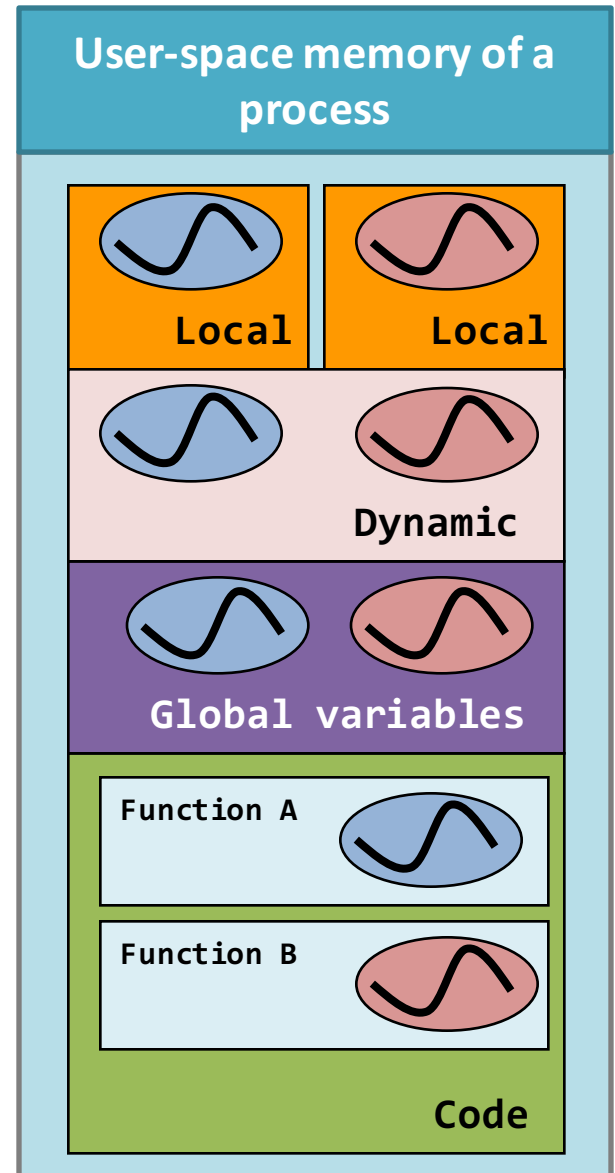


```
[examples@3150] cat pthread_game.c
```

Multi-thread – internals

Code

- All threads shares the same code.
- A thread starts with **one specific function**. and we name it the **thread function**. E.g., Functions A & B in the diagram.
- Of course, the thread function can invoke other functions or system calls.
- But, a thread could never return to the caller of the thread function.

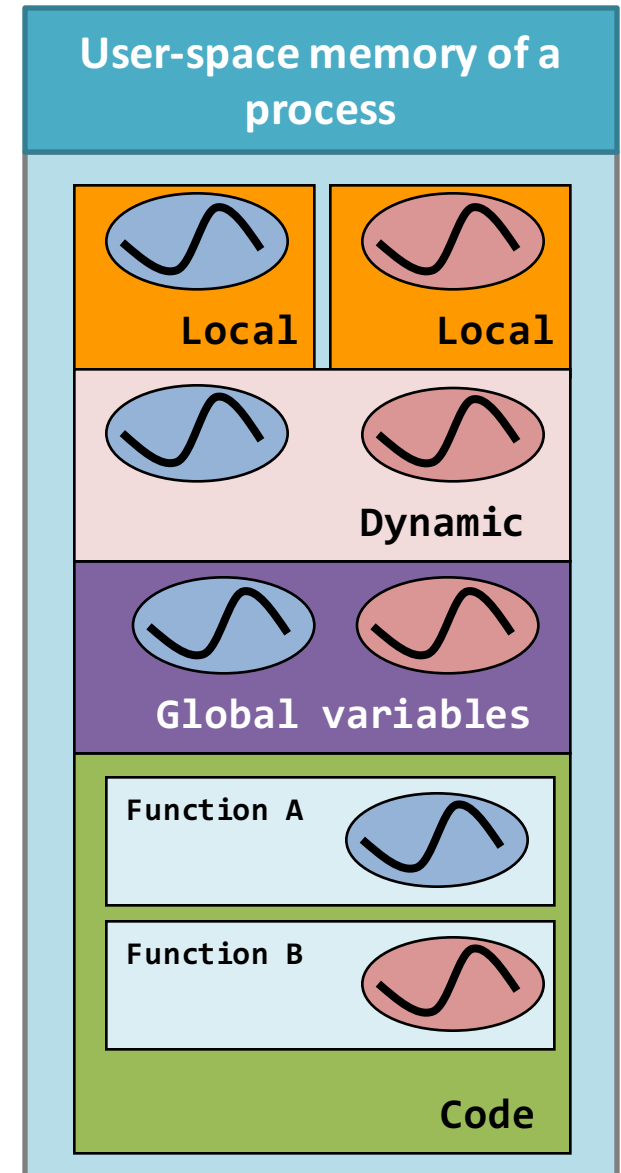


Multi-thread – internals

Dynamically allocated memory

Global variables

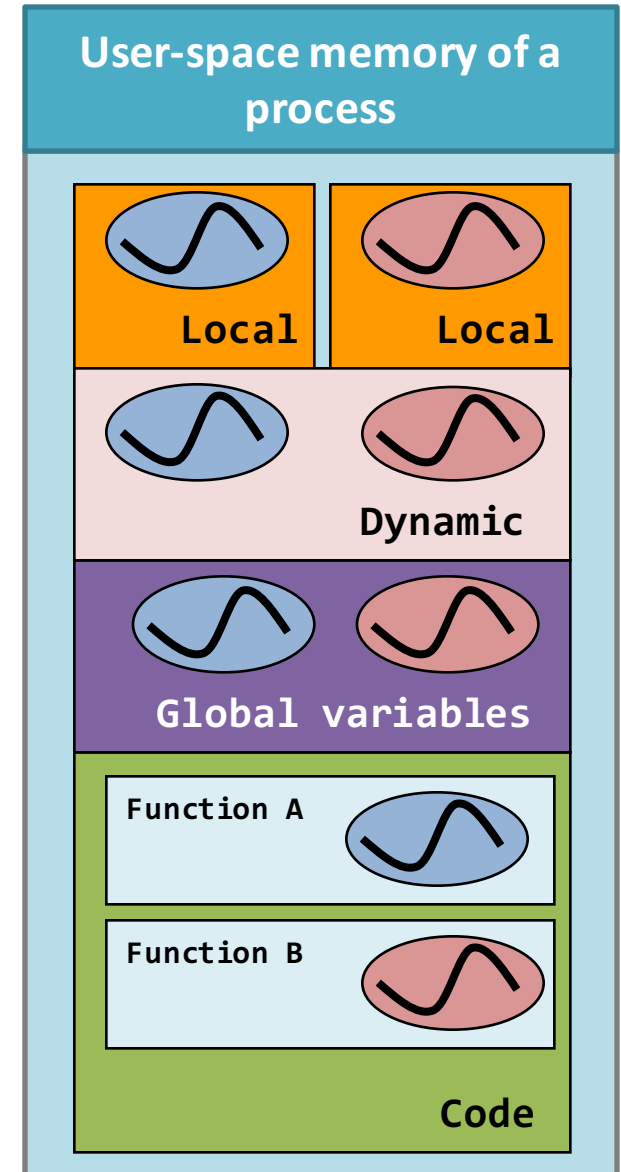
- All threads shares the same **global variable zone** and the same **dynamically allocated memory**.
- All threads can read from and write to both areas.
- So, will a multi-threaded process have the **race condition** problem built into the design?



Multi-thread – internals


Local variables

- Each thread has its own memory range for the local variables.
- So, the stack is the private zone for each stack.



Multi-threading

- Introduction.
- **Basic Programming.**



The pthread library

- We use a library called **pthread**, POSIX thread.

	Process	Thread
Creation	<code>fork()</code>	<code>pthread_create()</code>
I.D. Type	PID, an integer	“pthread_t”, a structure
Who am I?	<code>getpid()</code>	<code>pthread_self()</code>
Termination	<code>exit()</code>	<code>pthread_exit()</code>
Wait for child termination	<code>wait()</code> or <code>waitpid()</code>	<code>pthread_join()</code>
Kill?	<code>kill()</code>	<code>pthread_kill()</code>

Thread creation – `pthread_create()`

Thread Function

```
1 void * hello( void *input ) {  
2     printf(“%s\n”, (char *) input);  
3     pthread_exit(NULL);  
4 }
```

The `pthread_create()` function allows one argument to be passed to the thread function.

Main Function

```
5 int main(void) {  
6     pthread_t tid;  
7     pthread_create(&tid, NULL, hello, “hello world”);  
8     pthread_join(tid, NULL);  
9     return 0;  
10 }
```

This sets the thread function of the to-be-created thread as: `hello()`.

[examples@3150] cat pthread_hello_world.c

Thread creation – pthread_create()


Thread Function

```
1 void * hello( void *input ) {  
2     printf(“%s\n”, (char *) input);  
3     pthread_exit(NULL);  
4 }
```

Main Function

```
5 int main(void) {  
6     pthread_t tid;  
7     pthread_create(&tid, NULL, hello, “hello world”);  
8     pthread_join(tid, NULL);  
9     return 0;  
10 }
```

At the beginning,
there is only one
thread running: **the
main thread.**



Main Thread

[examples@3150] cat pthread_hello_world.c

Thread creation – pthread_create()

Thread Function


```
1 void * hello( void *input ) {  
2     printf(“%s\n”, (char *) input);  
3     pthread_exit(NULL);  
4 }
```

Main Function

```
5 int main(void) {  
6     pthread_t tid;  
7     pthread_create(&tid, NULL, hello, “hello world”);  
8     pthread_join(tid, NULL);  
9     return 0;  
10 }
```


The hello thread is created!

It is running *“together”* with the main thread.



Main Thread

pthread_create()



Hello Thread

[examples@3150] cat pthread_hello_world.c

Thread creation – `pthread_create()`

Thread Function

```
1 void * hello( void *input ) {  
2     printf(“%s\n”, (char *) input);  
3     pthread_exit(NULL);  
4 }
```

Main Function

```
5 int main(void) {  
6     pthread_t tid;  
7     pthread_create(&tid, NULL, hello, “hello world”);  
8     pthread_join(tid, NULL);  
9     return 0;  
10 }
```


Remember `wait()`
and `waitpid()`?

`pthread_join()`
performs similarly.



Blocked

Main Thread



Hello Thread

[examples@3150] cat pthread_hello_world.c

Thread creation – pthread_create()


Thread Function

```
1 void * hello( void *input ) {  
2     printf(“%s\n”, (char *) input);  
3     pthread_exit(NULL);  
4 }
```

Main Function

```
5 int main(void) {  
6     pthread_t tid;  
7     pthread_create(&tid, NULL, hello, “hello world”);  
8     pthread_join(tid, NULL);  
9     return 0;  
10 }
```

Termination of the target thread causes `pthread_join()` to return.



[examples@3150] cat pthread_hello_world.c

Thread creation – passing parameter?

Thread Function

```
1 void * do_your_job( void *input ) {  
2     printf("child = %d\n", *( (int *) input) );  
3     *((int *) input) = 20;  
4     printf("child = %d\n", *( (int *) input) );  
5     pthread_exit(NULL);  
6 }
```

Main Function

```
7 int main(void) {  
8     pthread_t tid;  
9     int input = 10;  
10    printf("main = %d\n", input);  
11    pthread_create(&tid, NULL, do_your_job, &input);  
12    pthread_join(tid, NULL);  
13    printf("main = %d\n", input);  
14    return 0;  
15 }
```

```
$ ./pthread_evil_1  
main = 10  
child = 10  
child = 20  
main = 20  
$
```



What?! Each thread should have a separated stack, right?

Why do we have such results!!!

[examples@3150] cat pthread_evil_1.c

Thread creation – passing parameter?

Well, we all know that the local variable “`input`” is in the stack for the main thread.

```
1 void * do_your_job( void *input ) {
2     printf("child = %d\n", *( (int *) input) );
3     *((int *) input) = 20;
4     printf("child = %d\n", *( (int *) input) );
5     pthread_exit(NULL);
6 }

7 int main(void) {
8     pthread_t tid;
9     int input = 10;
10    printf("main = %d\n", input);
11    pthread_create(&tid, NULL, do_your_job, &input);
12    pthread_join(tid, NULL);
13    printf("main = %d\n", input);
13    return 0;
14 }
```

Local
(main thread)

Dynamic

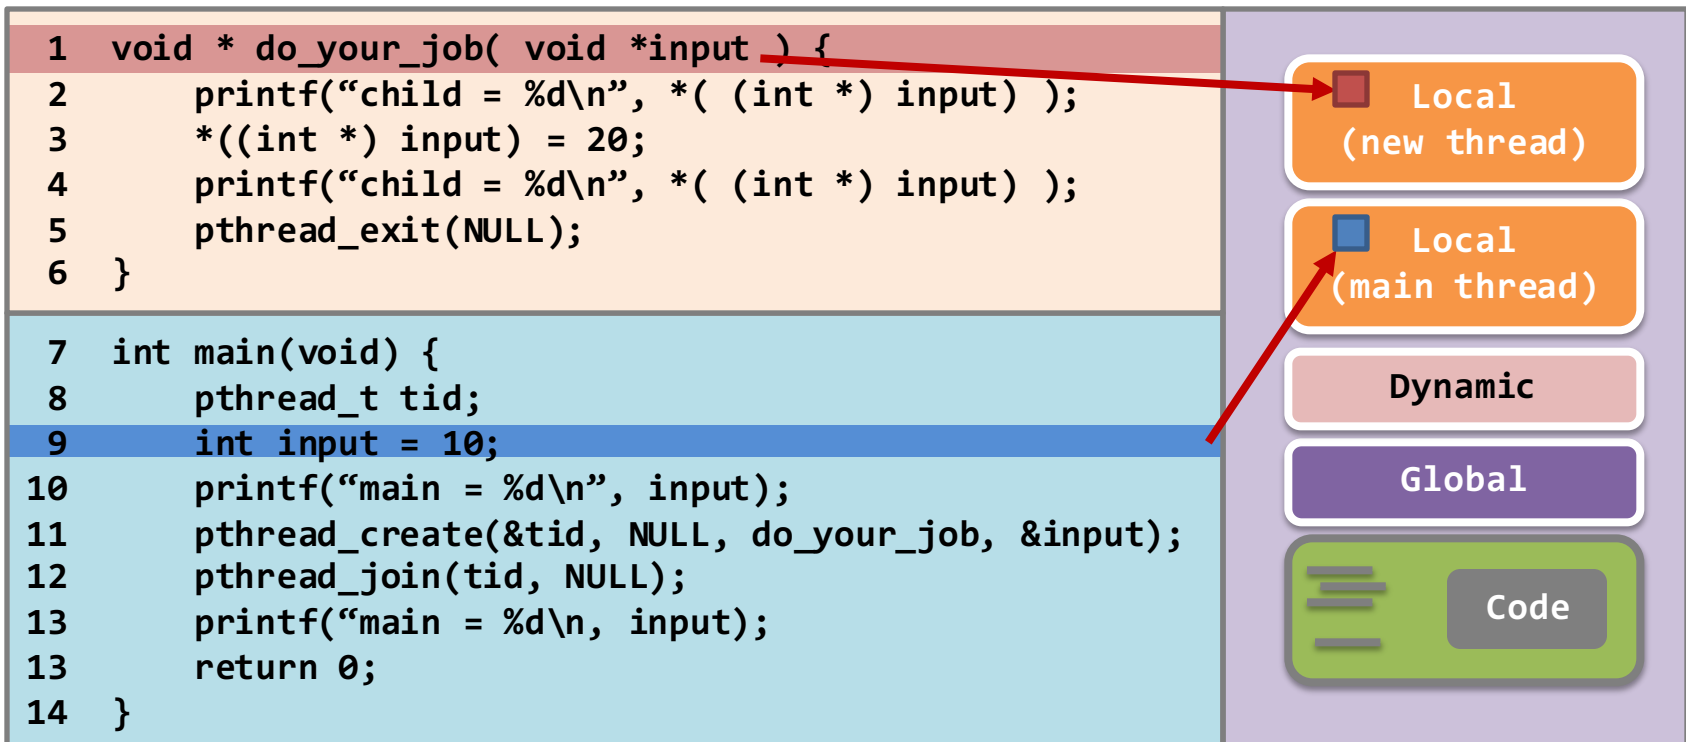
Global

Code

[examples@3150] cat pthread_evil_1.c

Thread creation – passing parameter?

Yet...the stack for the new thread is not on the another process, but is on the same piece of user-space memory as the main thread.




[examples@3150] cat pthread_evil_1.c

Thread creation – passing parameter?

The `pthread_create()` function only passes an address to the new thread.
Worse, the address is pointing to a variable in the stack of the main thread!

```
1 void * do_your_job( void *input ) {  
2     printf("child = %d\n", *( (int *) input) );  
3     *((int *) input) = 20;  
4     printf("child = %d\n", *( (int *) input) );  
5     pthread_exit(NULL);  
6 }  
  
7 int main(void) {  
8     pthread_t tid;  
9     int input = 10;  
10    printf("main = %d\n", input);  
11    pthread_create(&tid, NULL, do_your_job, &input);  
12    pthread_join(tid, NULL);  
13    printf("main = %d\n", input);  
13    return 0;  
14 }
```




[examples@3150] cat pthread_evil_1.c

Thread creation – passing parameter?

Therefore, the new thread can change the value in the main thread, and vice versa.
Is it brain-damaging?

```
1 void * do_your_job( void *input ) {  
2     printf("child = %d\n", *( (int *) input) );  
3     *((int *) input) = 20;  
4     printf("child = %d\n", *( (int *) input) );  
5     pthread_exit(NULL);  
6 }  
  
7 int main(void) {  
8     pthread_t tid;  
9     int input = 10;  
10    printf("main = %d\n", input);  
11    pthread_create(&tid, NULL, do_your_job, &input);  
12    pthread_join(tid, NULL);  
13    printf("main = %d\n", input);  
14    return 0;  
15 }
```



[examples@3150] cat pthread_evil_1.c

Thread creation – passing parameter?

Thread Function

```
1 void * do_your_job(void *input) {  
2   int id = *((int *) input);  
3   printf("My ID number = %d\n", id);  
4   pthread_exit(NULL);  
5 }
```

Doing backup!

Just want to avoid
accessing the stack
of the main thread.

Main Function

```
6 int main(void) {  
7   int i;  
8   pthread_t tid[5];  
9  
10  for(i = 0; i < 5; i++)  
11    pthread_create(&tid[i], NULL, do_your_job, &i);  
12  for(i = 0; i < 5; i++)  
13    pthread_join(tid[i], NULL);  
14  return 0;  
15 }
```

Challenge!

Is there any problem?
If yes, how to correct it?

[examples@3150] cat pthread_evil_2.c pthread_evil_2sol.c

Thread termination – passing return value?

Thread Function

```
1 void * do_your_job(void *input) {  
2     int *output = (int *) malloc(sizeof(int));  
3     srand(time(NULL));  
4     *output = ((rand() % 10) + 1) * (((int *) input));  
5     pthread_exit(output);  
6 }
```

Important! Similar to the normal function call, you cannot return a pointer to a local variable.

`void pthread_exit(void *return_value);`

Together with termination, a pointer to a global variable or a piece of dynamically allocated memory is returned to the main thread.

Main Function

```
7 int main(void) {  
8     pthread_t tid;  
9     int input = 10, *output;  
10    pthread_create(&tid, NULL, do_your_job, &input);  
11    pthread_join(tid, (void **) &output);  
12    printf("output = %d\n", *output);  
13    return 0;  
14 }
```

Using pass-by-reference, a pointer to the result is received in the main thread.

[examples@3150] cat pthread_exit.c

Multi-threading


- Introduction.
- Basic Programming.
- **Thread scheduling?**



Thread models – 1 of 2

- **Many-to-One Model**


- All the threads are mapped to one process structure in the kernel.
- Merit.
 - easy for the kernel to implement.
- Drawback.
 - when a blocking system call is called, all the threads will be blocked.
- **Example.** Old UNIX & [green thread](#) in some programming languages.



Thread models – 2 of 2

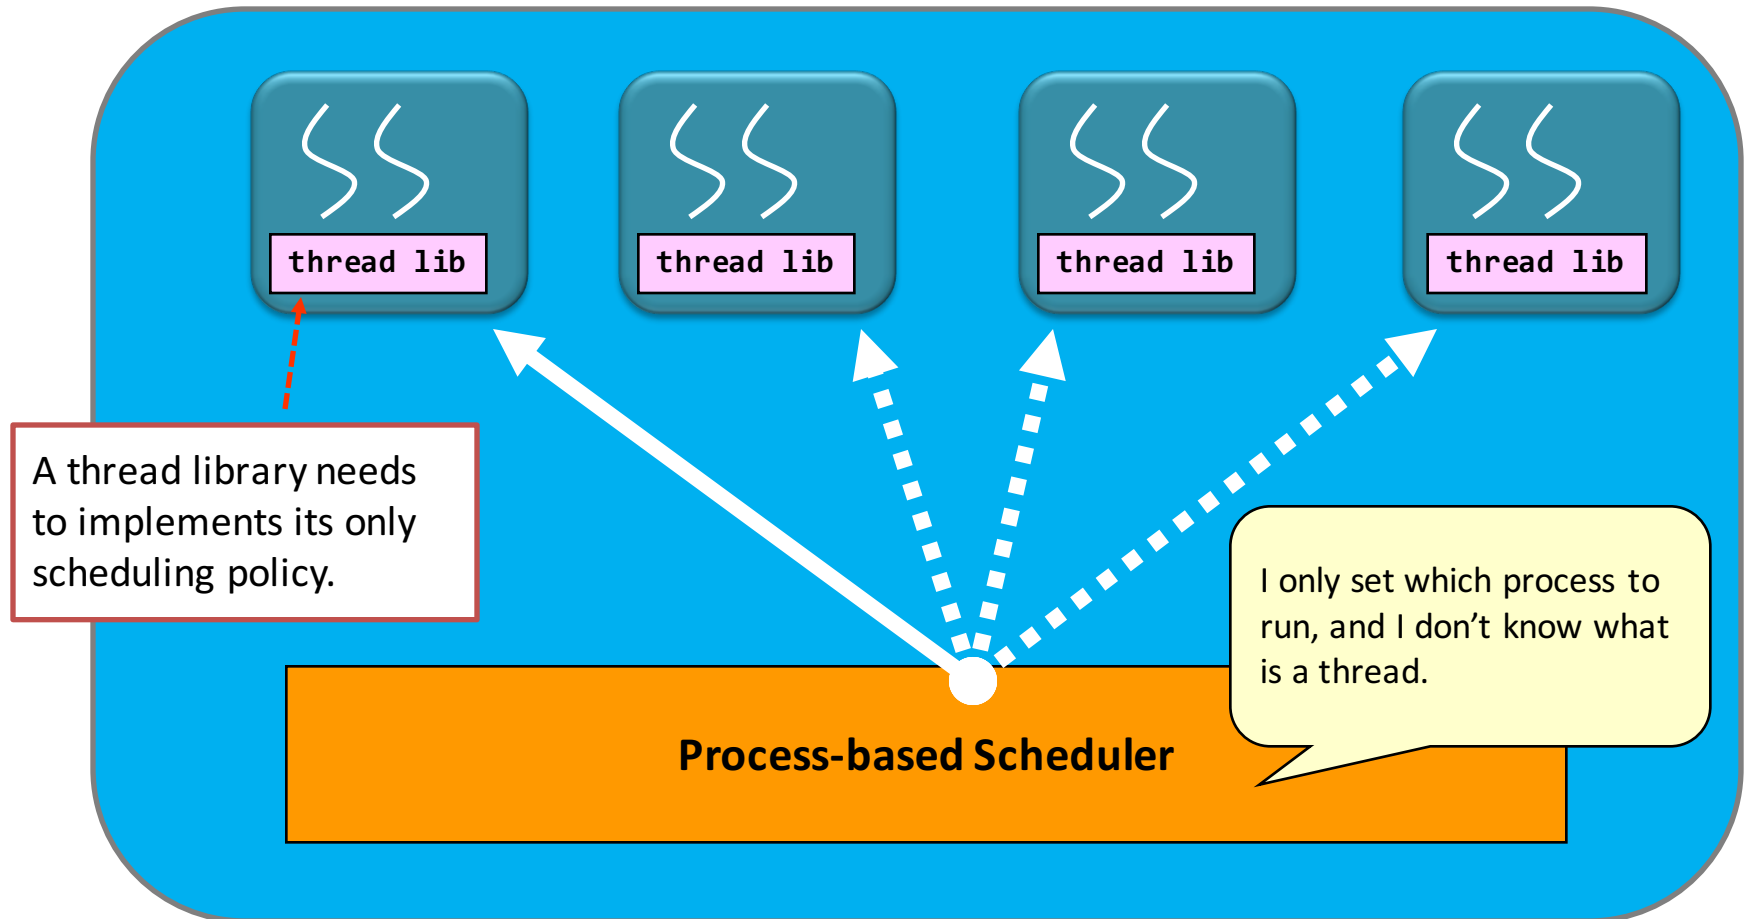
- **One-to-One Model**

- Each thread is mapped to a process or a thread structure.
- Merit:
 - calling blocking system calls only block those calling threads.
 - A high degree of concurrency.
- Drawback:
 - cannot create too many threads as it is restricted by the size of the kernel memory.
- **Example.** Linux and Windows follow this thread model.



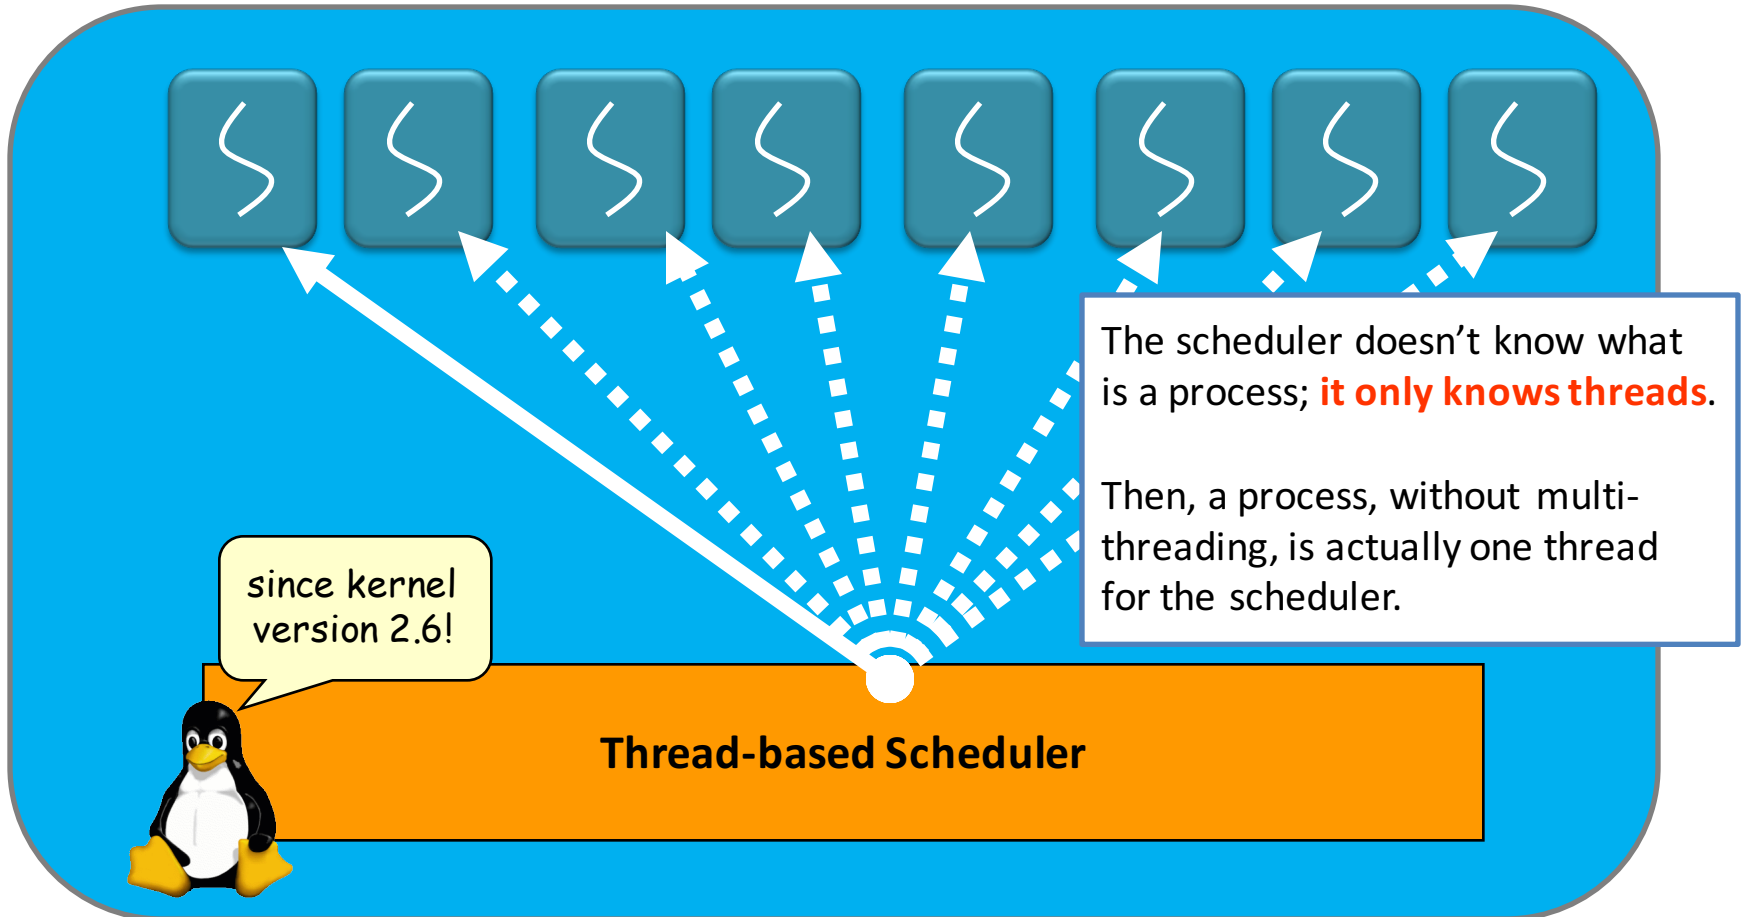
Scheduling – why & who cares?

- If a scheduler only interests in **processes**...




Scheduling – why & who cares?

- If a scheduler only interests in **threads**...



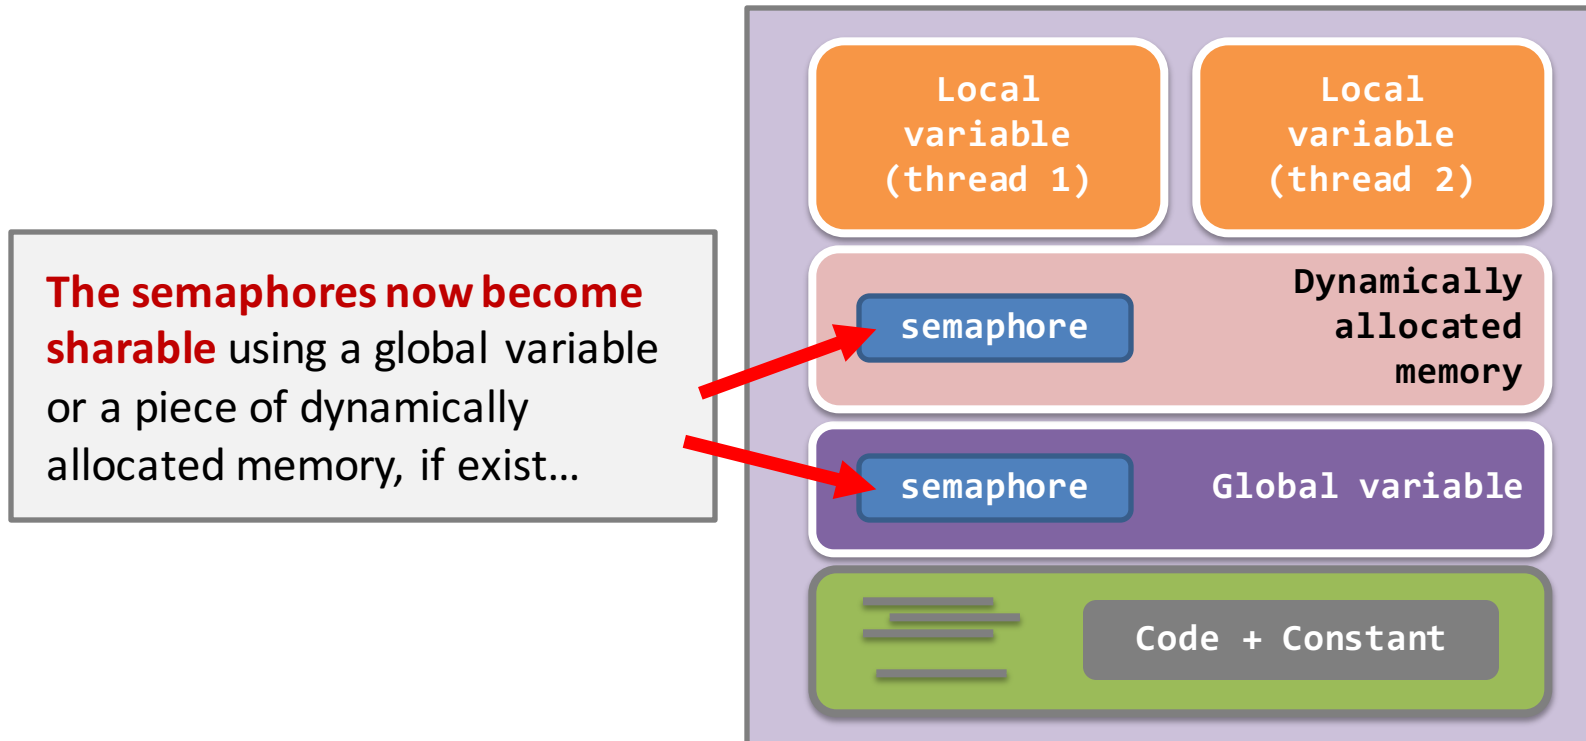
Multi-threading

- Introduction.
- Basic Programming.
- Thread scheduling?
- **Mutual exclusion
& synchronization!**



Mutual exclusion...using semaphore?

- We meet our old friend...
 - Dynamically allocated memory & global variables are shared objects.
 - But, there is **no semaphore** in the pthread library!



Mutual exclusion...using mutex!

Data type

`pthread_mutex_t`

This is a binary semaphore.

Initialization #1

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

This is the initialization at the declaration statement only.

Initialization #2

```
pthread_mutex_init(&mutex, NULL);
```

This is the initialization for any locations in a program.

Locking

```
pthread_mutex_lock(&mutex);
```

- If the mutex is not locked, the call locks the mutex.
- If the mutex is locked, the call blocks the calling thread.

Unlocking

```
pthread_mutex_unlock(&mutex);
```

- If the mutex is locked, the call unlocks the mutex. If there are threads blocked because of this mutex, those threads will resume.
- If the mutex is unlocked, the call does nothing.

Mutual exclusion...example

Data type

`pthread_mutex_t`


This is a binary semaphore.

Thread Function

```
1 while(TRUE) {  
2     pthread_mutex_lock(&mutex);  
3     if(shared < MAXIMUM)  
4         shared++;  
5     else {  
6         pthread_mutex_unlock(&mutex);  
7         break;  
8     }  
9     pthread_mutex_unlock(&mutex);  
10    sleep(rand() % 2);  
11 }
```

Shared object.

Purpose of the program?



[examples@3150] cat pthread_mutex.c

Synchronization...without semaphore?

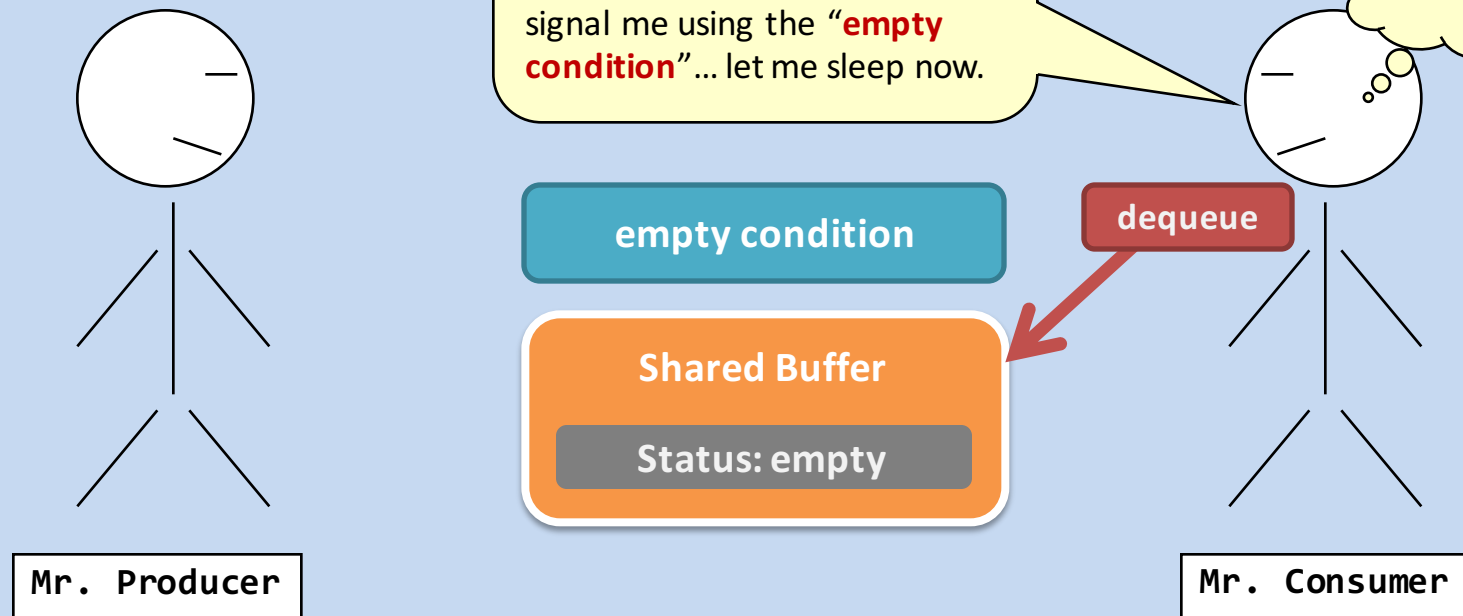
Data type

`pthread_cond_t`

This is a condition variable. Its main purpose is to **let a thread to wait** under certain conditions.

Case 1.

An illustration using
producer-consumer model



Synchronization...without semaphore?

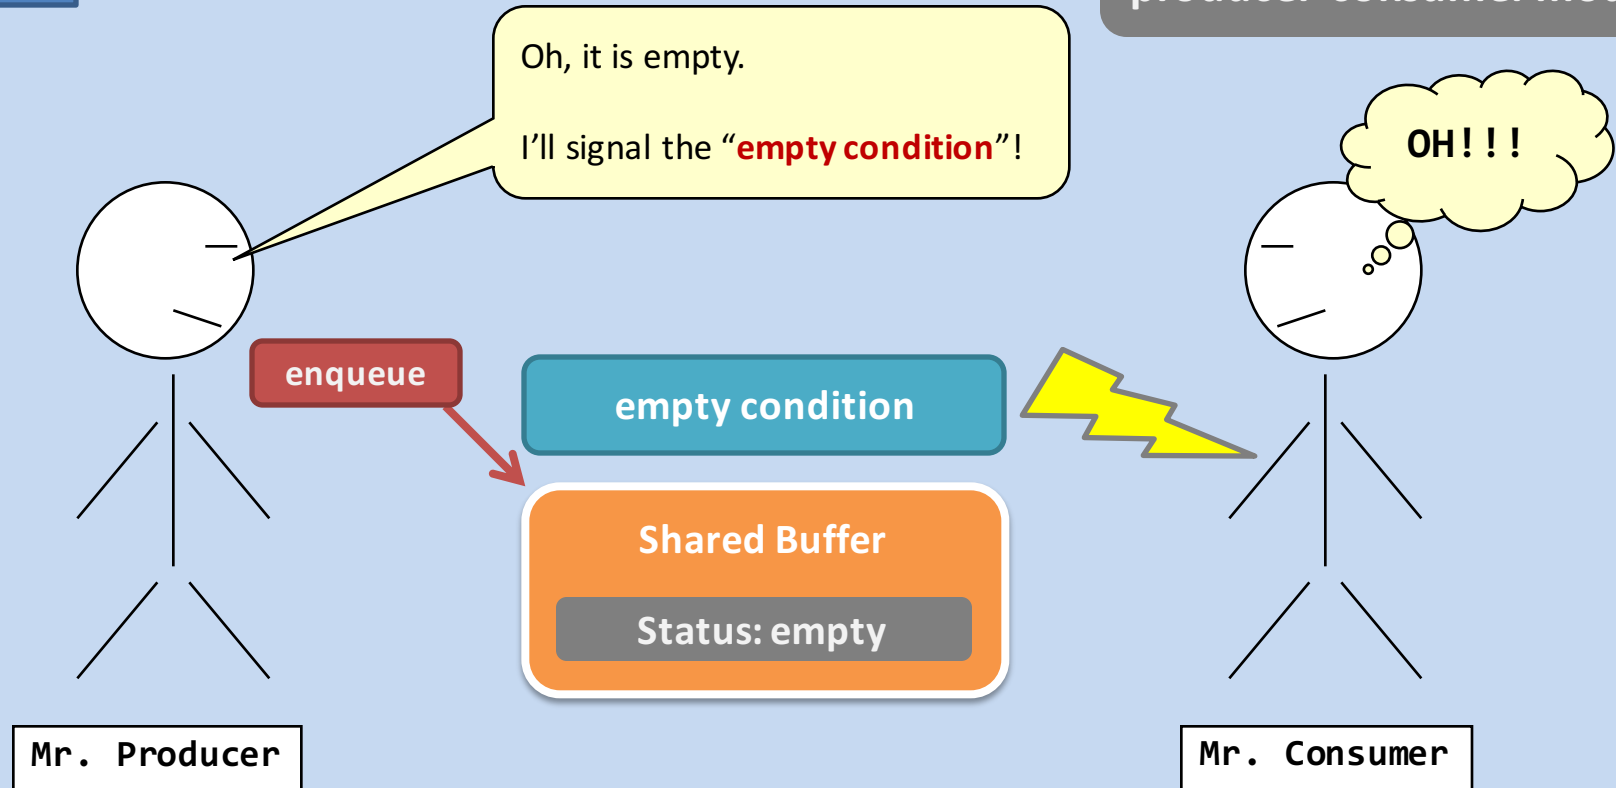
Data type

`pthread_cond_t`

This is a condition variable. Its main purpose is to **let a thread to wait** under certain conditions.

Case 1.

An illustration using producer-consumer model



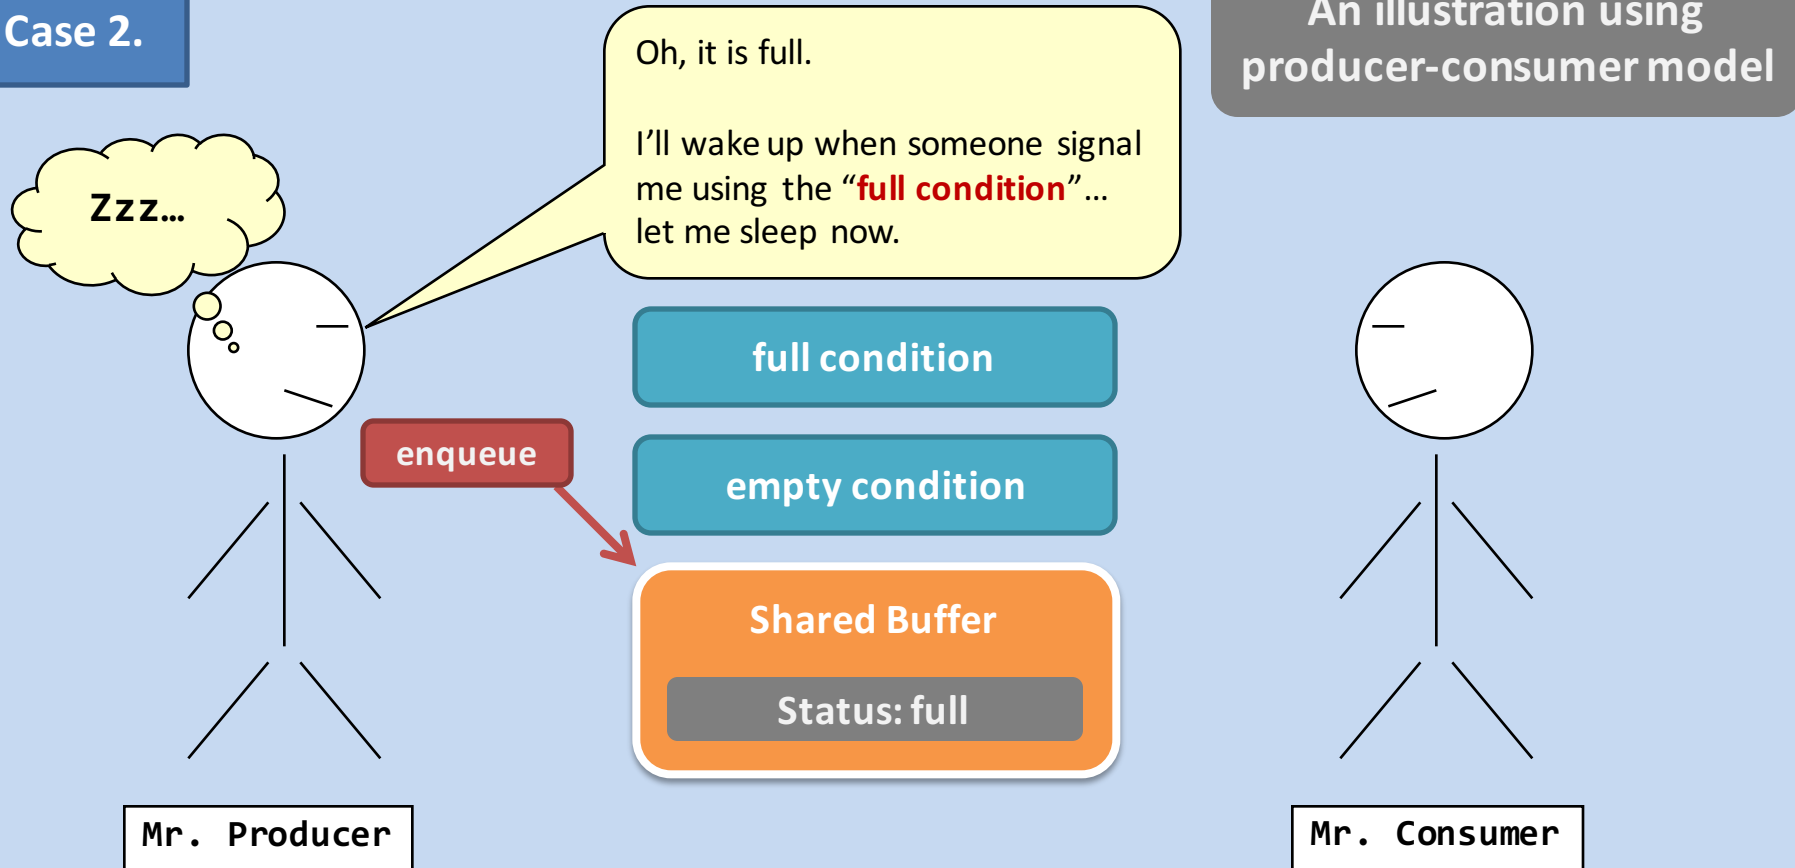
Synchronization...without semaphore?

Data type

`pthread_cond_t`

This is a condition variable. Its main purpose is to **let a thread to wait** under certain conditions.

Case 2.



Synchronization...without semaphore?

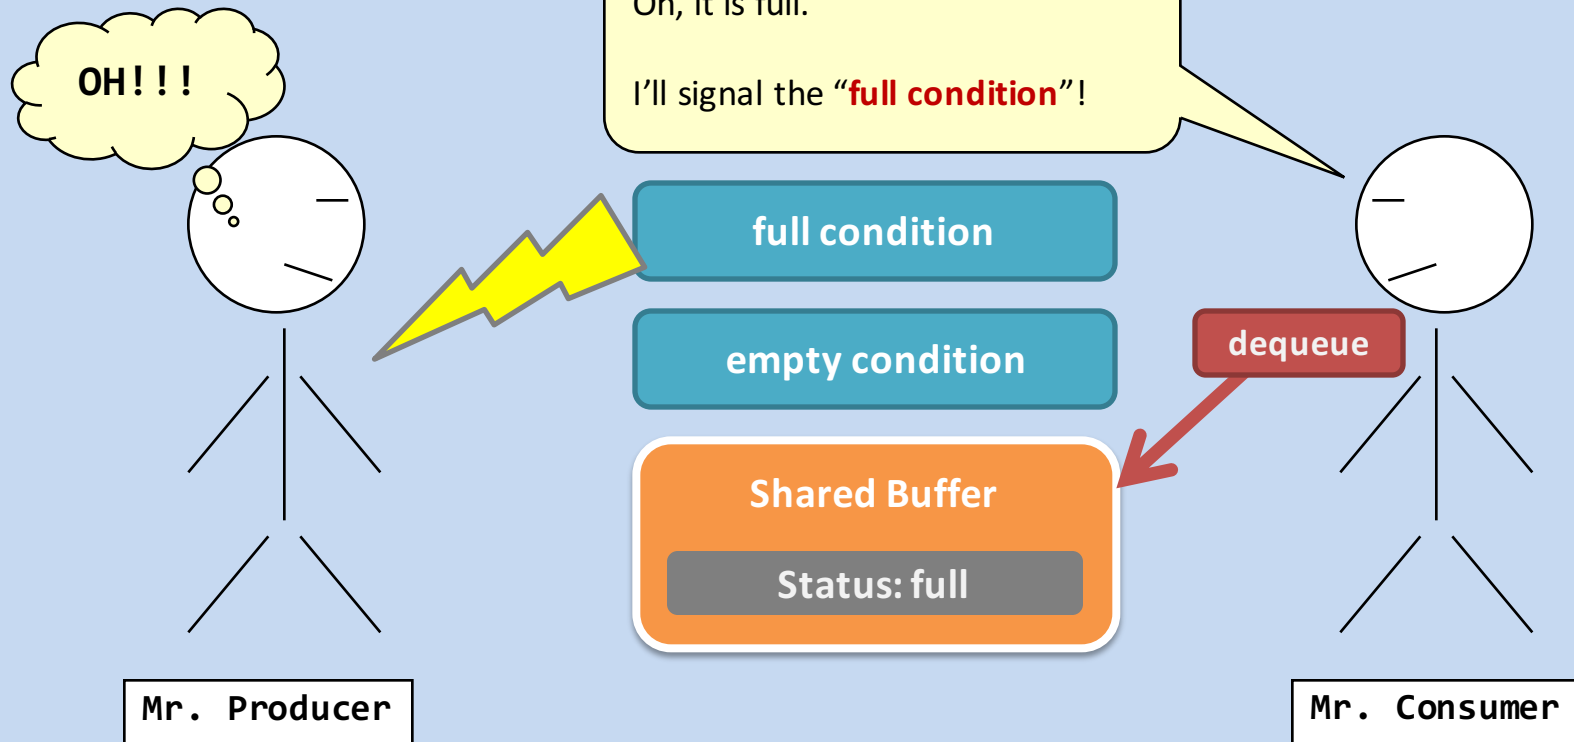
Data type

`pthread_cond_t`

This is a condition variable. Its main purpose is to **let a thread to wait** under certain conditions.

Case 2.

An illustration using producer-consumer model



Synchronization...using condition...

Data type

`pthread_cond_t`

This is a condition variable. Its main purpose is to **let a thread to wait** under certain conditions.

Initialization #1	<code>pthread_cond_t condition = PTHREAD_COND_INITIALIZER;</code>
Initialization #2	<code>pthread_cond_init(&condition, NULL);</code>
Waiting #1	<code>pthread_cond_wait(&condition, &mutex);</code>

This function is to set the thread to wait on a condition variable. **But, why is there a “mutex” variable?**

It is there to break the “hold-and-wait” deadlocking condition. Before a thread reaches the waiting statement, **it is usually inside a critical section.** Therefore...[think about it by yourself].

When the thread is about to be blocked, the state of the “**mutex**” variable will be **unlocked “atomically” and automatically.**

When the thread is unblocked, the “**mutex**” variable is **locked “atomically” and automatically.**

Synchronization...using condition...

Data type

`pthread_cond_t`

This is a condition variable. Its main purpose is to **let a thread to wait** under certain conditions.

Initialization #1	<code>pthread_cond_t condition = PTHREAD_COND_INITIALIZER;</code>
Initialization #2	<code>pthread_cond_init(&condition, NULL);</code>
Waiting #1	<code>pthread_cond_wait(&condition, &mutex);</code>
Waiting #2	<code>pthread_cond_timedwait(&condition, &mutex, &timeout);</code>

This is a fantastic function. At least, you can find it useful in getting out of an indefinite sleep.

The timeout variable is of the type “**struct timespec**”. It should store the real time that the timeout reaches.

struct timespec		
time_t	tv_sec	Seconds
long	tv_nsec	Nano-second, 1×10^{-9}

Synchronization...using condition...

Data type

`pthread_cond_t`

This is a condition variable. Its main purpose is to **let a thread to wait** under certain conditions.

Initialization #1	<code>pthread_cond_t condition = PTHREAD_COND_INITIALIZER;</code>
Initialization #2	<code>pthread_cond_init(&condition, NULL);</code>
Waiting #1	<code>pthread_cond_wait(&condition, &mutex);</code>
Waiting #2	<code>pthread_cond_timedwait(&condition, &mutex, &timeout);</code>
Wakeup #1	<code>pthread_cond_signal(&condition);</code> This function is to wake up a thread which is waiting on the “ condition ” variable. If there are more than one thread waiting, then at least one of those threads will wake up .
Wakeup #2	<code>pthread_cond_broadcast(&condition);</code> This function is to wake up all threads which is waiting on the “ condition ” variable.

Producer-consumer – in pthread

Global variables

```
1 pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
2 pthread_cond_t full = PTHREAD_COND_INITIALIZER;
3 pthread_cond_t empty = PTHREAD_COND_INITIALIZER;
4 /* the buffer structure */
```

Producer function

```
1 pthread_lock(&mutex);

2 if buffer is full ; then
3     pthread_cond_wait(&full, &mutex);

4 if buffer is empty, then
5     pthread_cond_signal(&empty);

6 produce & enqueue an item.

7 pthread_unlock(&mutex);
```

Consumer Function

```
1 pthread_lock(&mutex);

2 if buffer is empty ; then
3     pthread_cond_wait(&empty, &mutex);

4 if buffer is full, then
5     pthread_cond_signal(&full);

6 dequeue & consume an item.

7 pthread_unlock(&mutex);
```

[examples@3150] cat pthread_producer_consumer.c

How about other IPC problems?

- It seems that we have to **rethink** all the IPC solutions that we learnt before...
 - Is there any method to relieve such pain?

Semaphore data type

```
pthread_mutex_t mutex;  
pthread_cond_t  cond;  
int             value
```

Extra & Challenge.

Implement the solutions to those IPC problems using the pthread library!

void down(Semaphore *s)


```
1 pthread_mutex_lock(s->mutex);  
2 while(s->value == 0)  
3     pthread_cond_wait(&s->cond, &s->mutex);  
4 s->value--;  
5 pthread_mutex_unlock(&s->mutex);
```

void up(Semaphore *s)

```
1 pthread_mutex_lock(s->mutex);  
2 if(s->value == 0)  
3     pthread_cond_signal(&s->cond);  
4 s->value++;  
5 pthread_mutex_unlock(&s->mutex);
```

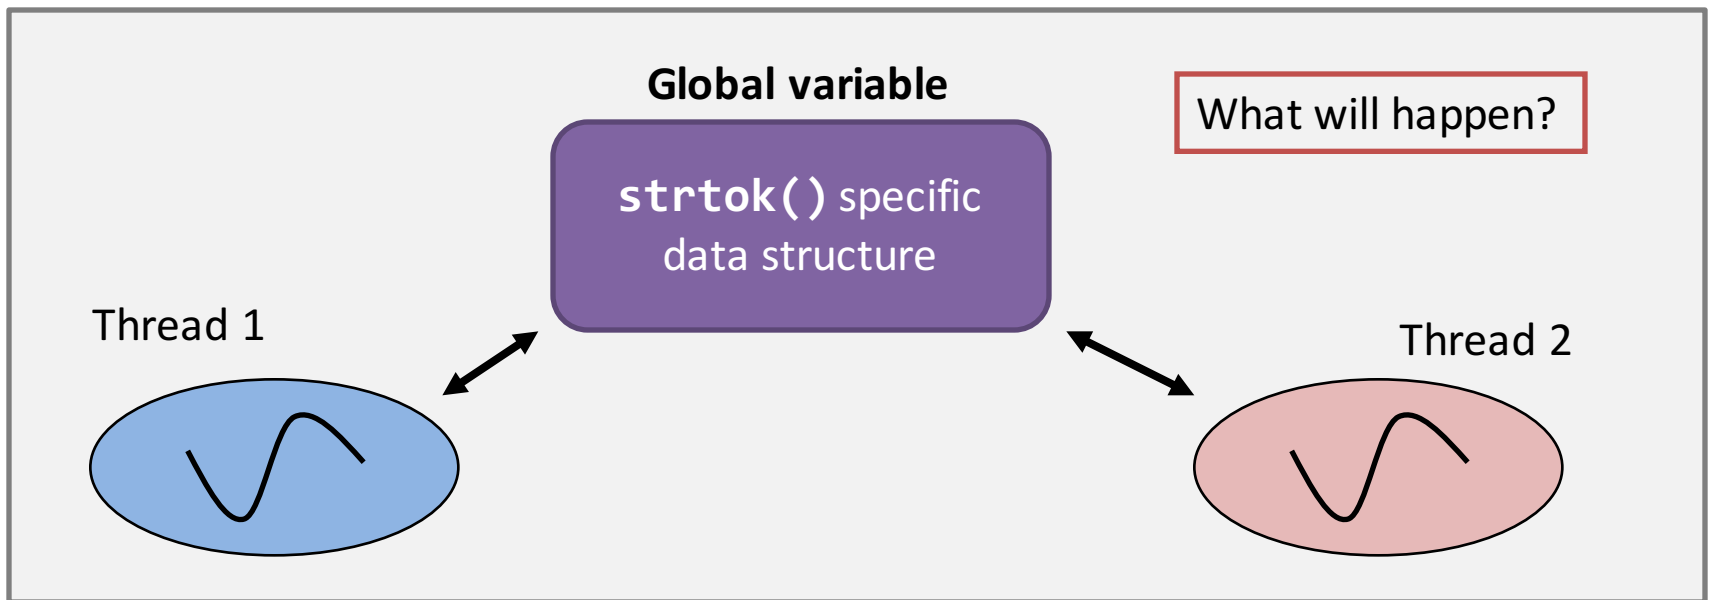
Multi-threading

- Introduction.
- Basic Programming.
- Thread scheduling?
- Mutual exclusion
& synchronization!
- **Thread safety?!**



Thread Safety – an important concern

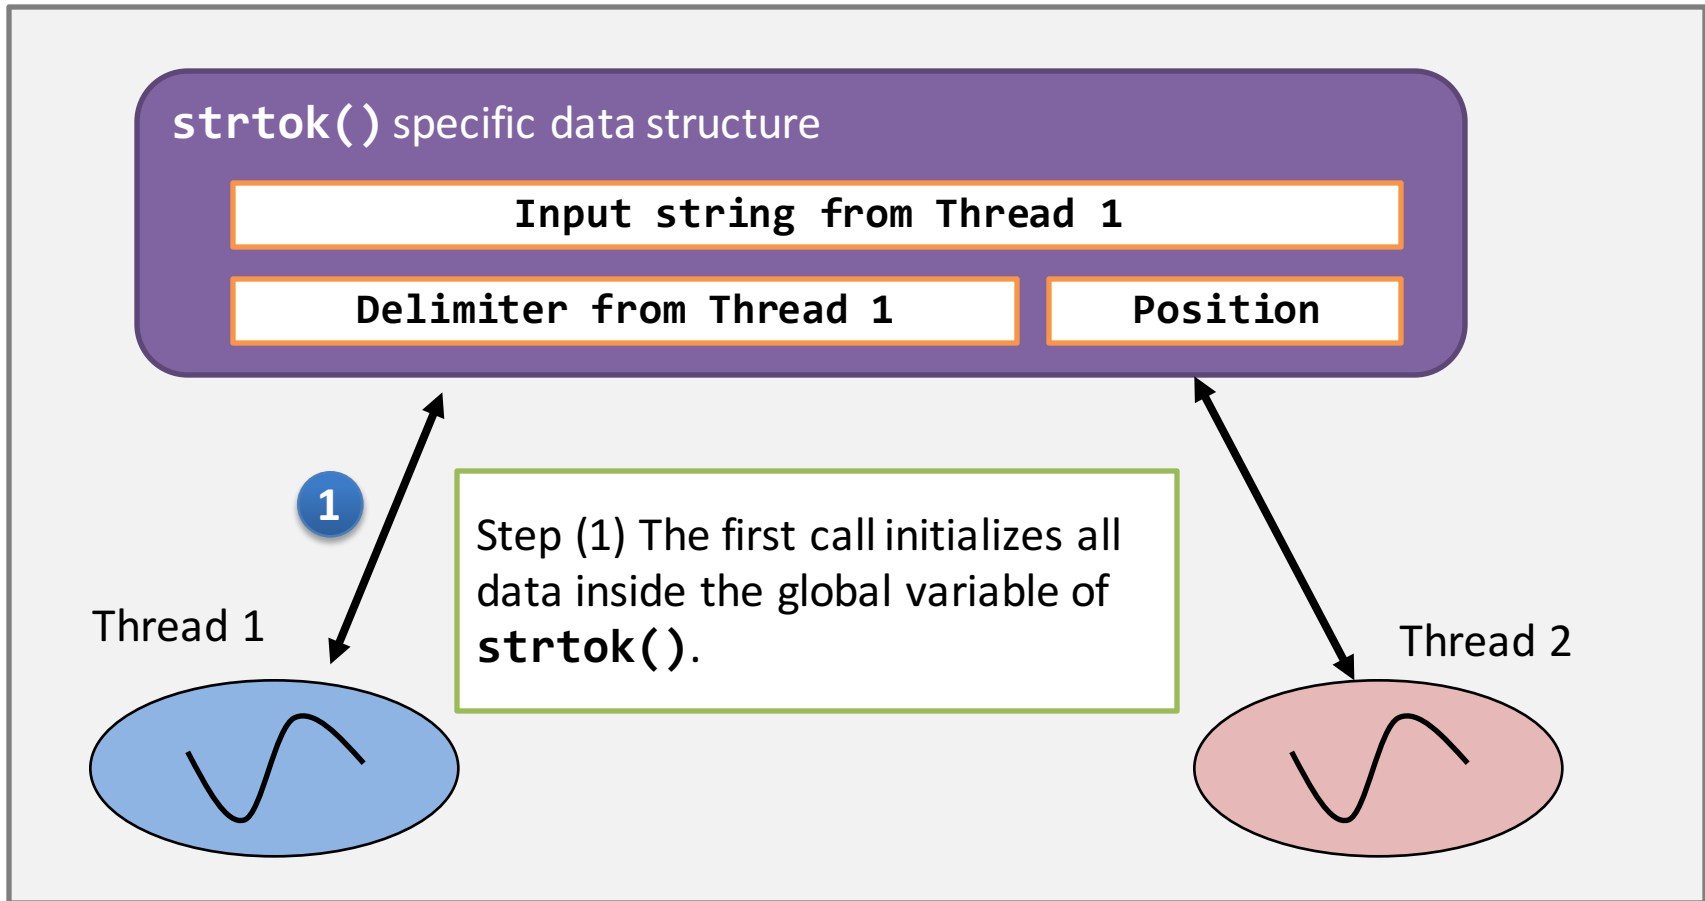
- Some functions, by design, has assumed a **single-threaded execution**. E.g., **strtok()**;
 - When such a function is deployed in a multi-threaded execution, we face a problem called **thread safety**.



[examples@3150] cat normal_strtok.c pthread_strtok.c

Thread Safety – an important concern

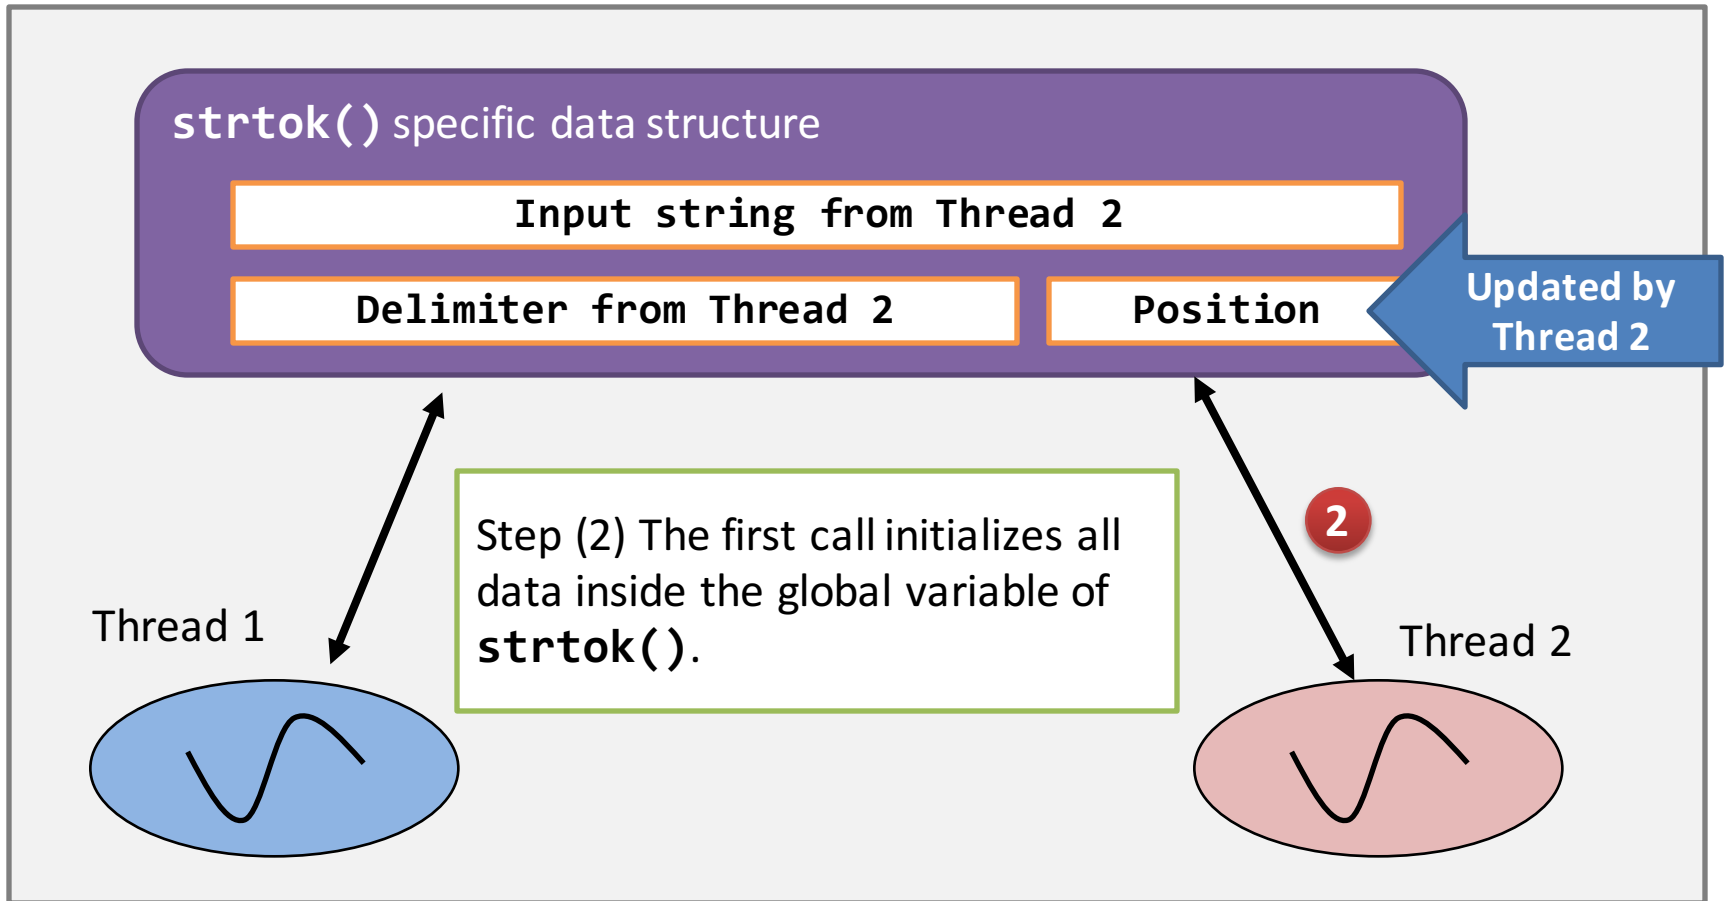
- Illustration of thread-unsafe function (1 of 3)



[examples@3150] cat normal_strtok.c pthread_strtok.c

Thread Safety – an important concern

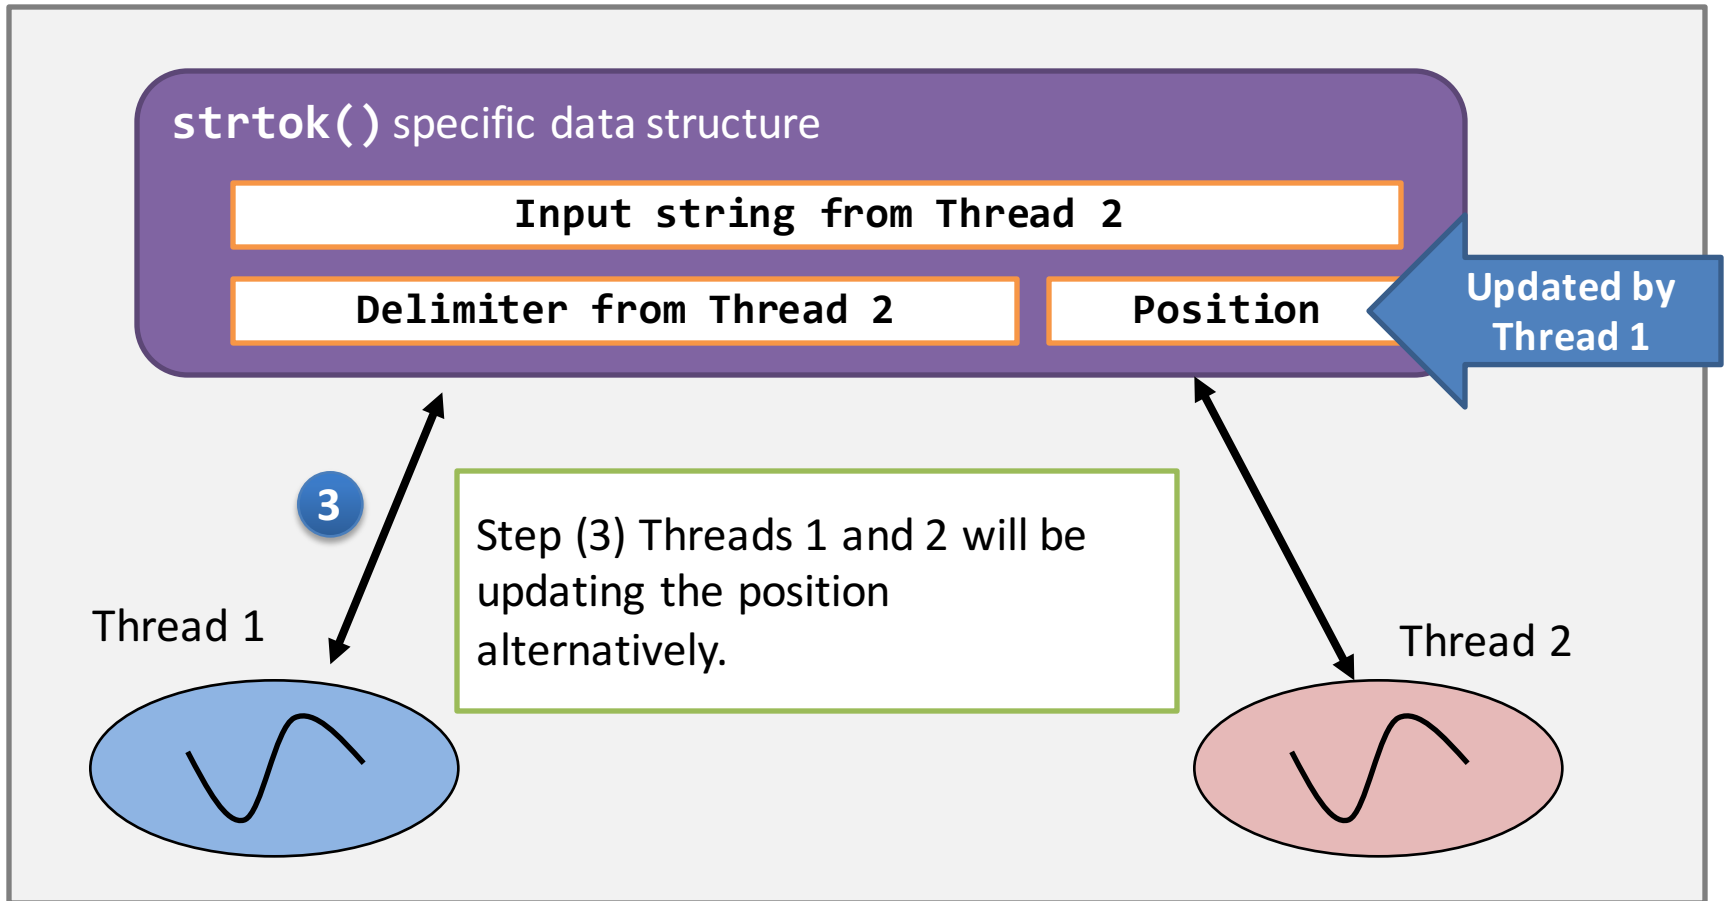
- Illustration of thread-unsafe function (2 of 3)



[examples@3150] cat normal_strtok.c pthread_strtok.c

Thread Safety – an important concern

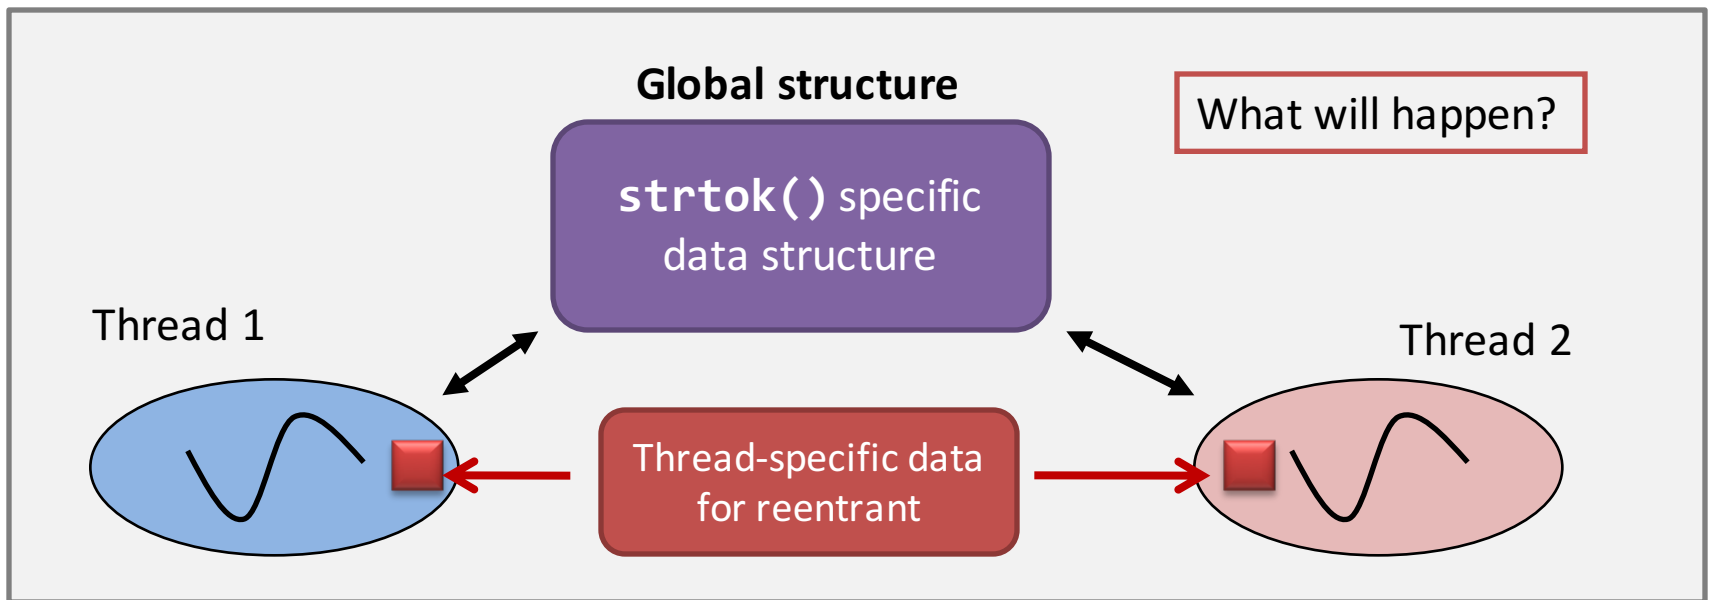
- Illustration of thread-unsafe function (3 of 3)



[examples@3150] cat normal_strtok.c pthread_strtok.c

Thread Safety – an important concern

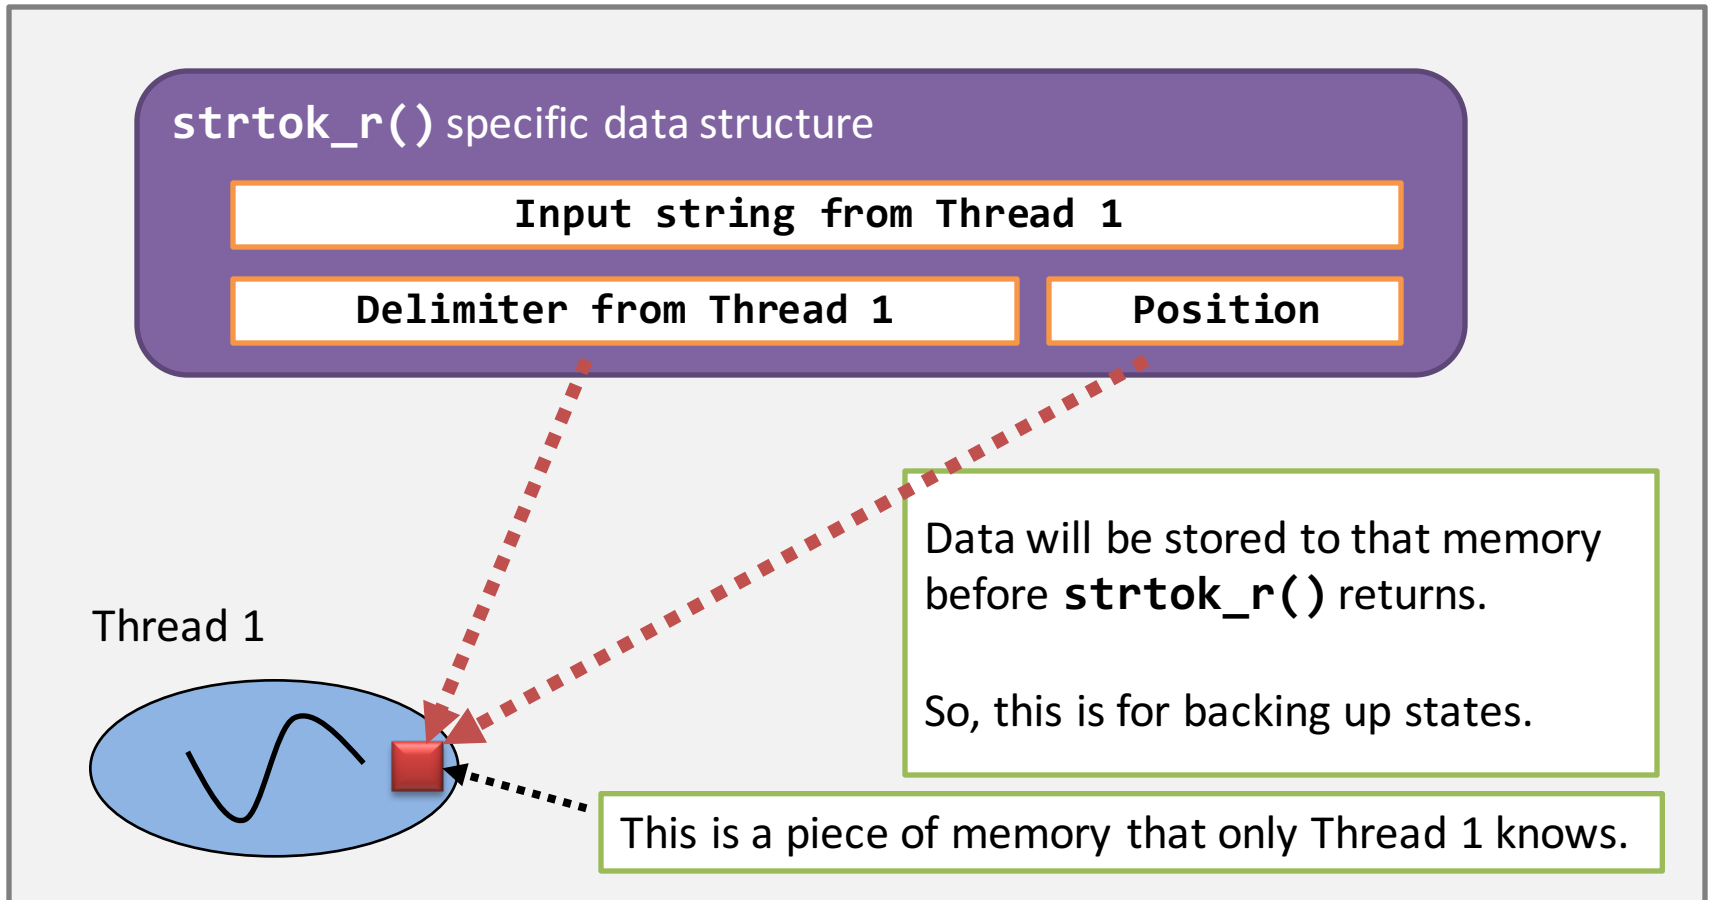
- A thread-safe function is a function that has **the reentrant property** implemented.
 - The **strtok()** thread-safe counterpart: **strtok_r()**.



[examples@3150] cat pthread_strtok_r.c

Thread Safety – an important concern

- Using re-entrant data.



[examples@3150] cat normal_strtok.c pthread_strtok.c

Conclusion

- Threading is a convenient tool to achieve the following purposes:
 - Multi-tasking within a process;
 - E.g., GUI is a thread and the main method is another one.
 - Simple shared memory environment.
- But, you have to take care of:
 - Mutual exclusion;
 - Synchronization; and
 - Thread safety.
- You should **take great care** when writing both multi-threaded and multi-processed programs.

Chapter Conclusion

- We have finished the **Chapter of Process Management**. Hope you enjoyed!

