

3150 - Operating Systems

Dr. WONG Tsz Yeung

Chapter 3, part 1 File Systems – Programmer Perspectives.

*- Here comes the chapter that you can play with something solid,
e.g., to corrupt a disk ...*

Outline

We'll teach this!

`fopen()` `fread()` `fwrite()` `fclose()`

Library Calls

We won't teach this!

`open()` `read()` `write()` `close()`

System Calls

NTFS-specific functions


Ext4-specific functions

FAT32-specific functions

ISO9660-specific functions

Kernel Functions

We will teach this!



Class Discussion: Question #1

OS == FS? Reasons?

- An OS supports a FS.
- An OS supports more than one FS.
- A FS can be used by more than one OS.
- But, “FS != OS”.


Class Discussion: Question #2

Storage Device == FS? Reasons?

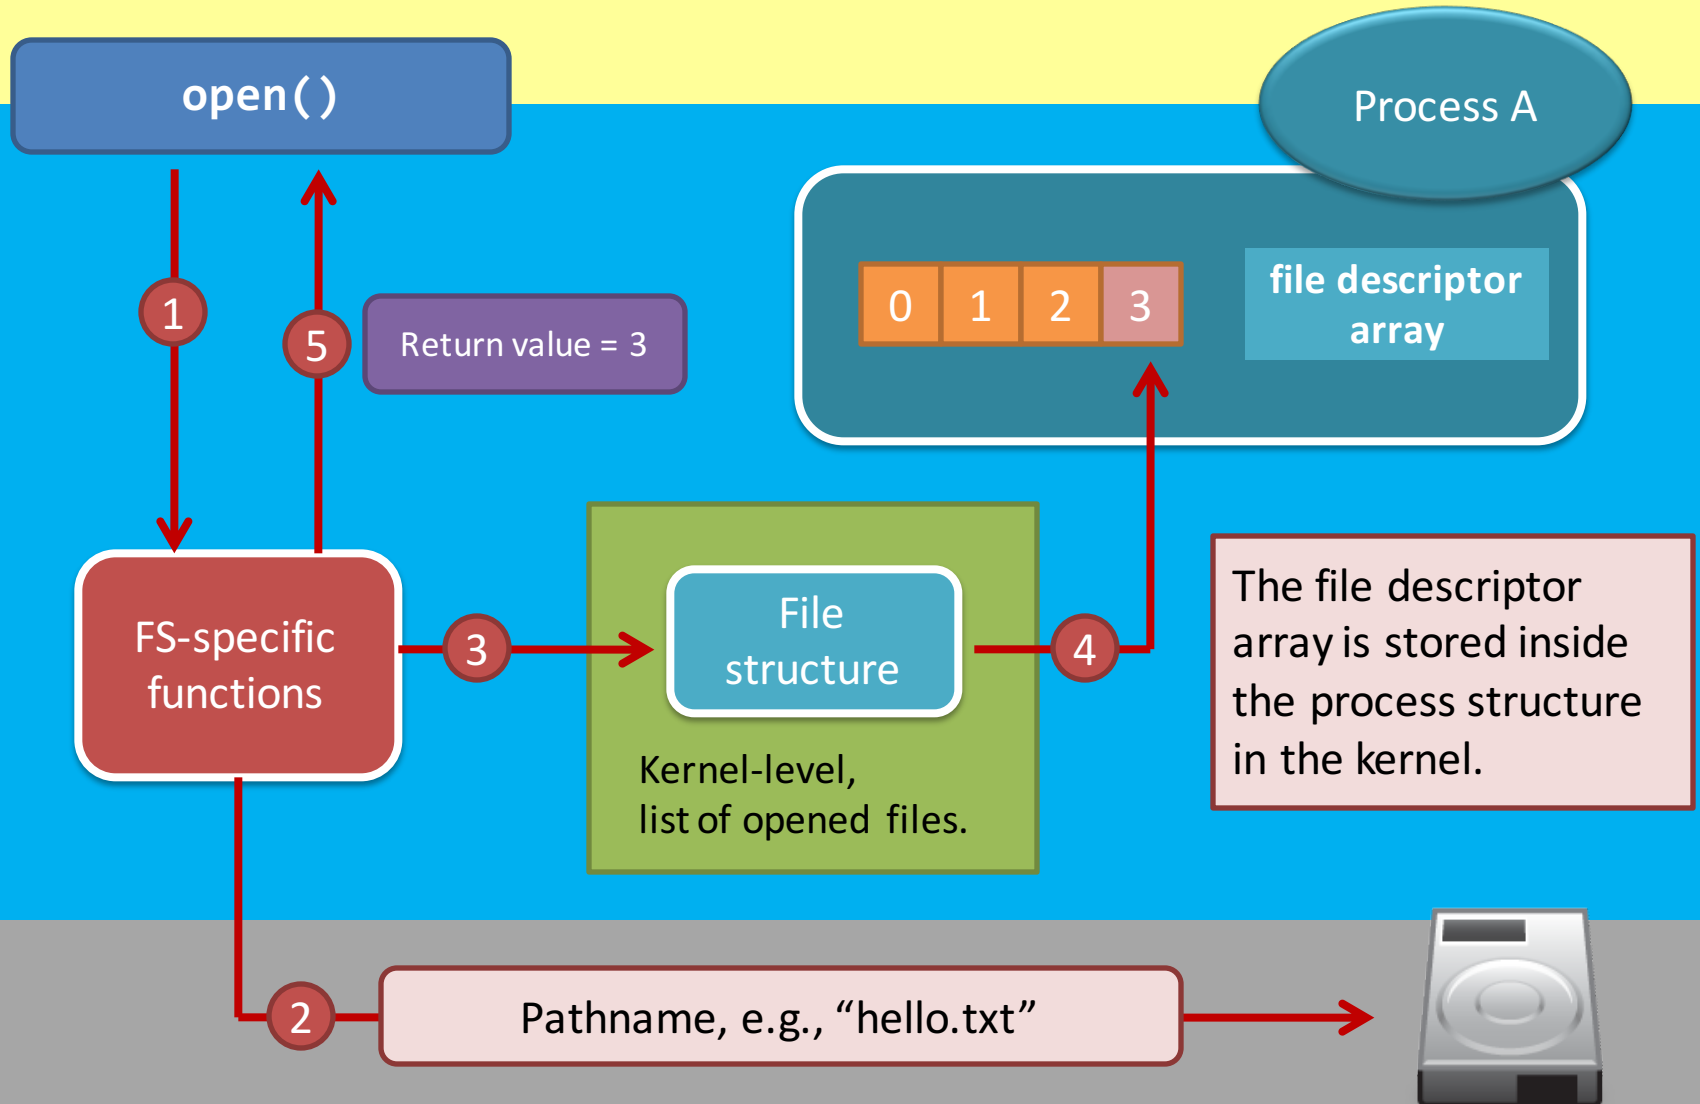
- A FS needs a storage device.
 - But, a device can be either physical or virtual.
- A storage device is just a container.
 - Doesn't need to know what FSes are stored.
 - Doesn't need to know how many FSes are stored (on different partitions).
- So, "Storage Device != FS"!

Looking at FS from the userspace

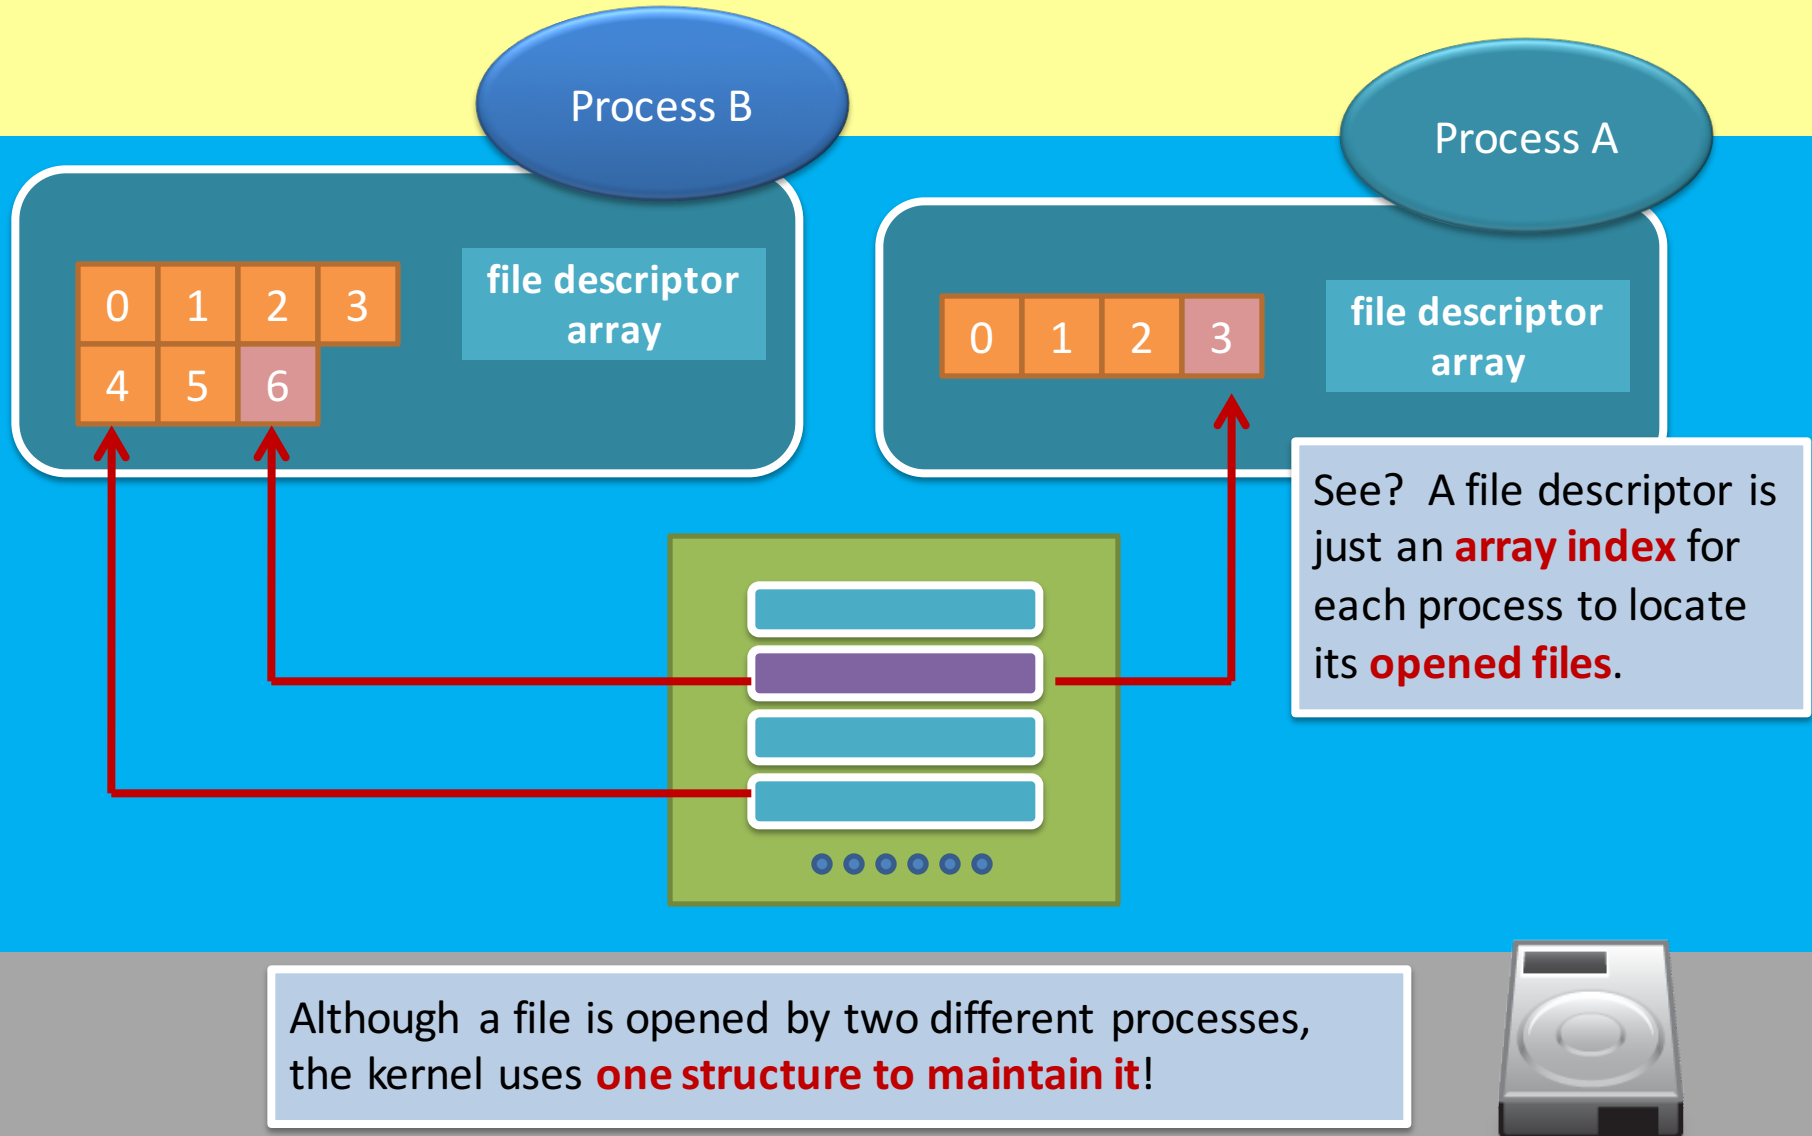
- GNU C Library call VS System call?



What is a file descriptor?




What is a file descriptor?



Library call VS System call


- What is **fopen()**? What is the type “**FILE**”?
 - First thing first, **fopen()** calls **open()**.
 - “**FILE**” is just a structure in defined in “**stdio.h**”.
 - However, **fopen()** creates memory for the “**FILE**” structure.
 - Fact: occupying space in the area of dynamically allocated memory, i.e., **malloc()**



What is inside the “FILE” structure?

- There is a lot of helpful data in **FILE**:
 - Two important things: the **file descriptor** and **a buffer**!

```
int main(void) {  
    printf("fd of stdin  = %d\n", fileno(stdin) );  
    printf("fd of stdout = %d\n", fileno(stdout) );  
    printf("fd of stderr = %d\n", fileno(stderr) );  
}
```



fileno() returns the file descriptor of the FILE structure.


The type of **stdin**, **stdout**, and **stderr** is “FILE *”

```
$ ./fileno  
fd of stdin  = 0  
fd of stdout = 1  
fd of stderr = 2  
$ _
```

[examples@3150] cat fileno.c

Looking at FS from the userspace

- GNU C Library call VS System call?
- Buffered I/O and efficiency.



read() & write()

- You know, I/O-related calls will invoke system calls.

int	read (int fd,	void *buffer,	int bytes_to_read)
-----	--------	---------	---------------	---------------------

From file to buffer.

int	write (int fd,	void *buffer,	int bytes_to_write)
-----	---------	---------	---------------	----------------------

From buffer to file.

Note: I modified the function prototypes.

Library calls that eventually invoke the
read() system call

scanf(), fscanf()

getchar(), fgetc()

gets(), fgets()

fread()

Library calls that eventually invoke the
write() system call

printf(), fprintf()

putchar(), fputc()

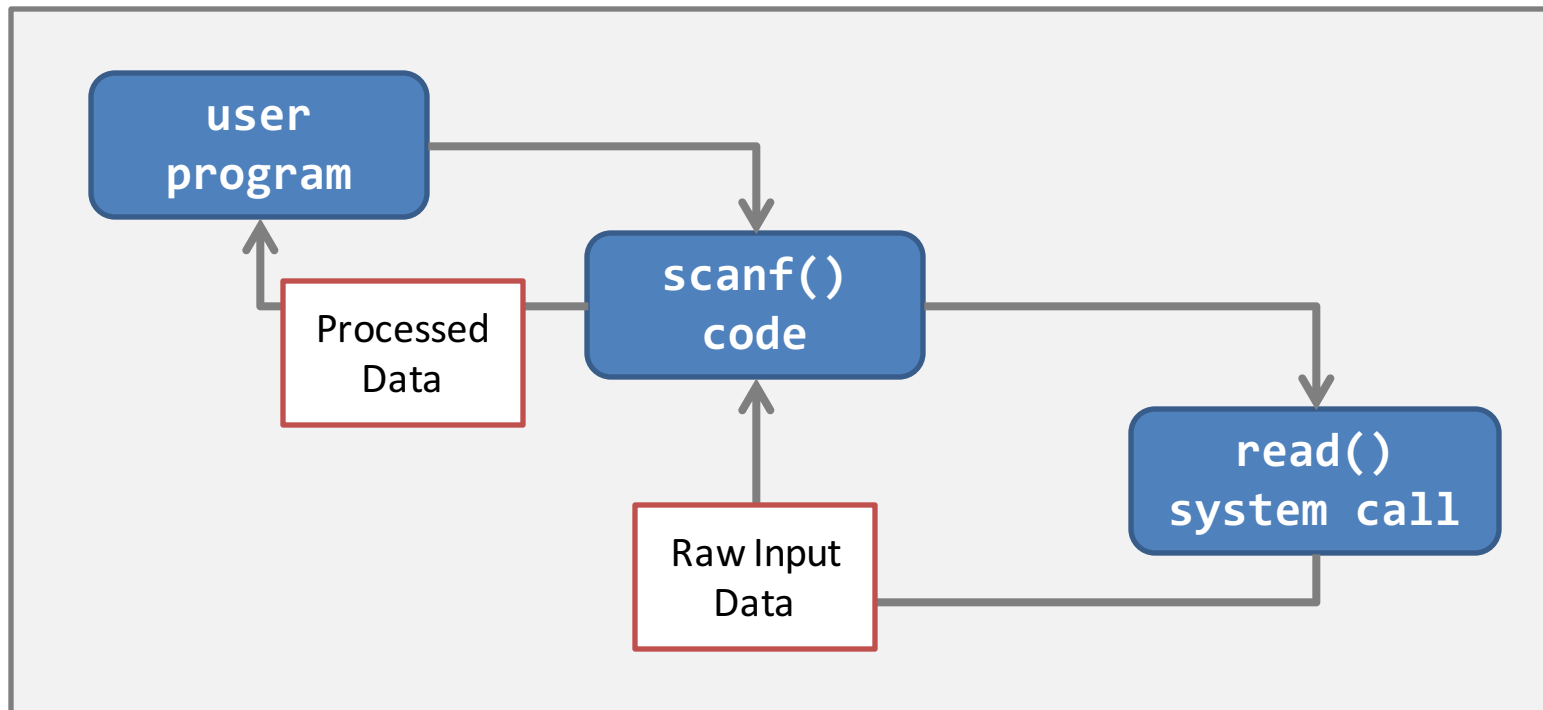
puts(), fputs()

fwrite()

read() & write()

- **scanf()** as an example!

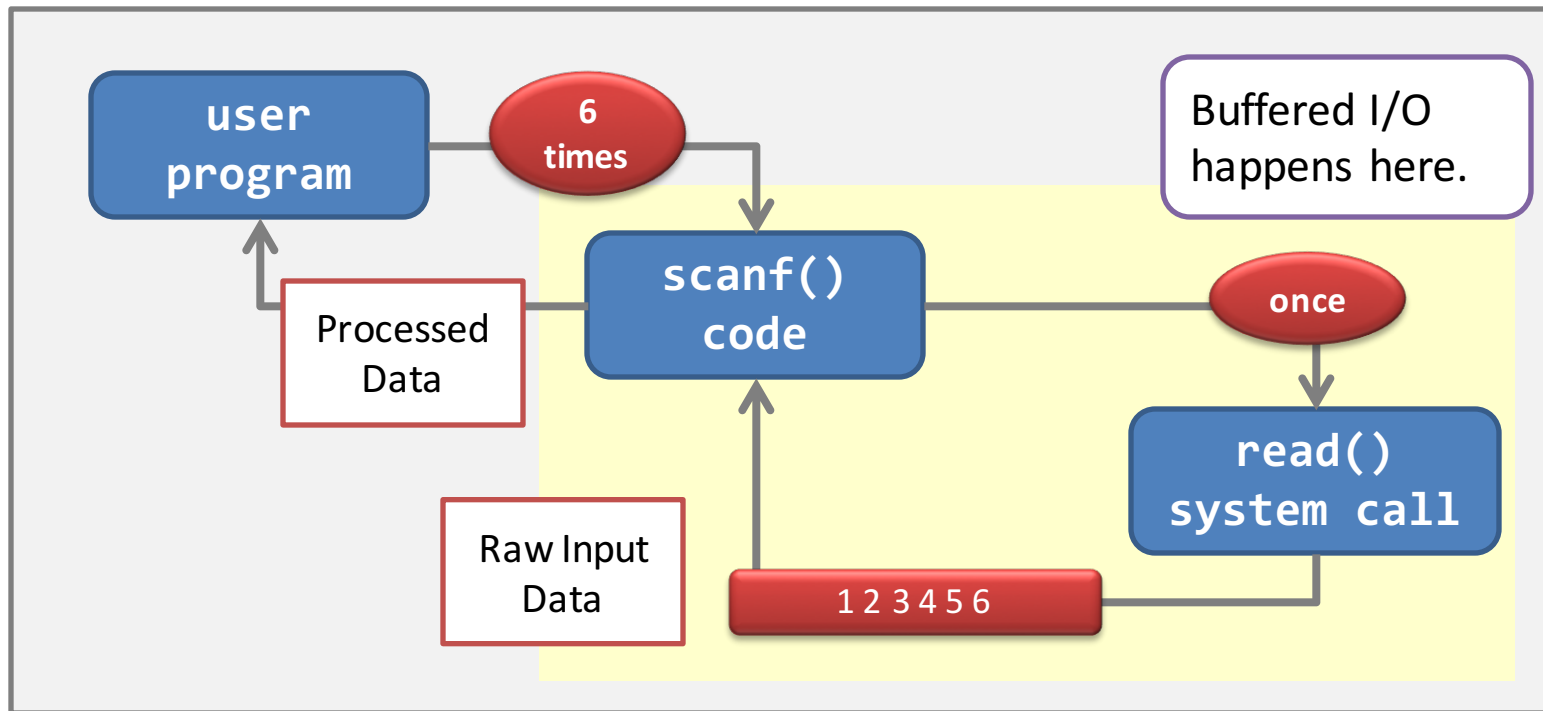
```
int main(void) {  
    int input;  
    while(1) {  
        scanf("%d", &input);  
        printf("%d\n", input);  
    }  
}
```



What is buffered I/O?

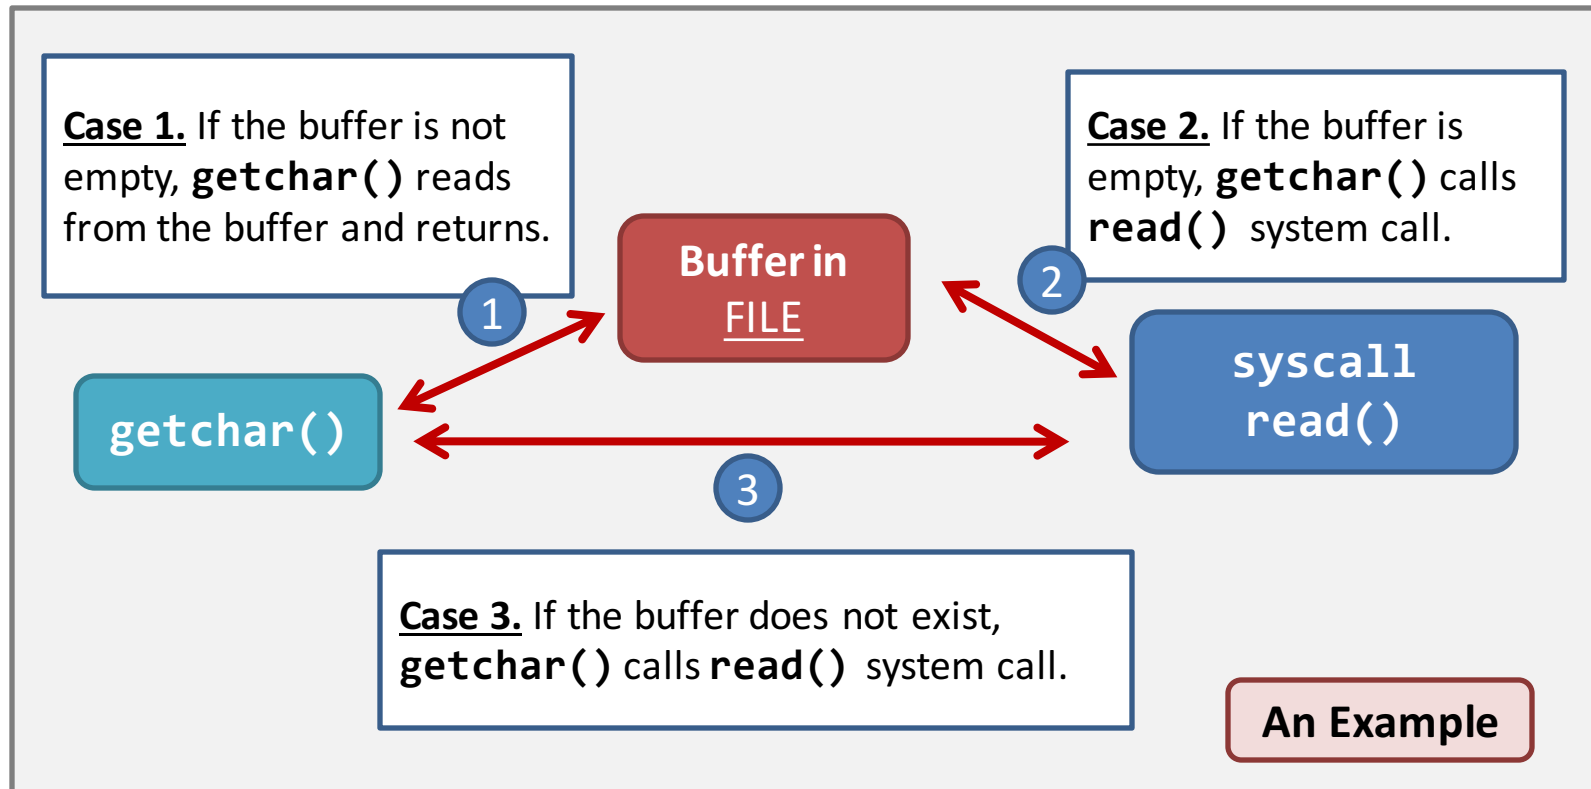
- If I input “**1 2 3 4 5 6**”, you will find 6 outputs, right?
 - But, are there 6 read() calls?

```
int main(void) {  
    int input;  
    while(1) {  
        scanf("%d", &input);  
        printf("%d\n", input);  
    }  
}
```



Buffered I/O and the “FILE” structure?

- There is a memory buffer in the FILE structure and **this cache (or buffer) to reduce the number of system calls!**



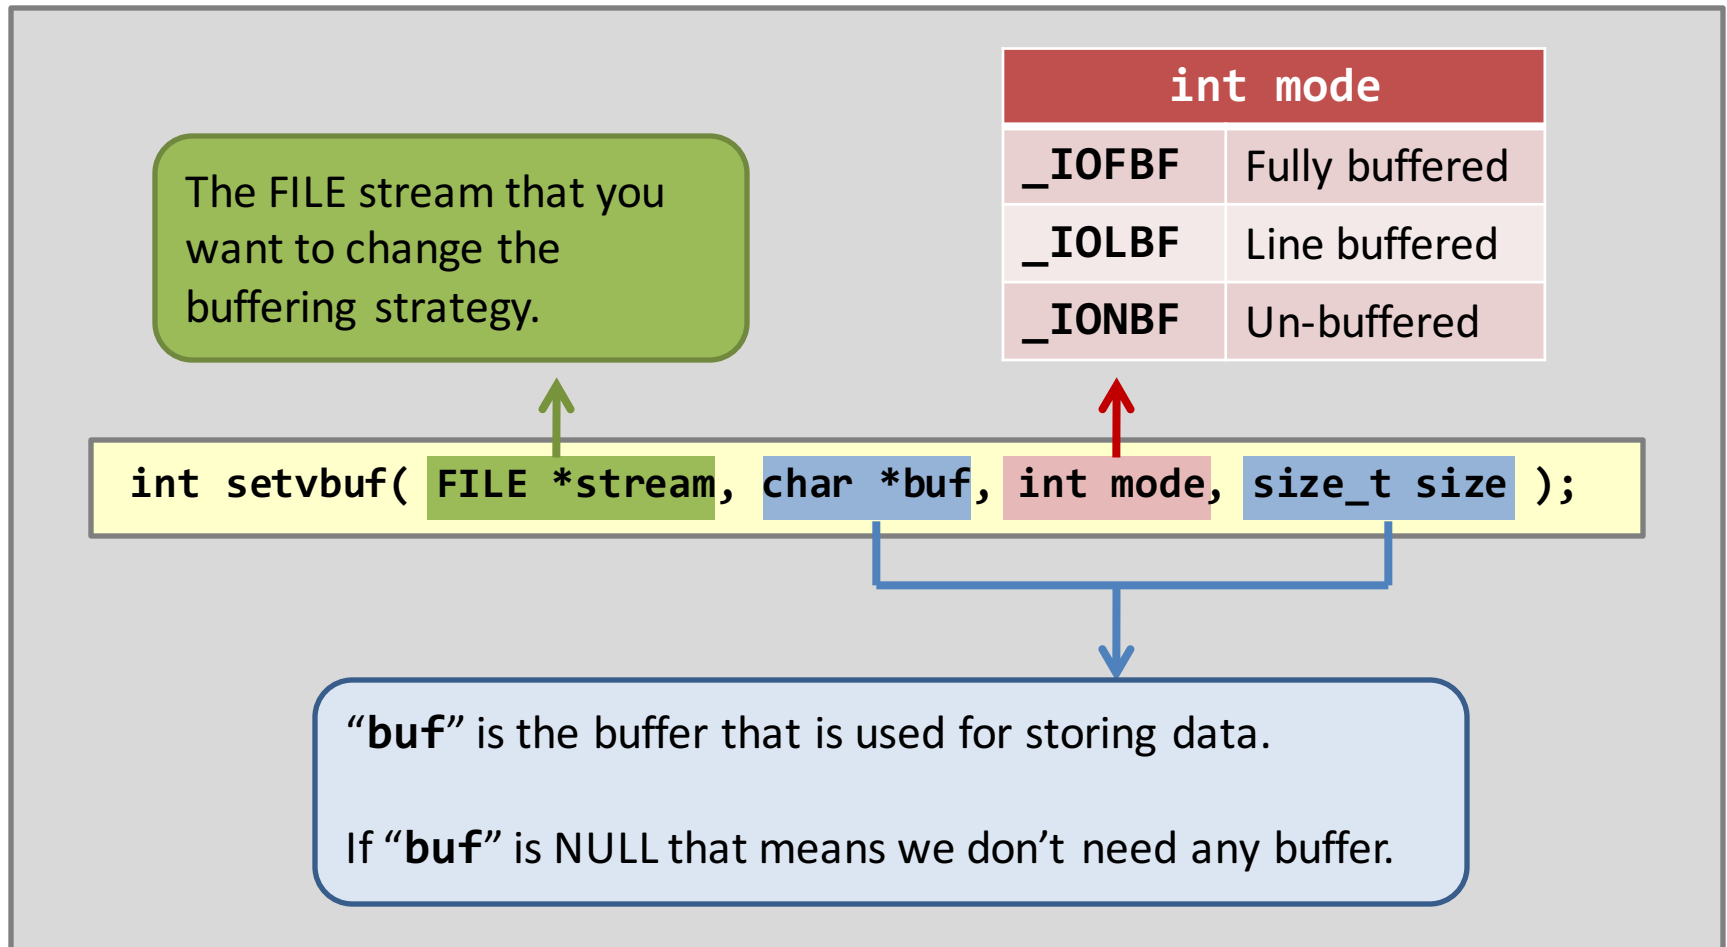
Buffered I/O – different modes

- 3 modes:

Modes	Read-related call e.g., <code>getchar()</code>	Write-related call e.g., <code>putchar()</code>
Fully-buffered	Data is read in one bulk and is stored in the buffer. Invoke the <code>read()</code> system call when the buffer becomes empty.	Data is written to the buffer. Invoke the <code>write()</code> system call when the buffer becomes full, or before the process terminates.
Line-buffered	Data is read into the buffer until the newline character is encountered.	Data is written to the buffer. When a newline character is encountered, <code>write()</code> system call is invoked.
Un-buffered	Directly translate every library call into a <code>read()</code> system call.	Directly translate every library call into a <code>write()</code> system call.

Buffered I/O – change the buffer

- There is a convenient call that controls **everything!**



Buffered I/O – change the buffer

- “**stdin**” and “**stdout**” are **line-buffered** by default.
- “**stderr**” is **un-buffered** by default.

```
char buf[1024];
int main(void) {
    .....
    printf("H");    sleep(1);
    printf("E");    sleep(1);
    printf("L");    sleep(1);
    printf("L");    sleep(1);
    printf("O");    sleep(1);
    printf("\n");   sleep(1);
}
```



```
setvbuf(stdout, NULL, _IONBF, 0);
```

```
setvbuf(stdout, buf, _IOLBF, sizeof(buf));
```

```
setvbuf(stdout, buf, _IOFBF, sizeof(buf));
```

[examples@3150] cat no_buf.c line_buf.c full_buf.c


- Now, you know the buffer is just a piece of memory.
 - So, you need to be careful when you are playing with **fork()** and **pthread_create()**.
- **Challenge.** What will be the output?

```
int main(void) {  
    printf("Hello");  
    fork();  
    printf("\n");  
    return 0;  
}
```

[examples@3150] cat fork.c

Looking at FS from the userspace

- GNU C Library call VS System call?
 - Buffered I/O and efficiency.
- What is the true meaning of EOF?



Library call VS System call – what is EOF?

- Well, the following is just one of the common usage of EOF:

```
int main(void) {  
    char c;  
    unsigned long long count = 0;  
    while(1) {  
        c = getchar();  
        if(c == EOF)  
            break;  
        else  
            count++;  
    }  
    printf("EOF! Read %lld bytes.\n", count);  
}
```

Do you know what
EOF really is?

[examples@3150] cat getchar_eof.c

Library call VS System call – what is EOF?

- First of all, **you can't find any "EOF character"** when using system calls.

```
int main(void) {  
    int ret;  
    char c;  
    unsigned long long count = 0;  
    while(1) {  
        ret = read(fileno(stdin), &c, 1);  
        if(ret == 0)  
            break;  
        else {  
            count += ret;  
            if(c == EOF)  
                printf("WoW!\n");  
        }  
    }  
    printf("Read %lld bytes.\n", count);  
}
```

No more bytes
to read.

Any "WoW!"?

[examples@3150] cat read_eof.c

Library call VS System call – what is EOF?

- Somewhere inside “**/usr/include/stdio.h**”:

```
#ifndef EOF
# define EOF (-1)
#endif
```


- That means: all those “**f*()**” functions *memorize* whether the end of file is reached or not!
 - If yes, it just returns -1 (EOF)!
 - If no, it either reads data from the buffer or system calls.
- **Main point: No EOF character in any files!**

Summary

- The GNU I/O library functions give you a lot of convenience:
 - Great functions: `fscanf()`, `fprintf()`
- Yet, this abstracts away many truths and introduces (not on purpose) mis-conceptions!
 - From now on, you should never said: “*An empty file has an EOF character at the end of the file!*”
 - Also, use *feof()* with great care!

Looking at FS from the userspace

- GNU C Library call VS System call?
 - Buffered I/O and efficiency.
 - What is the true meaning of EOF?
- **File and directory.**
 - **basics;**



Attributes

- First, both files and directories are very similar:
 - They **both** have two kinds of data: **attributes** and **data**.
 - Attributes are as important as data:
 - Can you read the data correctly *without the size attribute*?

The design of FAT32 does not include any security ingredients.

Common Attributes	FAT32	NTFS	Ext2/3/4
Name	✓	✓	✓
Size	✓	✓	✓
Permission		✓	✓
Owner		✓	✓
Access, creation, modification time	✓	✓	✓

Reading attributes

- The command is **stat**. You can find:
 - type, size, permission, etc.
- The system call counterpart includes:
 - **stat()**, **fstat()**, and **lstat()**.

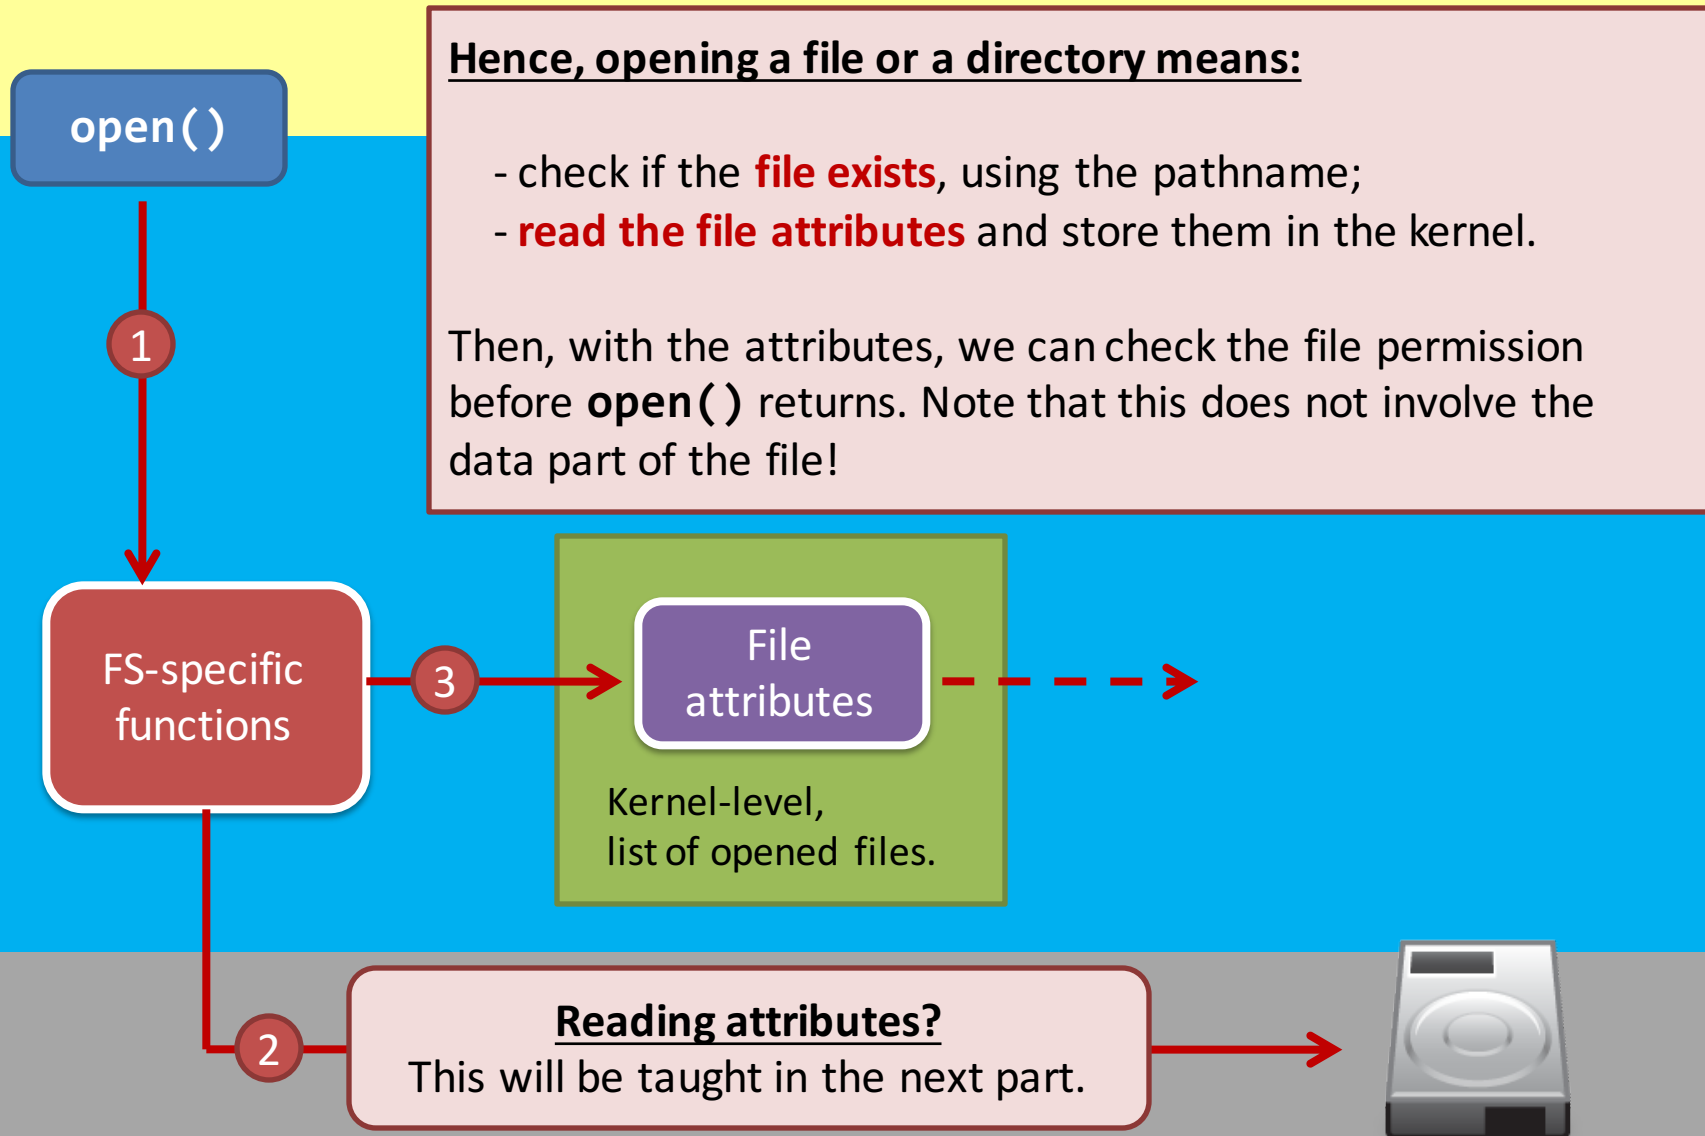
```
# stat /
File: '/'
Size: 4096      Blocks: 8      IO Block: 4096
Device: 802h/2050d Inode: 2      Links: 22
Access: (0755/drwxr-xr-x)  Uid: (  0/   root)  Gid: (  0/   root)
Access: 2008-11-01 13:53:35.000000000 +0800
Modify: 2008-11-01 13:42:30.000000000 +0800
Change: 2008-11-01 13:42:30.000000000 +0800
# _
```

Writing attributes?

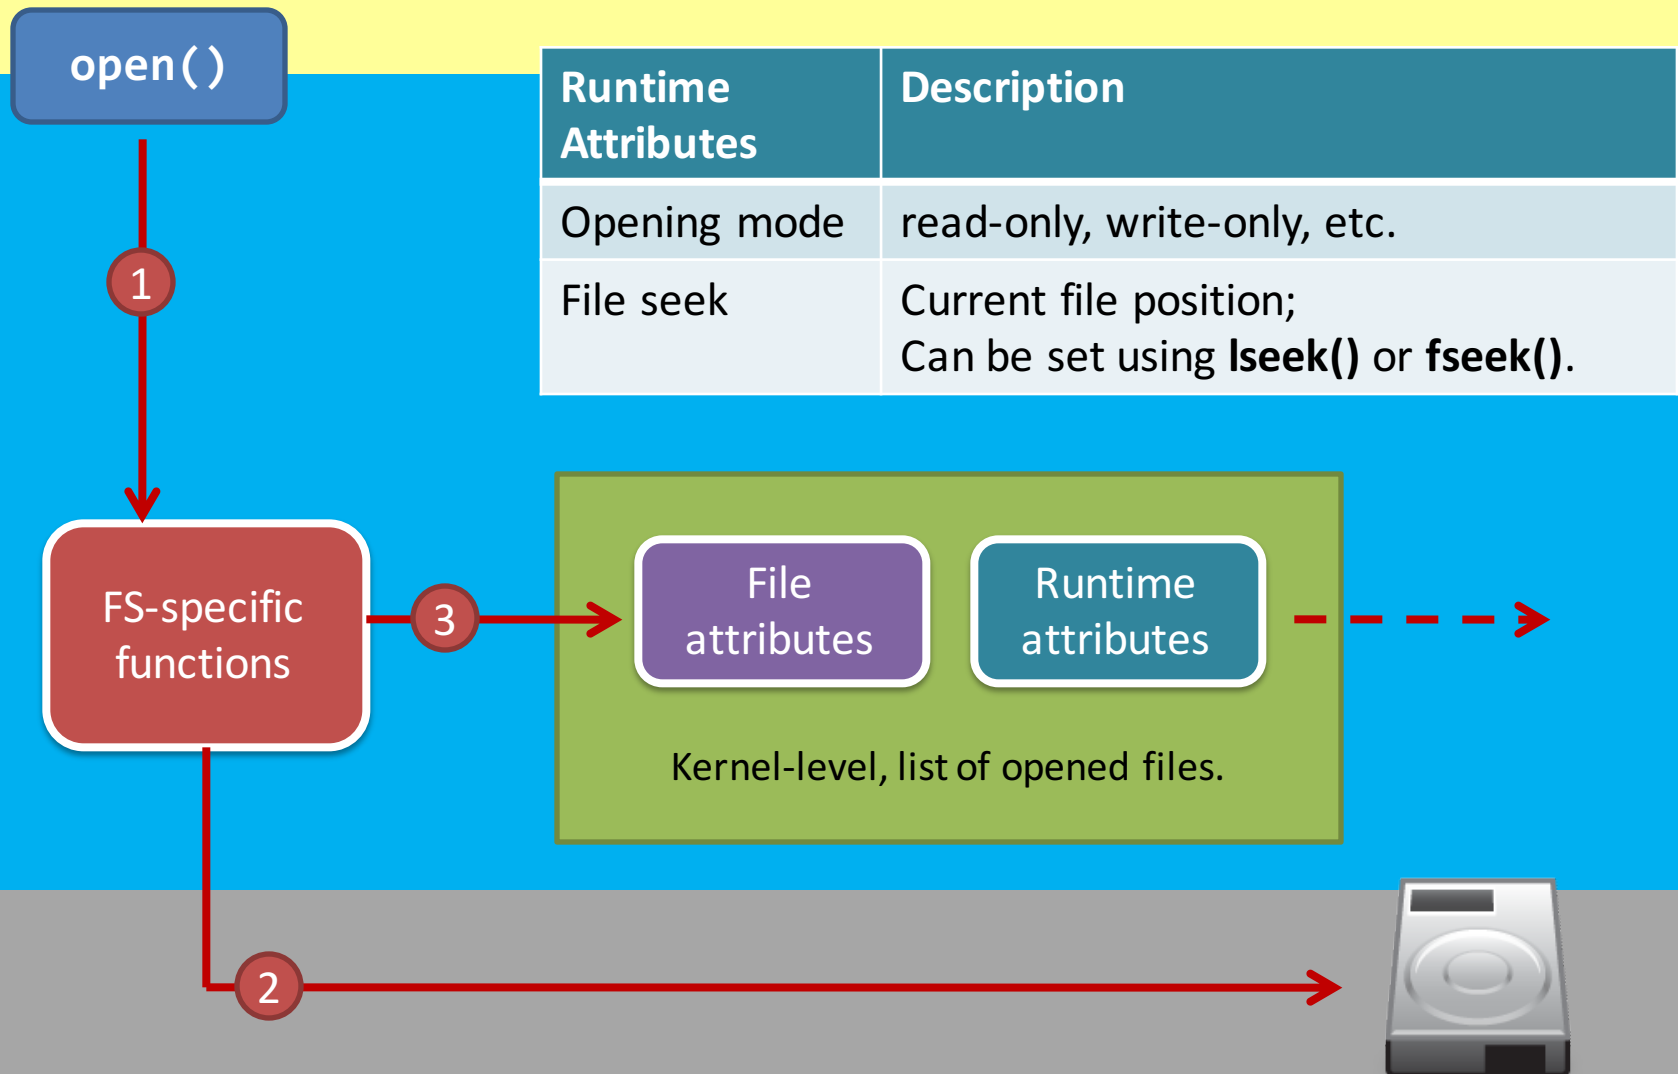
- Can you change those attributes directly?

Common Attributes	Way to change them?	
	Command?	Syscall?
Name	mv	rename()
Size	Too many tools to update files' contents	write(), truncate(), etc.
Permission	chmod	chmod()
Owner	chown	chown()
Access, creation, modification time	touch	utime()

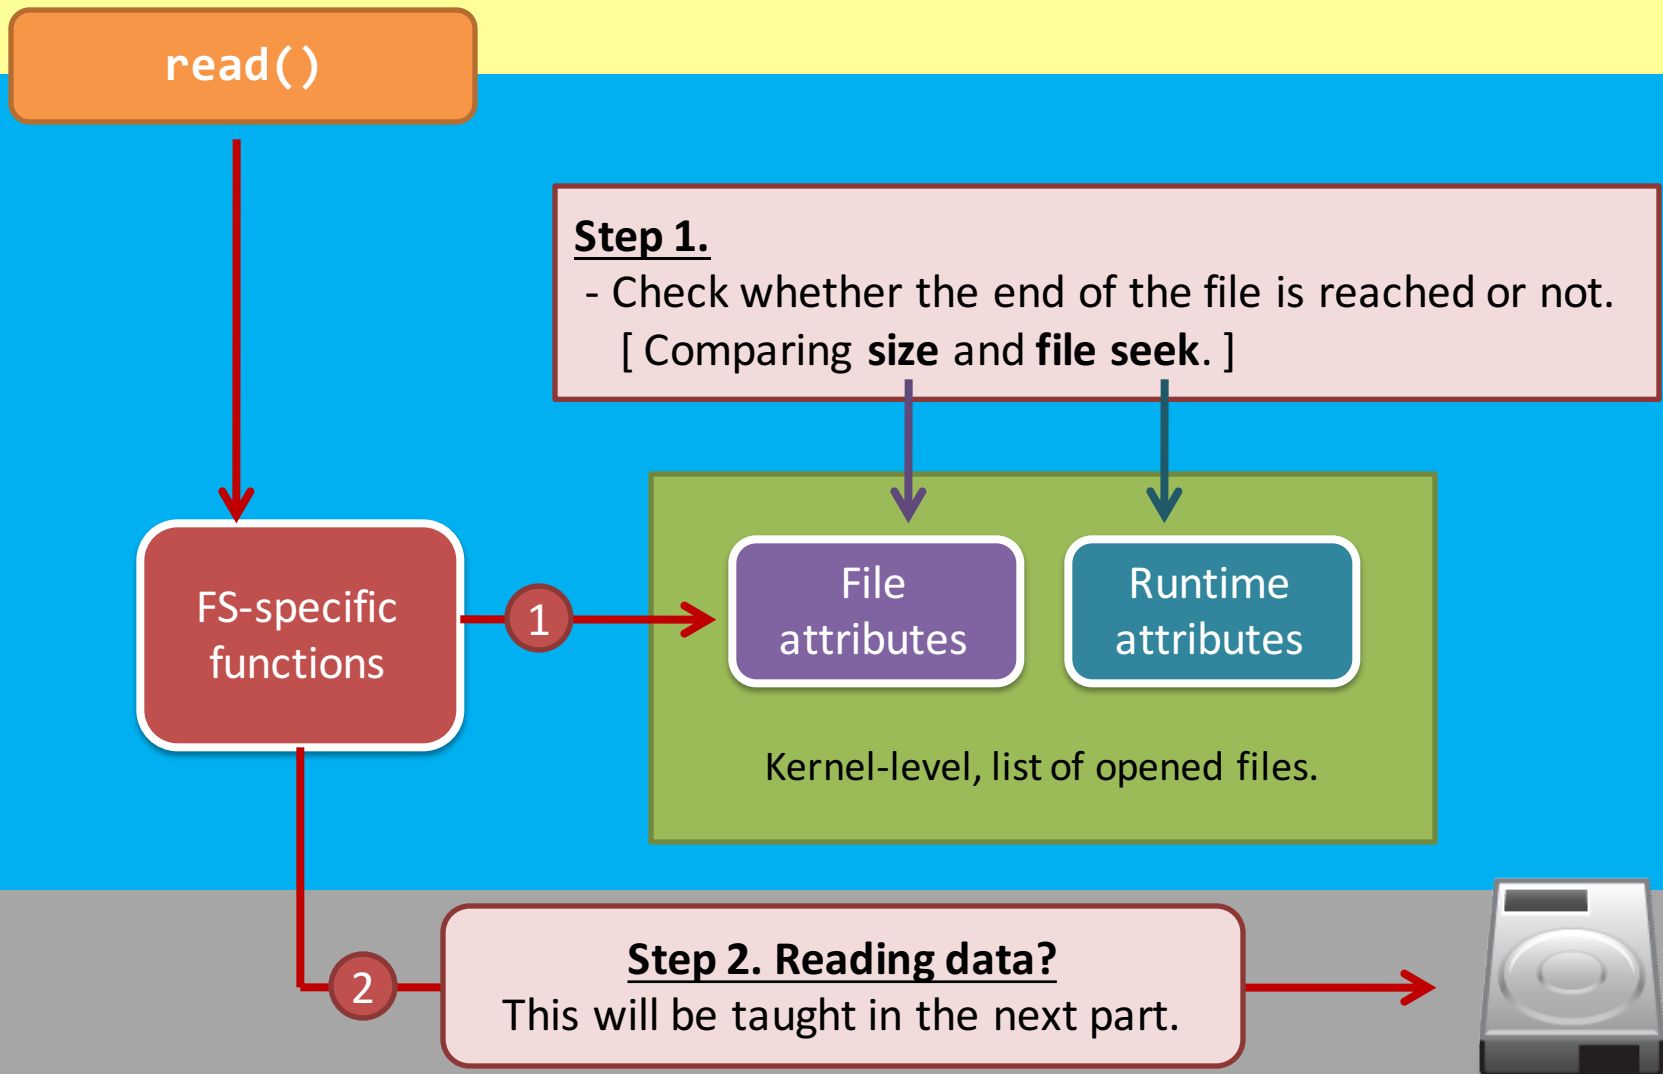
Still remember the flow of `open()`?



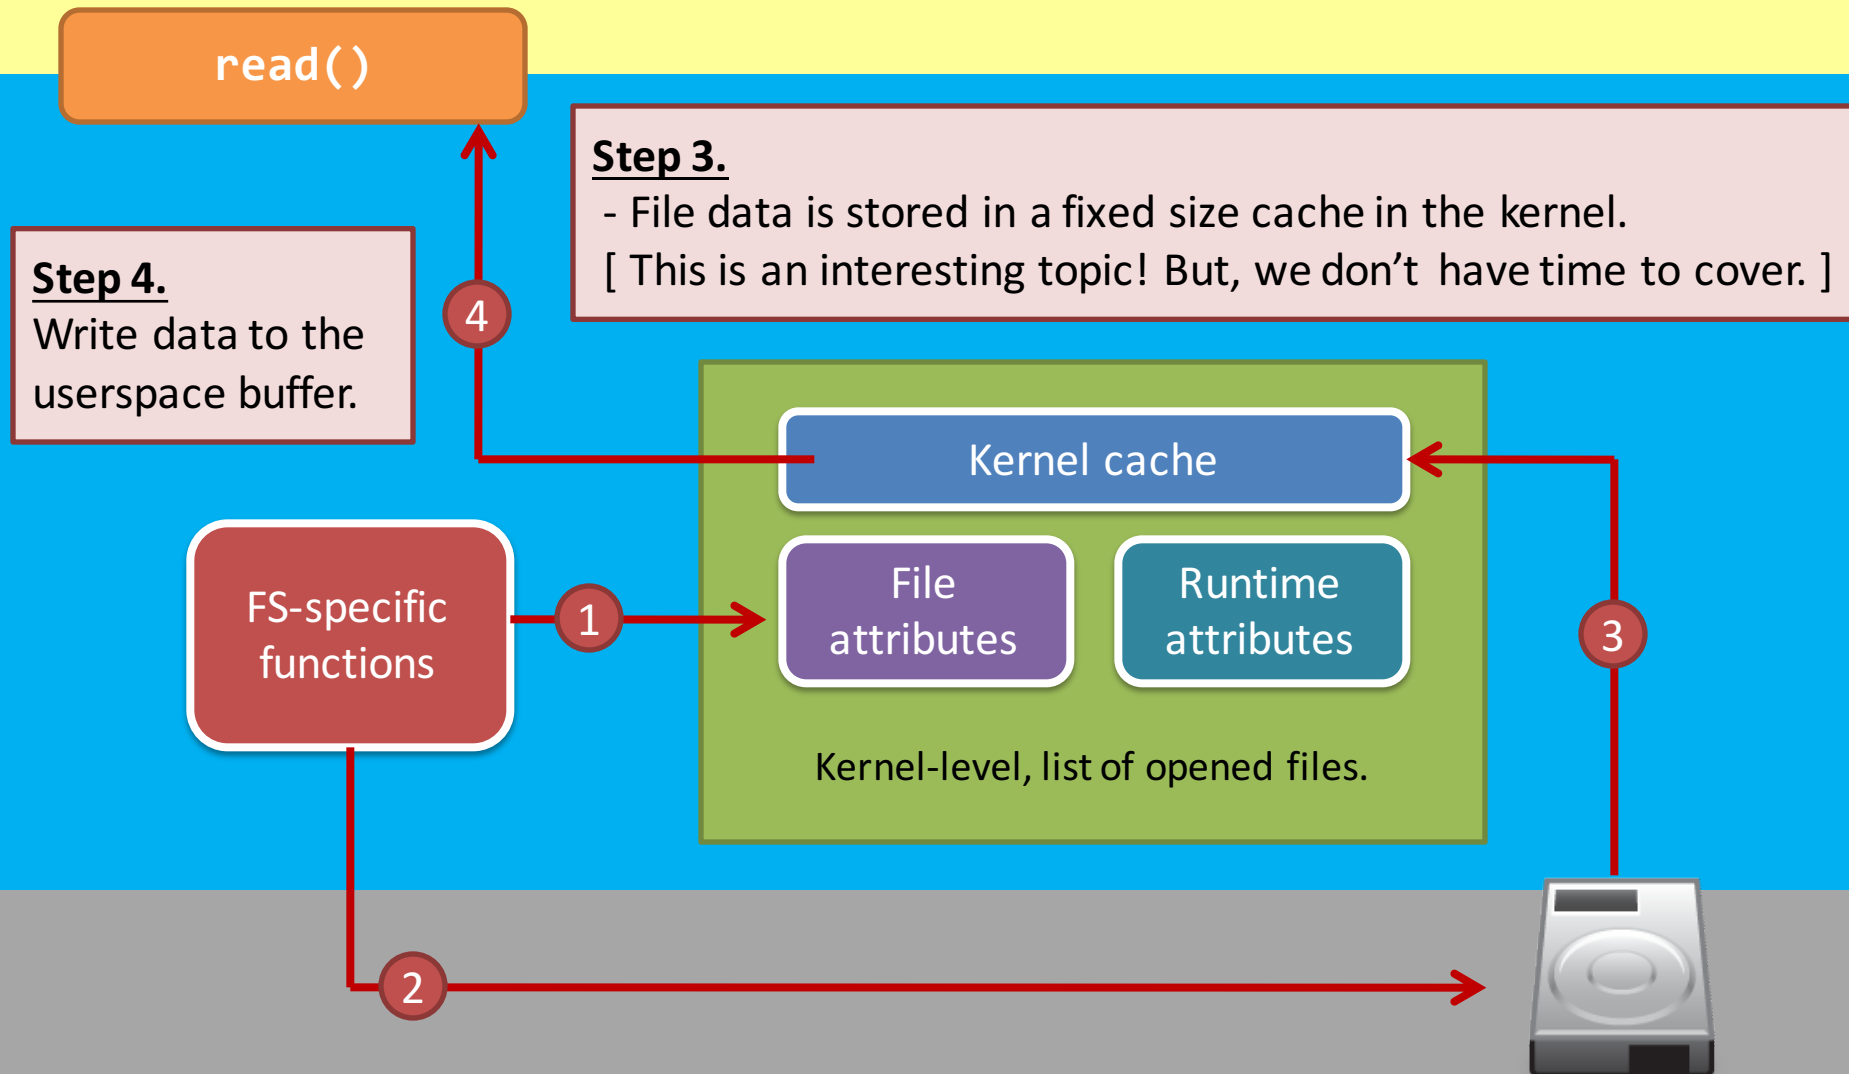
Still remember the flow of `open()`?



How about the **read()** system call?



How about the `read()` system call?



How about the `write()` system call?

`write()`

Step 1.

Write data to the kernel buffer.

1

Step 2.

According to the data length, (1) change in file size, if any, and (2) change in the file seek.

3

Step 3.

The call returns.

2

Kernel cache

2

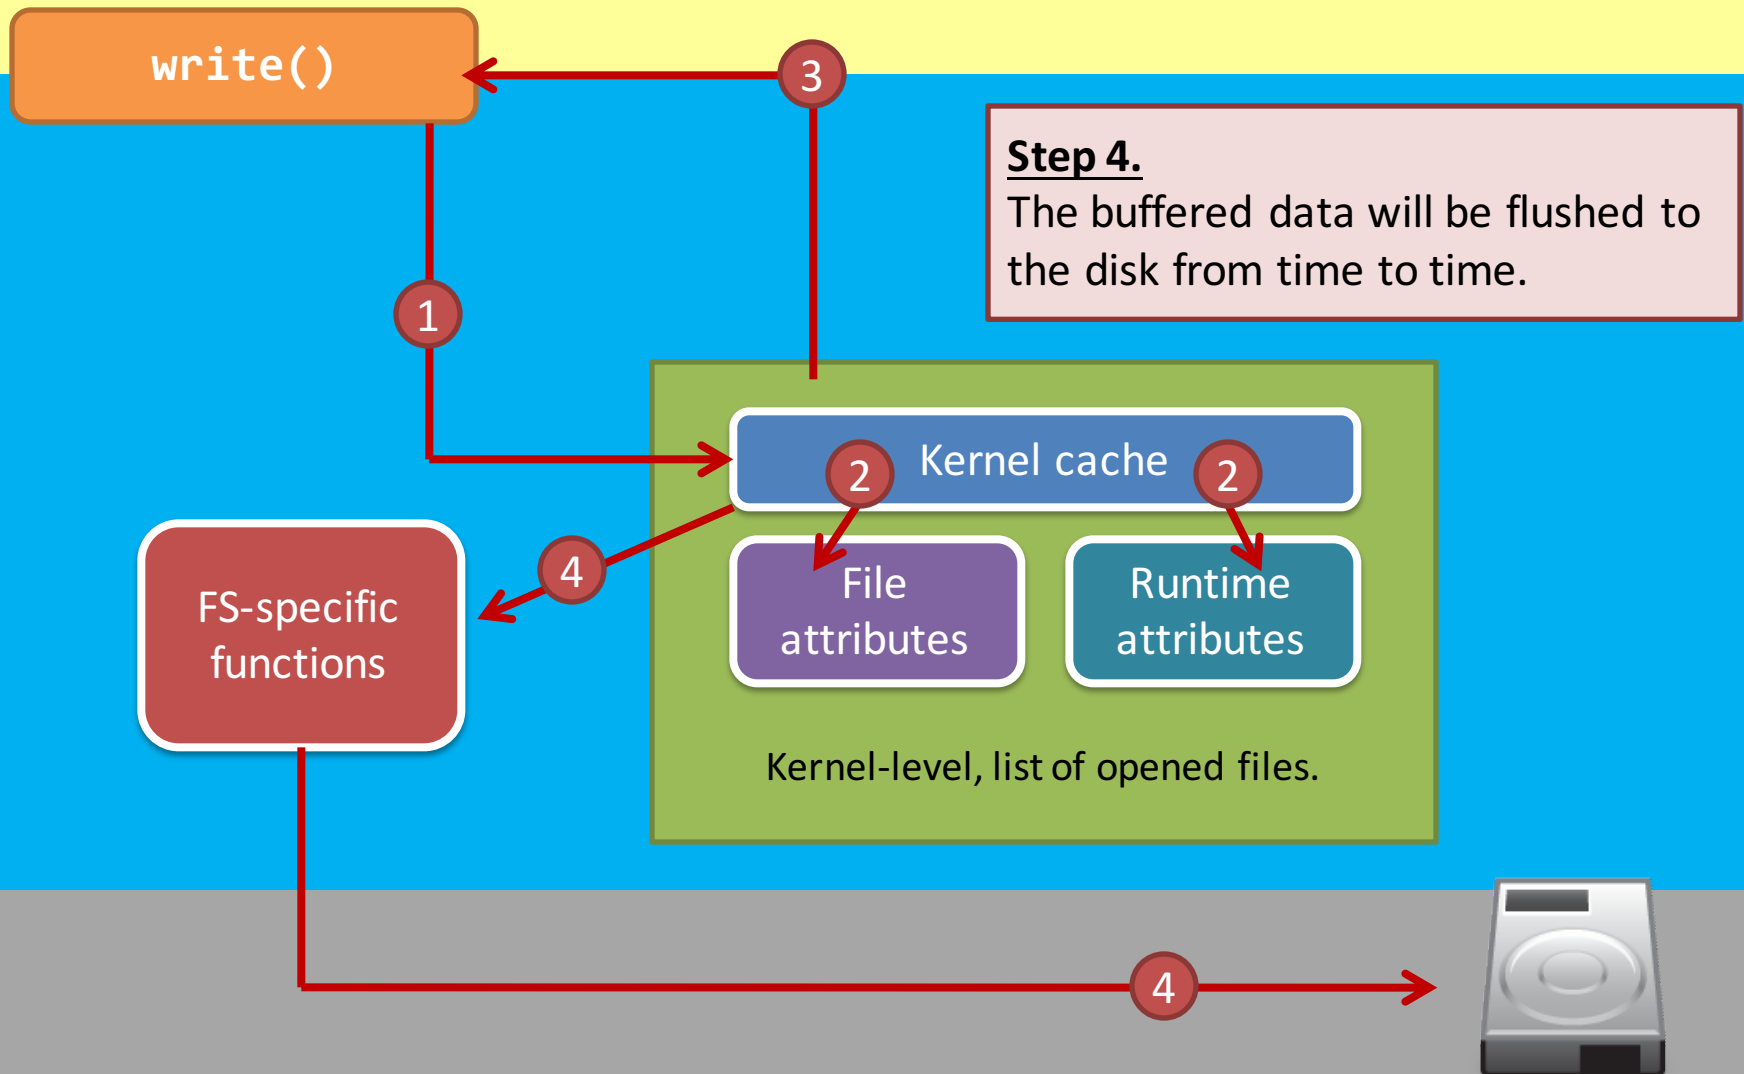
File attributes

Runtime attributes

Kernel-level, list of opened files.



How about the `write()` system call?




The kernel buffer cache implies...

- Class Discussion!
 - Increase reading performance?
 - Increase writing performance?
 - Can you answer me why you cannot press the reset button?
 - Can you answer me why you need to press the “eject” button before removing USB drives?

Looking at FS from the userspace

- GNU C Library call VS System call?
 - Buffered I/O and efficiency.
 - What is the true meaning of EOF?
- **File and directory.**
 - basics;
 - playing with directories.



Reading a directory

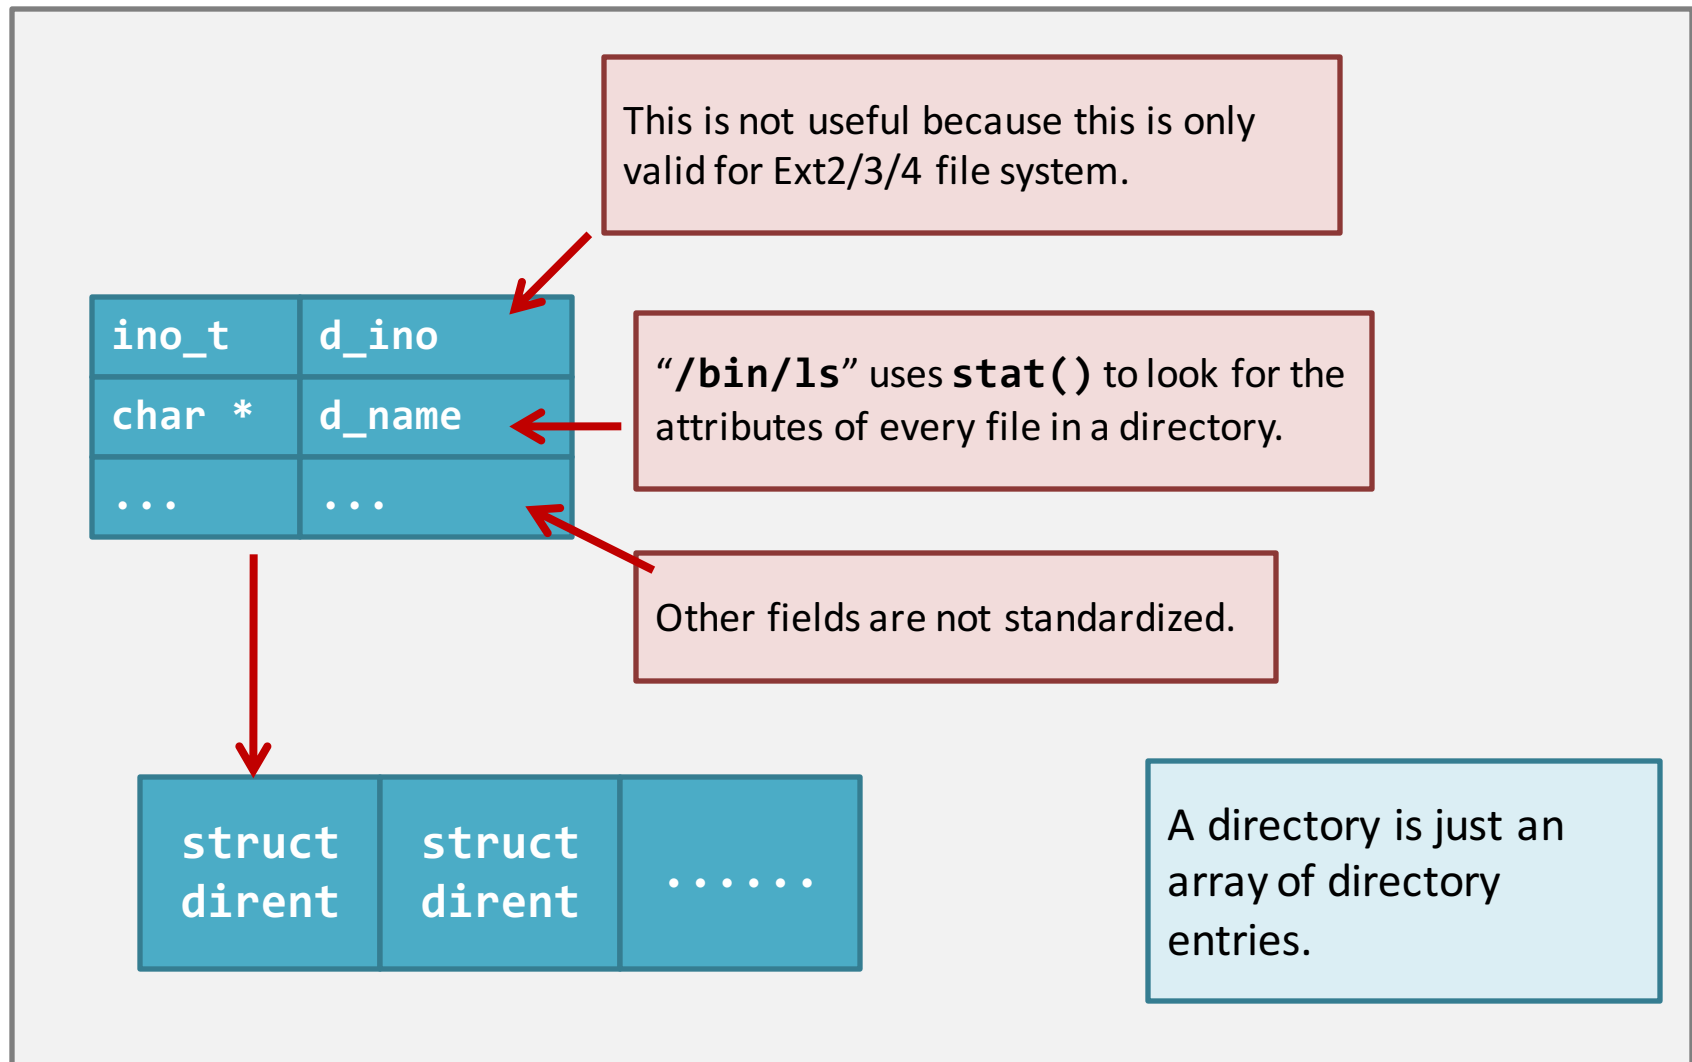
- Similarly, reading a directory also involves attributes and data.
 - Yet, you don't know how to play with in the userspace.

```
int main(int argc, char **argv) {
    DIR *dir;
    char *input = "/";
    struct dirent *entry;

    dir = opendir(input);                // open
    while( (entry = readdir(dir)) != NULL ) {    // read
        printf("%ld\t\t%s\n",
               (long) entry->d_ino,    // unique ID
               entry->d_name);        // name, max char: 255
    }
    closedir(dir);                        // close
}
```

[examples@3150] cat simple_ls.c

Directory entries – read



Directory entries – write

Removing an existing file also means writing to a directory. But, the write operation is to erase something...

The truth of file deletion: the system **will not** remove the corresponding directory entry in hosting directory completely.

[What evil things can you think of?]

Add a new file to a directory means appending data to the directory file.



This is a directory file: it contains an array of directory entries.

Directory entries – write

- Example screenshot.

```
$ ls -ld .  
drwxr-xr-x 11 tywong tywong 4096 2010-11-16 22:20 .  
  
$ ls -l final.xls  
ls: cannot access final.xls: No such file or directory
```

```
$ touch final.xls  
  
$ ls -l final.xls  
-rw-r--r-- 1 tywong tywong 0 2010-11-21 16:05 final.xls  
  
$ ls -ld .  
drwxr-xr-x 11 tywong tywong 4096 2010-11-21 16:05 .
```

```
$ rm final.xls  
rm: remove regular empty file `final.xls'? y  
  
$ ls -ld  
drwxr-xr-x 11 tywong tywong 4096 2010-11-21 16:06 .
```

Create a file updates the directory file.

Remove a file also updates the directory file.

Summary

- Through this part, we learn:
 - the truth about the calls that we usually use,
 - the content of a file is not the only entity, but also the file attributes.
- In the next part, we will go into the disk:
 - How and where to store the file attributes?
 - How and where to store the data?
 - How to manage a disk?