

# CSCI3150 – Tutorial 5

---

ASSIGNMENT 1, PHASE I REVISITED

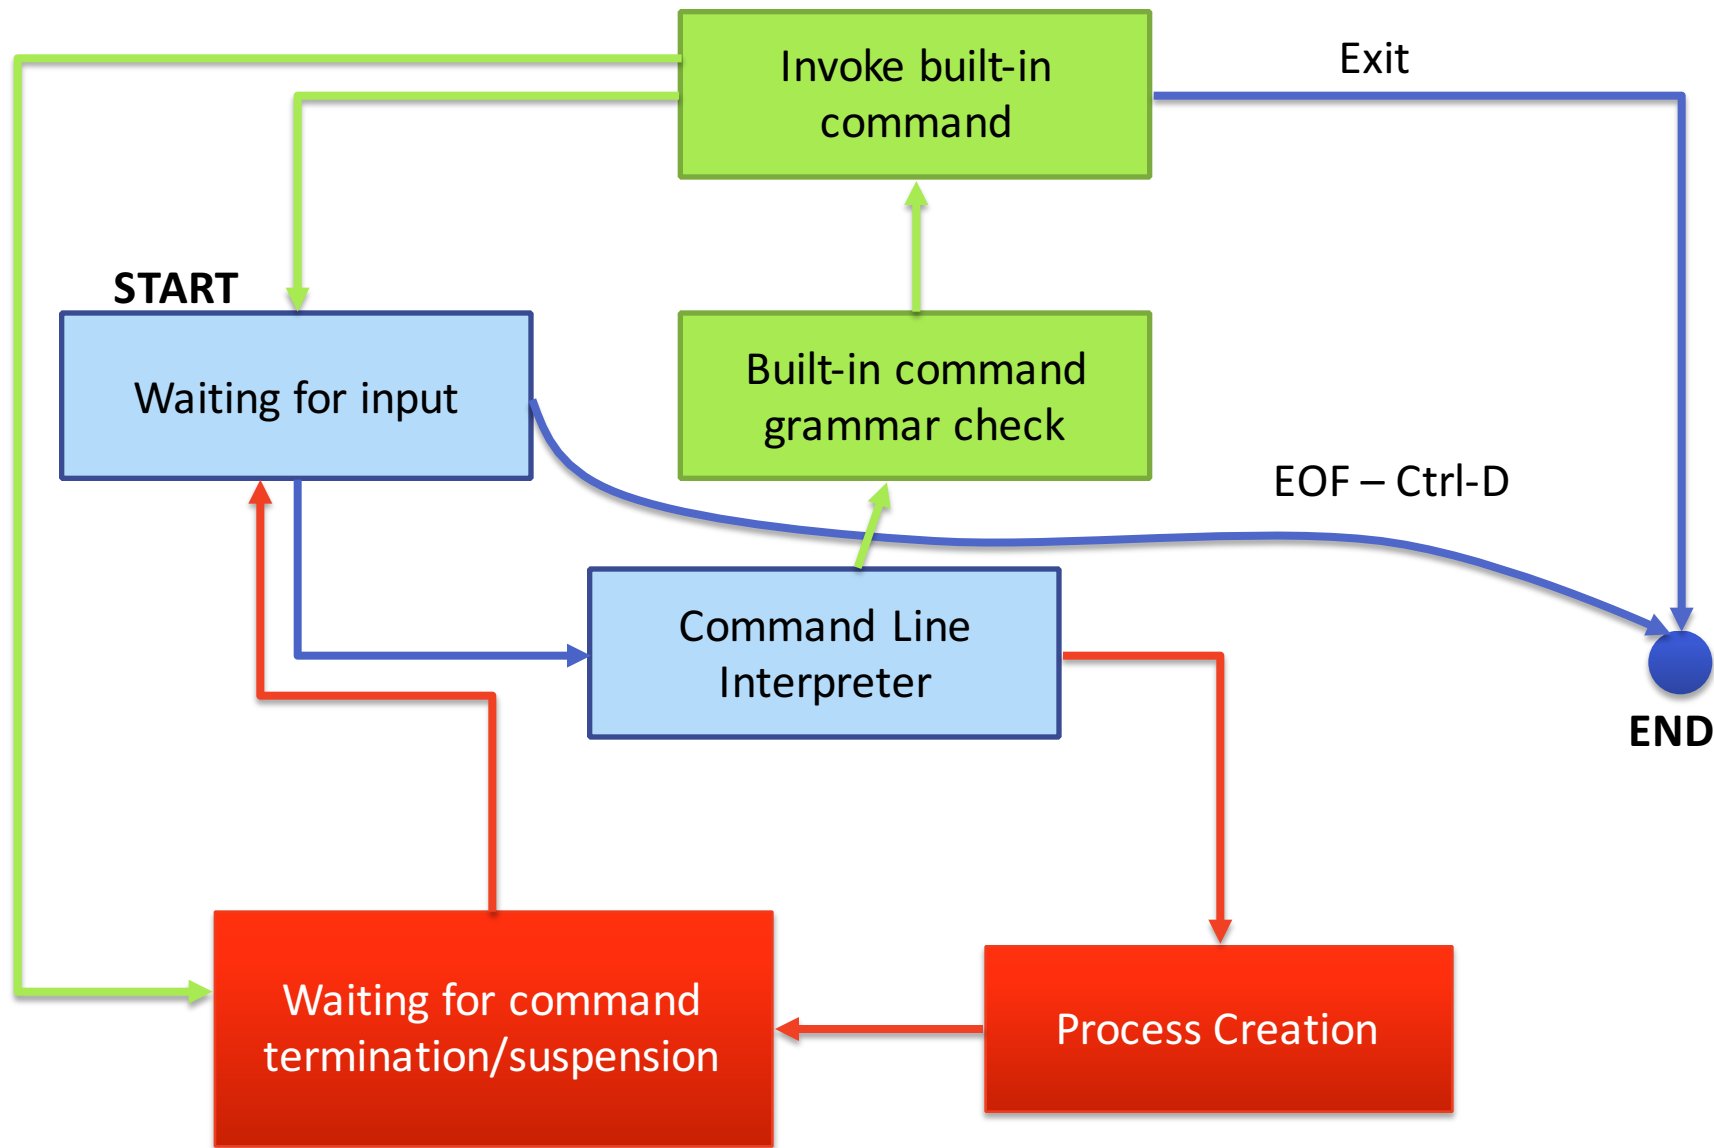
Calvin Kam <[hckam@cse.cuhk.edu.hk](mailto:hckam@cse.cuhk.edu.hk)>

# Phase 1 – Shell

---

You will need to write a workable “shell” that can

1. Accept user inputs
2. Analyze the input command line and distinguish if it is built-in commands or others
3. If it is built-in commands:
  1. Change working directories, or exit the program
4. If it is other commands:
  1. Set the search path
  2. Expand the wildcard expression
  3. Fork a new process, set the signal handlers and execute.
  4. Wait for the termination of the child.



# Phase 1 – Shell

---

You will need to write a workable “shell” that can

1. Get user inputs
2. Analyze the input command line and distinguish if it is built-in commands or others
3. If it is built-in commands:
  1. Change working directories, or exit the program
4. If it is other commands:
  1. Set the search path
  2. Expand the wildcard expression
  3. Fork a new process, set the signal handlers, and execute.
  4. Wait for the termination of the child.

# Get the User Input

<stdio.h>

```
[3150 Shell:/home/hckam]$ ls -al
```

- Get the whole input including spaces.... “ls -al” whether than “ls” alone.
- fgets() is recommended, as it is easy to use and get all the things you need.

```
#include <stdio.h>
#include <string.h>

int main(int argc, char *argv[]){
    char buf[255];
    printf("[I am Shell]$:");
    fgets(buf, 255, stdin);
    buf[strlen(buf)-1] = '\0';
    printf("%s\n", buf);
    return 0;
}
```

# Get the User Input

---

If the user presses Ctrl-D (End of file, NOT SIGNAL!!!), then your shell should terminate.



We can detect the EOF by checking the return value of `fgetc()`!

```
char buf[255];

if(fgets(buf,255,stdin) != NULL)
    printf("[%s]\n",buf);

printf("\nThe program has terminated!\n");
return 0;
```

# Phase 1 – Shell

---

You will need to write a workable “shell” that can

1. Get user inputs
2. Analyze the input command line and distinguish if it is built-in commands or others
3. If it is built-in commands:
  1. Change working directories, or exit the program
4. If it is other commands:
  1. Set the search path
  2. Expand the wildcard expression
  3. Fork a new process, set the signal handlers, and execute.
  4. Wait for the termination of the child.

# Tokenize the input

<string.h>

After you get the input, the string is “chopped” into words (tokens) in order to be examined (for example, check whether it is a built-in command) and build the argument list.

A useful string function - `char *strtok(char *str, const char *delim);`

```
#include <string.h> //Needed by strtok()
char buf[] = "Hello World tywong sosad";
char *token = strtok(buf, " ");
while(token != NULL)
{
    printf("%s\n", token);
    token = strtok(NULL, " ");
}
```



Hello

World

tywong

sosad



# Built-in Commands

---

In this assignment, we only have two built-in commands!

For the kind of **built-in commands**, you need to write your own code.

1. Change Directory (cd)
2. Exit the shell (exit)
  - Break the while loop, and we are done : )

# Phase 1 – Shell

---

You will need to write a workable “shell” that can

1. Get user inputs
2. Analyze the input command line and distinguish if it is built-in commands or others
3. If it is built-in commands:
  1. Change working directories, or exit the program
4. If it is other commands:
  1. Set the search path
  2. Expand the wildcard expression
  3. Fork a new process, set the signal handlers, and execute.
  4. Wait for the termination of the child.

# Get Current Directory

<unistd.h>  
<limits.h>

In the prompt, we need to know the current working directory.

Initially the path is where the program is located.

The useful function is - `char *getcwd(char *buf, size_t size);`

PATH\_MAX is  
the maximum  
characters  
allowed in the  
pathname.

```
#include <stdio.h>
#include <limits.h> // Needed by PATH_MAX
#include <unistd.h> // Needed by getcwd()

int main(int argc, char *argv[]){
    char cwd[PATH_MAX+1];
    if(getcwd(cwd, PATH_MAX+1) != NULL){
        printf("Current Working Dir: %s\n", cwd);
    }
    else{
        printf("Error Occured!\n");
    }
    return 0;
}
```

# Change the working directory

## [cd]

<unistd.h>

It is a **built-in** command and it changes the current working directory.

The useful system call is `int chdir(const char *path);`

```
char buf[PATH_MAX+1];
char input[255];
if(getcwd(buf,PATH_MAX+1) != NULL){
    printf("Now it is %s\n",buf);
    printf("Where do you want to go?:");
    fgets(input,255,stdin);
    input[strlen(input)-1] = '\0';
    if(chdir(input) != -1){
        getcwd(buf,PATH_MAX+1);
        printf("Now it is %s\n",buf);}
    else{
        printf("Cannot Change Directory\n");
    }
}
```

Remember to remove  
the newline character!

# Phase 1 – Shell

---

You will need to write a workable “shell” that can

1. Get user inputs
2. Analyze the input command line and distinguish if it is built-in commands or others
3. If it is built-in commands:
  1. Change working directories, or exit the program
4. If it is other commands:
  1. Set the search path
  2. Expand the wildcard expression
  3. Fork a new process, set the signal handlers, and execute.
  4. Wait for the termination of the child.

# Set the search path

---

**REQUIREMENT:** (1) The shell allows file names (2) Search for the program

The sequence of the search path is determined by the environment variable **\$PATH**.

`/bin → /usr/bin → . (current directory)`

To do : set the environment variable to our desired search path

# setenv()

<stdlib.h>

```
int setenv(const char *name, const char *value, int overwrite)
```

This function can change the environment variable in the system temporarily until the termination of the process.

```
char *command1[] = {"shutdown", NULL};
printf("Running shutdown.. it is in /sbin :P \n\n");
setenv("PATH", "/bin:/usr/bin:.", 1);
execvp(*command1, command1);

if(errno == ENOENT)
    // Print No command found message
else
    // Print unknown error message
return 0;
```

Now the shell only searches the specified path 🐱

# Phase 1 – Shell

---

You will need to write a workable “shell” that can

1. Get user inputs
2. Analyze the input command line and distinguish if it is built-in commands or others
3. If it is built-in commands:
  1. Change working directories, or exit the program
4. If it is other commands:
  1. Set the search path
  2. Expand the wildcard expression
  3. Fork a new process, set the signal handlers, and execute.
  4. Wait for the termination of the child.



# Wildcard expansion

---

It is the shell's responsibility to expand to wildcard character in command lines.



```
$ ls *.c  
$ ls a.c b.c c.c
```

Requirements:

1. Expands the wildcard expansion into series of tokens, and they are in lexicographical order.
2. If the file is not found, keep the expression unchanged.

# glob()

---

To find the filenames with the given pattern, you can use glob().

glob() function utilizes a special structure called **glob\_t**

```
typedef struct {  
    size_t    gl_pathc;    /* Count of paths matched so far */  
    char      **gl_pathv;  /* List of matched pathnames. */  
    size_t    gl_offs;     /* Slots to reserve in gl_pathv. */  
} glob_t;
```



**Reserve 1 space for the command name.**

# glob()

---

```
#include <glob.h>

int glob(const char *pattern, int flags,
         int (*errfunc) (const char *epath, int eerrno),
         glob_t *pglob);
```

1. lexicographical order :
  - Be default it is sorted, no need to do extra things.
2. Place the command before the expanded tokens.
  - Use GLOB\_DOOFFS
3. Returns the original tokens upon no matches
  - Use GLOB\_NOCHECK
4. Call glob() on each token
  - Remember to add GLOB\_APPEND on the 2<sup>nd</sup> and onward call.

# glob()

```
1 #include <stdio.h>
2 #include <glob.h>
3 #include <unistd.h>
4
5 int main(int argc, char *argv[])
6 {
7     glob_t globbuf;
8
9     globbuf.gl_offs = 1;
10
11     glob("*.c", GLOB_DOOFFS | GLOB_NOCHECK, NULL, &globbuf);
12     glob("*.txt", GLOB_DOOFFS | GLOB_NOCHECK | GLOB_APPEND, NULL, &globbuf);
13     glob("*.java", GLOB_DOOFFS | GLOB_NOCHECK | GLOB_APPEND, NULL, &globbuf);
14     globbuf.gl_pathv[0] = "ls";
15
16     execvp(globbuf.gl_pathv[0], globbuf.gl_pathv);
17
18     return 0;
19 }
```



Remember  
APPEND!



Put the command name at front

# Phase 1 – Shell

---

You will need to write a workable “shell” that can

1. Get user inputs
2. Analyze the input command line and distinguish if it is built-in commands or others
3. If it is built-in commands:
  1. Change working directories, or exit the program
4. If it is other commands:
  1. Set the search path
  2. Expand the wildcard expression
  3. Fork a new process, set the signal handlers, and execute.
  4. Wait for the termination of the child.

# Signals

---

To emulate the real “shell”, your shell should handle some signals properly.

According to the specification, the following signals should be ignored.

Signals	Handling routines
SIGINT (Ctrl + C)	Ignore
SIGKILL (Default for kill)	Ignore
SIGQUIT (Ctrl + \)	Ignore
SIGTSTP (Ctrl + Z)	Ignore

However, your children should restore the above signal handling routines. **REMEMBER!!!!!!!!!!**

# Signals

```
<stdio.h>
<signal.h>
<unistd.h>
<sys/types.h>
<sys/wait.h>
```

```
1  int main(int argc, char *argv[]) {
2      signal(SIGINT, SIG_IGN);
3      signal(SIGQUIT, SIG_IGN);
4      signal(SIGTERM, SIG_IGN);
5      signal(SIGTSTP, SIG_IGN);
6      if(!fork()) {
7          signal(SIGINT, SIG_DFL);
8          signal(SIGQUIT, SIG_DFL);
9          signal(SIGTERM, SIG_DFL);
10         signal(SIGTSTP, SIG_DFL);
11         printf("[%d] I am Child..\n", getpid());
12         while(1) {}
13     }
14     else{
15         wait(NULL);
16         printf("[%d] I am super parent, kill me if you can\n", getpid());
17         while(1) {}
18     }
19     return 0;
20 }
```

Remember

# fork() + exec()

<unistd.h>

We know that we need to fork a new process to use the system call exec().

```
#include <stdio.h>
#include <unistd.h> // Needed By fork(),sleep()
int main(int argc,char *argv[]){
    while(1)
    {
        printf("\nPress Enter to execute ls...");
        while(getchar() != '\n');
        if(!fork()) {
            char *arglist[] = {"ls",NULL};
            execvp(*arglist,arglist);
        }
        else{
            sleep(1);
        }
    }
    return 0;
}
```

Is it okay? 🐻



# How to handle the child properly

<unistd.h>  
<sys/wait.h>  
<sys/types.h>

The shell should wait until the child dies and handle the event.

The useful system call is **pid\_t wait(int \*status);**

```
while(1) {
    printf("\nPress Enter to execute ls...");
    while(getchar() != '\n');
    pid_t child_pid;
    if(!(child_pid = fork())) {
        char *arglist[] = {"ls", NULL};
        execvp(*arglist, arglist);
    }

    else{
        wait(NULL);
    }
}
```

# Last Reminder

---

```
if(!fork()) {  
    char *arg[] = {"Hello", NULL};  
    execvp(*arg, arg);  
    printf("Error!\n");  
}  
else {  
    wait(NULL);  
}
```

In case of error in exec call(), exec will return and the program continues execution. You need to handle that.

# Enjoy : )

---

Phase 1 **DEADLINE:** 23:59, 2016 Mar 4 (Fri)

INDIVIDUAL work. ; )